Claude Certified Associate – Foundations

(CCAO-F) — Module 2 Practice Exam

Prompting & Task Execution

102 scenario-style and recall questions, multiple choice (A/B/C/D), one correct answer each.

Covers: Anatomy of an Effective Prompt (the five-component stack) · Task Decomposition for Complex Requests · Iterating to Improve Output · Adapting Strategy by Task Type (Analysis, Research, Drafting, Brainstorming).

Instructions: Select the single best answer for each question. An answer key with brief explanations is provided at the end of this document. This practice material is based on user-provided course screenshots and is intended for self-study only — it is not an official Anthropic exam product.

A. Module Introduction

1. In the AI Fluency Framework referenced by this module, what is the name of the competency covering 'telling Claude precisely what you want'?

A. Iteration
B. Description
C. Decomposition
D. Delegation

2. The module opens with two versions of the same request: 'write something about our Q3 results' versus a version specifying audience, key results, format, and length. What produced the better draft?

A. A more capable model was used for the second request
B. The prompt, not the model, changed — specification produced the improvement
C. The second request was run three times and the best version was chosen
D. The first request used a different entry point

3. According to the module's framing, where does consistent output quality come from?

A. Cleverness or luck
B. Prompt structure, not cleverness or luck
C. Always using the highest model tier
D. Running the same prompt many times and picking the best result

4. Which four skills does the module state cover the entire Prompting section of the exam?

A. Entry-point selection, model selection, Memory management, Skill creation
B. Learning the five components, decomposing complex work, iterating on the failed component, matching strategy to task type
C. Context management, usage limits, Incognito mode, Project scoping
D. Research mode, web search, citations, and source verification

5. Per the disclaimer in the module introduction, how should learners treat examples and scenarios used in the course?

A. As guaranteed to be based on real, named companies
B. As illustrative and often fictitious, not implying endorsement in either direction
C. As legally binding professional advice
D. As permanently accurate with no need to verify against current documentation

B. Anatomy of an Effective Prompt

6. What does 'naming the components' of a prompt turn prompting from, according to the lesson?

A. A creative writing exercise into a technical one
B. Guesswork into a checklist you can run before sending any non-trivial request
C. A one-step process into a multi-step process
D. An entry-point decision into a model decision

7. How many core components make up the prompt 'component stack' taught in this module?

A. Three
B. Four
C. Five
D. Six

8. Which component is defined as 'who you want Claude to be for this task,' setting the vocabulary, depth, and assumptions Claude brings?

A. Context
B. Role
C. Task
D. Constraints

9. Which component is described as 'the background Claude cannot know unless you provide it' and is the component professionals most often omit?

A. Role
B. Context
C. Constraints
D. Output format

10. Which component is 'the specific action, stated as a clear instruction,' with one primary verb stated unambiguously?

A. Task
B. Role
C. Output format
D. Constraints

11. Which component covers the boundaries — length, tone, what to include, what to leave out, what to avoid?

A. Output format
B. Constraints
C. Context
D. Role

12. Which component defines 'the shape of the result' — a table, a bulleted list, a three-paragraph memo, a draft email?

A. Task
B. Constraints
C. Output format
D. Context

13. Do all five components need to be present in every prompt?

A. Yes, every prompt requires all five or it will fail
B. No — a quick question needs only a task and maybe a constraint; a client deliverable needs all five
C. No, only Role and Task are ever necessary
D. Yes, but only when using Research mode

14. What is the 'Description' competency, as applied to the component stack?

A. The habit of making each component explicit instead of assuming Claude will infer it
B. A synonym for the Output format component
C. A feature that automatically writes prompts for the user
D. The process of describing a Project's knowledge base

15. Unless a source is connected through a Connector, what can Claude see?

A. Everything in the user's inbox and organization by default
B. Only what has been explicitly provided in the prompt or connected/allowed access
C. Only the most recent Project's knowledge base, regardless of permissions
D. The user's entire browsing history

16. What is described as 'the single most common reason a prompt underperforms for new users'?

A. Using the wrong model tier
B. A context gap — information that lives only in the user's head or outside a connected source
C. Using Chat instead of a Project
D. Requesting the wrong output format

17. If output is generic or off-base, which component gap is the most likely cause?

A. Missing context
B. An ambiguous task verb
C. Absent constraints
D. Wrong output format only

18. If output answers the wrong question entirely, which component gap is the most likely cause?

A. Missing context
B. An ambiguous task verb (the task component)
C. Absent constraints
D. Wrong role

19. If output is the wrong length or tone, which component gap is the most likely cause?

A. Missing role
B. Missing context
C. Absent constraints
D. Missing output format only, never constraints

20. In the worked build example, why did the weak prompt ('Write a summary of our quarterly operations.') produce an unusable result?

A. The model used was too weak for the task
B. The output was factually wrong
C. The prompt specified almost nothing — no audience, no figures, no format, no sense of what mattered
D. Claude refused to answer the weak prompt

21. In the strong version of the worked build prompt, what role was specified?

A. A financial analyst
B. An operations analyst
C. A policy reviewer
D. A marketing manager

22. In the strong version of the worked build prompt, what was the specified output format?

A. A five-paragraph essay
B. A single table with ten columns
C. A short headline followed by three bullet points, each one sentence
D. An unstructured stream-of-consciousness summary

23. In the worked build comparison, what changed between the weak and strong prompts?

A. The model and the data both changed
B. Same model, same data — only the specification (the prompt) changed
C. Only the model changed; the prompt stayed the same
D. The entry point changed from Chat to a Project

24. What is the recommended habit to build before sending any prompt that matters?

A. Always ask Claude to try harder
B. Run the five components in your head: role, context, task, constraints, and format
C. Switch to the most expensive model tier available
D. Send the prompt three times and average the results

25. According to the module, roughly how much time does thirty seconds of specification routinely save?

A. It has no measurable effect
B. Several rounds of correction
C. Exactly one round of correction, never more
D. It only saves time on Haiku

26. A marketer's prompt 'Write a post about our new feature' produces generic copy. Which single change will most improve the result?

A. Ask Claude to try harder and be more creative
B. Add context (audience, channel, the feature's key benefit) and the output format
C. Switch to a more capable model and resend the same prompt
D. Run the same prompt several times and pick the best output

C. Task Decomposition

27. Why do some requests need to be decomposed rather than specified as a single instruction?

A. Decomposition is always faster regardless of task size
B. Some requests are too large to specify as a single instruction; packing multiple stages into one prompt produces shallow work on every stage
C. Claude cannot process prompts longer than one sentence
D. Decomposition is only needed when using Research mode

28. How is decomposition defined in the lesson?

A. Asking Claude the same question multiple times for variety

B. Splitting a multi-part problem into discrete, ordered steps, then running them in sequence rather than asking for everything at once

C. Reducing a prompt down to a single word

D. Switching between multiple model tiers mid-task

29. In the single-prompt vendor evaluation example ('Evaluate these three vendors and tell me which to pick'), what is the main problem?

A. Claude refuses to answer the prompt at all

B. Claude must invent criteria, apply them, weigh trade-offs, and recommend all in one pass, doing all four shallowly with no visible reasoning

C. The prompt is too short to be processed

D. The prompt requires Code Execution to run at all

30. What is Step 1 in the decomposed vendor evaluation sequence?

A. Score vendors

B. Raise trade-offs

C. Derive criteria from the requirements document and weigh them

D. Recommend a vendor

31. What is Step 2 in the decomposed vendor evaluation sequence?

A. Derive criteria

B. Score each vendor against those criteria using the supplied materials

C. Recommend, with reasoning tied to weighted criteria

D. Raise trade-offs where vendors diverge most

32. What is Step 3 in the decomposed vendor evaluation sequence?

A. Derive criteria

B. Score vendors

C. Raise the trade-offs where vendors diverge most

D. Recommend a vendor with reasoning

33. What is Step 4 in the decomposed vendor evaluation sequence?

A. Score vendors

B. Raise trade-offs

C. Derive criteria

D. Recommend, with the reasoning tied back to the weighted criteria

34. What advantage does decomposition provide beyond simply improving output quality on each step?

A. It eliminates the need for any human review

B. Each step produces a checkable intermediate result, catching errors early and making the process auditable

C. It guarantees the model will never make a mistake

D. It removes the need to specify a Role or Context in any step

35. If the criteria derived in Step 1 of a decomposed task turn out to be wrong, when is this caught under a decomposed approach versus a single-prompt approach?

A. It is caught only after the final recommendation, same as a single prompt

B. It is caught before scoring, rather than after the recommendation

C. It cannot be caught at all under decomposition

D. It is caught only if the model tier is Opus

36. When should sequential steps that build on each other be kept in one conversation rather than split across several?

A. Never — every step should always be a separate conversation

B. When each step needs to see the prior results

C. Only when using the Haiku model

D. Only when Code Execution is disabled

37. According to the lesson, when should a task move to a separate conversation rather than continuing in the same one?

A. Only if the user changes their mind about the topic entirely

B. When a step is genuinely independent, or when the conversation has grown long enough that early context is degrading

C. Every time a new step begins, without exception

D. Only when switching from Chat to Research

38. In the 'Decompose a Parallel Case' scenario, a communications manager must turn a 20-page policy change into an internal announcement, an FAQ, and an executive briefing. What is Step 1 of the model decomposition?

A. Draft the staff announcement first

B. Extract the substantive changes from the policy document and what each one means in practice

C. Draft the FAQ for staff

D. Draft the executive briefing

39. In that same scenario, what is Step 2 of the model decomposition?

A. Draft the executive briefing

B. Confirm the extraction is complete and accurate before building anything on top of it

C. Draft the FAQ

D. Draft the staff announcement

40. In that same scenario, what are Steps 3 through 5?

A. Extract changes, confirm extraction, and recommend

B. Draft the staff announcement, draft the FAQ, and draft the executive briefing — in that order

C. Score vendors, raise trade-offs, and recommend

D. Derive criteria, weigh them, and finalize the policy

41. Why does the model decomposition extract and confirm the policy changes (Steps 1–2) before drafting any of the three deliverables?

A. Because executive briefings must always be written last regardless of content

B. Drafting any deliverable before the change list is confirmed risks propagating the same misreading into three documents; the shared, high-stakes extraction is sequenced first

C. Because Claude cannot draft more than one document per conversation

D. Because FAQs always take longer to write than announcements

42. An analyst needs Claude to evaluate five suppliers against a requirements document and recommend one. Which approach produces the most reliable result?

A. Ask for the evaluation and recommendation in a single prompt

B. Decompose into derive-criteria, score-each, identify trade-offs, then recommend

C. Ask Claude to recommend first, then justify the choice afterward

D. Run five separate conversations, one per supplier, with no shared criteria

43. What connects the decision of when to split a decomposed task into a separate conversation to earlier course material?

A. It connects to the model-selection lesson (Haiku vs. Sonnet vs. Opus)

B. It connects directly to the context-management skills from Module 1

C. It connects to the Incognito mode lesson

D. It has no connection to any prior lesson

44. Which best describes why a vendor evaluation is called 'four tasks wearing one sentence'?

A. Because it literally requires four separate documents to be uploaded

B. Because a single request ('evaluate these vendors and tell me which to pick') actually bundles deriving criteria, scoring, weighing trade-offs, and recommending into one instruction

C. Because it needs four different model tiers to complete

D. Because it must be run in exactly four separate Projects

45. What is the primary risk of packing a genuinely multi-stage task into a single prompt?

A. The prompt will be rejected by Claude automatically

B. Shallow execution on every stage, with the reasoning behind any conclusion hidden from the user

C. It becomes impossible to use Code Execution afterward

D. It automatically triggers Incognito mode

D. Iterating to Improve Output

46. What is the core skill taught in this lesson, as distinguished from simply rewriting a weak prompt from scratch?

A. Always regenerating with a different model tier

B. Reading the output to diagnose which component fell short, then fixing that one thing

C. Sending the same prompt repeatedly until the output improves by chance

D. Abandoning any prompt that doesn't work on the first try

47. According to the deficiency-diagnosis table, what is the likely cause when output is generic or off-base?

A. The task verb was ambiguous
B. The context was thin
C. A constraint was missing
D. The wrong output format was requested

48. According to the deficiency-diagnosis table, what is the likely cause and fix when output answers the wrong question?

A. The context was thin; add background
B. The task verb was ambiguous; sharpen the instruction
C. A constraint was missing; add it
D. Nothing can be done; discard and restart

49. According to the deficiency-diagnosis table, what is the fix when output is the wrong length, tone, or shape?

A. Sharpen the task instruction
B. Add the missing constraint or format
C. Switch to a different entry point
D. Increase the model tier

50. According to the deficiency-diagnosis table, what should be done when output is close but misses on one section?

A. Discard the whole draft and start over with a new prompt
B. Iterate on that section only; do not discard a draft that is mostly right
C. Switch models and regenerate everything from scratch
D. Ignore the miss since it is a minor issue

51. Why does rewriting the entire prompt from scratch tend to be counterproductive when only one component was off?

A. It always produces a shorter output
B. It loses the parts that worked, and makes it impossible to tell which change fixed the problem
C. It automatically switches the model tier
D. It removes the ability to iterate further

52. What is the recommended targeted-revision approach when output disappoints?

A. Change the one component the output pointed to, resend, and compare
B. Regenerate using a completely different task and role
C. Always increase the constraint list to the maximum
D. Switch to Research mode regardless of the task

53. In the live iteration cycle example, what was diagnosed as missing after Round 1's prompt ('Write a follow-up email to the client about the delayed deliverable')?

A. The role and the output format
B. Context (no reason, no new date) and a tone constraint
C. The task verb
D. Nothing — Round 1 already produced an ideal result

54. In that same example, what was diagnosed as still missing after Round 2, which added the cause, new date, and a tone constraint?

A. The context was still too thin
B. Only the subject line was missing
C. The role was undefined
D. The task verb was still ambiguous

55. How many rounds did the live iteration cycle example take, and what changed in each round?

A. Five rounds, rewriting the entire prompt each time
B. Three rounds, each changing exactly the component the previous output exposed
C. One round, since the first prompt was already correct
D. Ten rounds with random variations

56. What signals that iteration has converged and it's time to stop?

A. Each additional round produces marginal change rather than improvement
B. The output has been rewritten at least five times
C. The model tier has been upgraded to Opus
D. The user has spent more than ten minutes on the prompt

57. What is the actual goal of the iteration process, according to the lesson on knowing when to stop?

A. A theoretically perfect prompt
B. A usable result
C. The shortest possible prompt
D. A prompt that requires zero human review

58. What broader skill, covered in depth in Module 7, uses the same diagnostic discipline as prompt iteration?

A. Choosing between model tiers
B. Troubleshooting any underperforming workflow
C. Setting up Incognito mode
D. Creating a new Project from scratch

59. A drafted memo came back accurate but far too long and too formal for its audience. What is the most efficient next step?

A. Rewrite the entire prompt from scratch
B. Switch models and regenerate
C. Adjust the constraint and format components (length and tone) and resend
D. Manually cut the memo down and abandon the prompt

60. Why are output deficiencies described as 'prompt diagnostics' rather than simply failures?

A. Because each kind of disappointing output points back to a specific, identifiable component gap that can be fixed

B. Because Claude labels every output with an error code

C. Because the model automatically explains what went wrong

D. Because deficiencies always indicate the wrong model tier was used

61. What is lost when a person discards a mostly-correct draft and starts over with an entirely new prompt after one section falls short?

A. Nothing is lost; this is always the most efficient approach

B. The parts of the draft that were already working, and the ability to isolate which change fixes the remaining issue

C. Only the output format, never the content

D. The ability to use Code Execution going forward

62. Which statement best summarizes the relationship between rewriting and targeted revision?

A. They are functionally identical approaches

B. Targeted revision changes only the diagnosed component; wholesale rewriting risks losing what already worked and obscures which change caused any improvement

C. Wholesale rewriting is always faster and therefore always preferred

D. Targeted revision is only possible when using Opus

63. In the live iteration cycle, what was requested as the final Round 3 change?

A. A complete rewrite of the email's content

B. A subject line that signals resolution, not just delay

C. A switch to a more formal tone

D. An extension of the word count limit

64. What word count constraint was specified in Round 2 of the live iteration example?

A. Under 50 words

B. Under 120 words

C. Under 300 words

D. No word limit was specified

E. Adapting Strategy by Task Type

65. What does the lesson mean by 'the component stack applies to every prompt, but the emphasis shifts with the task'?

A. A different set of components is used for each task type

B. The same five components apply everywhere, but which ones need to be tightened versus loosened changes by task type

C. Only Analysis and Research use the component stack; Drafting and Brainstorming do not

D. The component stack is replaced entirely for creative tasks

66. What does Analysis want, in terms of specificity and creative latitude?

A. Loose constraints and high creative latitude

B. Tight constraints and explicit criteria; low creative latitude, high specification

C. No constraints at all, since analysis should be exploratory

D. Only a Role component, nothing else

67. What does Research want, according to the strategy lesson?

A. High creative latitude and loose scope

B. Clear scope and source discipline: a defined question, defined boundaries, and clarity on whether current sources are required

C. No need to define a question, only a topic

D. Exclusively deep multi-source investigation, never quick lookups

68. What does Drafting want, in terms of control versus latitude?

A. Full control over every word choice, with no latitude at all

B. Audience, tone, and format specified, with room for Claude to find the phrasing — medium latitude

C. No specification of audience or tone, only a topic

D. The same tight constraints as Analysis

69. What does Brainstorming want, according to the strategy lesson?

A. Tight constraints and low creative latitude, the same as Analysis

B. Loose constraints and high latitude; over-specifying kills the divergence being sought

C. No goal statement at all, only constraints

D. The exact same treatment as Research

70. In the 'Strategy quick reference' table, what should be tightened for Analysis tasks?

A. Word choice

B. Criteria, standards, scope

C. Quantity and direction

D. Synthesis approach

71. In the 'Strategy quick reference' table, what should be loosened for Analysis tasks?

A. Criteria

B. Scope

C. Phrasing

D. Standards

72. In the 'Strategy quick reference' table, what should be tightened for Research tasks?

A. Word choice and tone

B. Question, sources, citations

C. Quantity and direction

D. Goal and guardrails only

73. In the 'Strategy quick reference' table, what should be loosened for Research tasks?

A. The synthesis approach
B. The citation requirement
C. The scope of sources allowed
D. The definition of the question

74. In the 'Strategy quick reference' table, what should be tightened for Drafting tasks?

A. Criteria and standards
B. Audience, tone, and format
C. Goal and guardrails only
D. Quantity and direction

75. In the 'Strategy quick reference' table, what should be loosened for Drafting tasks?

A. Format
B. Tone
C. Word choice
D. Audience

76. In the 'Strategy quick reference' table, what should be tightened for Brainstorming tasks?

A. Criteria, standards, and scope, same as Analysis
B. Goal and guardrails only
C. Question, sources, and citations
D. Audience, tone, and format

77. In the 'Strategy quick reference' table, what should be loosened for Brainstorming tasks?

A. Quantity and direction
B. Sources and citations
C. Audience and tone
D. Criteria and standards

78. In the Analysis mini-demo ('Compare these two vendor contracts on payment terms, termination rights, and liability caps...'), what characterizes the prompt?

A. Loose criteria and an undefined output shape
B. Tight criteria and a defined output (a three-row table), with no room to wander
C. An open-ended request with no comparison points specified
D. A request for the widest possible range of contract interpretations

79. In the Research mini-demo, what is emphasized about citations produced from Claude's training memory alone, versus those from a grounded source?

A. They are equally reliable and never need independent verification
B. Citations from training memory alone can look equally confident as grounded citations, but should be independently verified
C. Training-memory citations are always more accurate than web-search citations
D. Citations are never needed for Research tasks

80. In the Drafting mini-demo ('Draft a 150-word LinkedIn post announcing our new reporting feature, aimed at operations managers, in a confident but not salesy voice'), what is fixed and what is left open?

A. Everything is fixed, including exact phrasing
B. Audience, length, and tone are fixed; phrasing is left open
C. Nothing is fixed; the whole prompt is open-ended
D. Only the topic is fixed; everything else is open

81. In the Brainstorming mini-demo ('Give me 20 angles for a campaign around faster month-end close. Range widely; do not self-edit yet.'), what does this illustrate?

A. Tight constraints and low latitude, same as Analysis
B. Only the goal and one guardrail are set; constraints come later, after the range exists
C. The need for citations and source discipline
D. A request for a single best idea to save time

82. What is the risk of using one fixed prompting style across Analysis, Research, Drafting, and Brainstorming tasks?

A. There is no risk; one style works equally well for all four
B. It costs quality on every task it does not fit
C. It only affects Drafting tasks, not the other three
D. It automatically triggers a model downgrade

83. A team wants the widest possible range of campaign concepts before narrowing. How should the prompt be calibrated?

A. Apply tight constraints on length, format, and terminology up front
B. Give the goal and a few guardrails, then ask for volume and range before filtering
C. Request a single best idea to save time
D. Specify the exact creative direction so Claude stays on message

84. What is the 'underlying move' that stays the same across all four task types, according to the lesson's conclusion?

A. Always applying the maximum number of constraints possible
B. Deciding where control is needed and where range is needed, then setting constraints accordingly
C. Always using Research mode regardless of task type
D. Ignoring the component stack for creative tasks

85. What does the lesson say separates a competent prompter from one who gets the same mediocre output on every task?

A. Using the most expensive model tier available
B. Matching strategy to task type
C. Writing the longest possible prompt every time
D. Avoiding the use of constraints entirely

F. Checkpoint, Exercise & Mixed Review

86. In the checkpoint exercise, the prompt '"Make this better." (pasted alongside a draft email)' is missing which components most obviously?

A. Only Role

B. Context, Task specificity, Constraints, and Output format are all essentially undefined

C. Only Output format

D. Nothing is missing; the prompt is complete

87. In the checkpoint exercise, the prompt '"Give me everything you know about supply chain risk."' is most clearly missing which component?

A. Role only

B. Constraints and Task specificity — the request is unbounded with no clear deliverable shape

C. Output format only, nothing else

D. Context only

88. In the checkpoint exercise, the prompt '"Write a professional document about the project."' most clearly illustrates a gap in which components?

A. Role and Constraints only, nothing else

B. Context (what project, what audience), Task specificity, and Output format

C. Only the Output format is missing

D. No components are missing

89. In the checkpoint exercise, the prompt '"Brainstorm names, but each must be one word, under eight letters, avoid these twelve terms, and match our exact brand voice guidelines"' illustrates what kind of problem?

A. Missing context only

B. Over-constraining a Brainstorming task, which kills the divergence brainstorming is supposed to produce

C. A perfectly calibrated prompt with no issues

D. Missing Role and Task components

90. In the 'Repair the Underperforming Prompt' exercise, the weak prompt was 'Summarize the customer feedback and tell me what to do.' What output did it produce?

A. A precise, ranked list of the top three issues with representative quotes

B. A generic five-bullet list of themes with vague advice, nothing tied to the actual data, no priorities, nothing actionable

C. An error message, since the prompt was too short to process

D. A fully formatted executive report with citations

91. In that same exercise, what was the author's actual goal?

A. A generic overview of customer sentiment with no ranking

B. The top three issues by frequency from 200 survey responses, each with a representative quote, ranked so she can decide what to fix this quarter

C. A single-sentence summary with no supporting detail

D. A comparison of customer feedback against a competitor's feedback

92. In that same exercise, which component(s) did the weak prompt fail to specify, based on the gap between the weak output and the stated goal?

A. Role and Output format were the only two problems

B. Context (the 200 responses and their nature), Task specificity (top three by frequency, ranked), and Output format (quote-supported, ranked list) were all under-specified

C. Nothing was missing; the output already matched the goal

D. Only the Constraints component was missing

93. In the exercise's fragment-matching stage, which fragment was explicitly identified as 'does not belong' to any of the five components?

A. 'You are a product analyst'

B. 'Attached are 200 customer survey responses'

C. 'Make the summary sound professional and polished'

D. 'Format as a ranked list, most frequent first'

94. In the exercise's fragment-matching stage, which fragment best fits the Task component?

A. 'You are a product analyst'

B. 'Attached are 200 customer survey responses'

C. 'Identify the three most frequently raised issues, ranked by how many responses mention each'

D. 'Format as a ranked list, most frequent first'

95. In the exercise's fragment-matching stage, which fragment best fits the Constraints component, and what does it specify?

A. 'You are a product analyst' — specifies the role only

B. 'For each issue include one representative verbatim quote and the approximate share of responses it appears in; use code execution to count accurately rather than estimating' — specifies boundaries and a verification requirement

C. 'Attached are 200 customer survey responses' — specifies context only

D. 'Format as a ranked list, most frequent first' — specifies only the output shape

96. According to Key Takeaway 1 of Module 2, what carries almost every professional prompt?

A. Cleverness and creative phrasing

B. Five components: role, context, task, constraints, format — run before sending anything that matters

C. The choice of model tier alone

D. The length of the prompt

97. According to Key Takeaway 2 of Module 2, what is the most common cause of generic output?

A. A model limitation

B. A context gap, not a model limitation

C. An overly long prompt

D. Using Chat instead of a Project

98. According to Key Takeaway 3 of Module 2, when do multi-stage requests succeed?

A. When they are compressed into a single instruction

B. When each step produces a checkable result and the high-stakes foundation is built first

C. Only when Research mode is used

D. Only when a Project is used instead of Chat

99. According to Key Takeaway 4 of Module 2, how should iteration be approached?

A. Rewrite the whole prompt every round regardless of the diagnosis

B. Read the output as a diagnostic, change the one thing it points to, and stop when rounds stop improving the result

C. Always run at least five iteration rounds regardless of quality

D. Iterate only when using the Opus model

100. According to Key Takeaway 5 of Module 2, what should be decided for each task, and then acted on?

A. Which model tier is cheapest

B. Where control is needed and where range is needed, then set constraints to fit — since analysis wants constraints and brainstorming wants latitude

C. Whether to use Chat or a Project

D. Whether Memory should be enabled

101. A single prompt is used to both decompose a complex vendor evaluation AND apply the wrong task-type strategy (treating it like a Brainstorming task with loose constraints). What would this combination most likely produce?

A. A perfectly reasoned, auditable recommendation

B. A shallow or unfocused result, since Analysis-type tasks need tight criteria and standards, which loose constraints would undermine

C. No output at all, since Claude cannot process mismatched strategies

D. The same result as a properly decomposed and correctly calibrated prompt

102. Which best summarizes the overall throughline connecting the Anatomy of a Prompt, Task Decomposition, Iteration, and Adapting Strategy lessons in Module 2?

A. Each lesson is unrelated and covers a completely separate exam domain

B. A shared component framework (role, context, task, constraints, format) underlies specification, is applied stage-by-stage in decomposition, is diagnosed and repaired through iteration, and is weighted differently depending on task type

C. Only model selection determines prompt quality; the components are optional

D. Decomposition and iteration are the same skill taught twice

Answer Key & Explanations

1. B — Description is named as the competency of telling Claude precisely what you want, and it is the backbone of Module 2.

2. B — The module explicitly states that the model did not get smarter between the two requests — the prompt did; specification produced the improved draft.

3. B — The module states plainly that consistent output quality comes from prompt structure, not cleverness or luck.

4. B — The module states that learning the five components, decomposing complex work, iterating on the component that failed, and matching strategy to task type cover the entire Prompting section of the exam.

5. B — The disclaimer states examples/scenarios are illustrative and often fictitious; if a company or product is mentioned, it doesn't mean they endorse Anthropic or vice versa.

6. B — Naming the components turns prompting from guesswork into a checklist that can be run before sending any non-trivial request.

7. C — Five components — Role, Context, Task, Constraints, and Output format — carry almost all the weight in a professional prompt.

8. B — Role defines who you want Claude to be for the task — e.g., a financial analyst, editor, or policy reviewer — setting vocabulary, depth, and assumptions.

9. B — Context is the background Claude cannot infer — audience, situation, prior decisions, source material — and it's the component most often left out.

10. A — Task is the specific action — 'summarize,' 'compare,' 'draft,' 'identify' — stated as one clear, unambiguous primary verb.

11. B — Constraints are the boundaries that keep the output usable without heavy editing: length, tone, inclusions, exclusions, and things to avoid.

12. C — Output format defines the shape of the result; stating it up front saves an iteration round.

13. B — Not every prompt needs all five; the skill is noticing which components a given task requires — a quick question may need just a task and maybe a constraint, while a client deliverable needs all five.

14. A — Description is the habit of making each component explicit rather than assuming Claude will infer it.

15. B — Claude cannot see the user's inbox, org chart, or last week's meeting unless connected through a Connector, and even then it sees only what has been allowed.

16. B — Context gaps — information that lives only in the user's head or outside a connected source — are named as the single most common reason a prompt underperforms for new users.

17. A — Missing context typically produces generic output, since Claude lacks the background needed to make the response specific.

18. B — An ambiguous task verb tends to produce the wrong action/answer, since the specific instruction wasn't clearly stated.

19. C — Absent constraints (length, tone, boundaries) typically produce output of the wrong length or tone.

20. C — The output was not wrong, but unusable, because the prompt specified almost nothing about audience, figures, format, or priorities.

21. B — The strong prompt example specifies: 'You are an operations analyst,' establishing the Role component.

22. C — The strong prompt specifies formatting as a short headline followed by three bullet points, each one sentence.

23. B — Same model, same data — the difference was entirely in the specification, i.e., making the five components explicit.

24. B — The habit to build is running the five components in your head before sending: role, context (what Claude cannot infer), task, constraints, and output format.

25. B — The module states that thirty seconds of specification routinely saves several rounds of correction.

26. B — Generic output signals a context and format gap; adding audience/channel/key-benefit context and specifying the output format directly addresses the cause.

27. B — When a task has several distinct stages, packing it into one prompt produces shallow work on every stage — decomposition breaks it into a sequence Claude can execute well.

28. B — Decomposition means splitting a multi-part problem into discrete, ordered steps and running them in sequence, rather than requesting everything in a single pass.

29. B — Packed into a single prompt, Claude has to invent criteria, apply them, weigh trade-offs, and recommend all at once — doing all four shallowly, with the reasoning hidden.

30. C — Step 1 is 'Derive criteria': from the requirements document, derive the evaluation criteria that matter and weigh them.

31. B — Step 2 is 'Score vendors': score each vendor against the derived criteria using the supplied materials.

32. C — Step 3 is 'Raise trade-offs': surface where vendors diverge most so the decision-maker can see the trade-offs explicitly.

33. D — Step 4 is 'Recommend': make the recommendation, with reasoning explicitly tied back to the weighted criteria derived in Step 1.

34. B — Decomposition produces a checkable intermediate result at each stage, so errors (e.g., wrong criteria) are caught before scoring rather than after the recommendation, and the process becomes auditable.

35. B — Because each step produces a checkable intermediate result, a flawed Step 1 (wrong criteria) is caught before scoring, not after the final recommendation.

36. B — Keep sequential steps that build on each other in one conversation so each step can see the prior results.

37. B — Move to a separate conversation when a step is genuinely independent, or when the conversation has grown long enough that early context is degrading — connecting directly to Module 1's context-management skills.

38. B — Step 1 is extracting the substantive changes from the policy document and what each one means in practice — the shared foundation for all three deliverables.

39. B — Step 2 confirms the extraction is complete and accurate before any of the three deliverables are built on top of it.

40. B — Steps 3–5 are: draft the staff announcement (general audience), draft the FAQ (anticipating staff questions), and draft the executive briefing (compressed to decisions and impact).

41. B — Steps 1 and 2 build a verified foundation that all three deliverables draw on; drafting first risks propagating the same misreading into three separate documents, so the shared extraction is sequenced first.

42. B — This mirrors the vendor-evaluation decomposition pattern: deriving criteria first, then scoring, then trade-offs, then a reasoned recommendation, produces auditable, higher-quality results.

43. B — The lesson explicitly states that the judgment of when to split into a new conversation connects directly to the context-management skills taught in Module 1.

44. B — The phrase reflects that the seemingly simple request actually bundles four distinct stages of work (derive, score, weigh, recommend) into a single sentence.

45. B — Packing a multi-stage task into one prompt produces shallow work on every stage and hides the reasoning behind the final output.

46. B — The lesson frames iteration as diagnosis: reading the output to identify which specific component fell short, then fixing only that one thing.

47. B — Generic or off-base output is attributed to thin context — add the background Claude could not infer.

48. B — When output answers the wrong question, the likely cause is an ambiguous task verb, and the fix is to sharpen the instruction.

49. B — Wrong length, tone, or shape points to a missing constraint or output-format component — the fix is to add it.

50. B — When output is close but misses in one section, the guidance is to iterate on that section only rather than discarding a mostly-correct draft.

51. B — Rewriting the whole prompt loses the parts that worked and makes it impossible to isolate which change actually fixed the issue.

52. A — The recommended approach is to change only the specific component the output's flaw points to, resend, and compare results.

53. B — Round 1's output was generic and defensive; the diagnosis was thin context (no reason, no new date) and no tone constraint.

54. B — Round 2 produced a strong, accountable note; the only remaining gap was the missing subject line.

55. B — Three rounds were used, with each round changing exactly the component that the previous round's output had exposed as weak.

56. A — Iteration has converged when additional rounds yield only marginal change; at that point, further prompting produces less benefit than a quick manual edit.

57. B — The lesson states explicitly: the goal is a usable result, not a perfect prompt — recognizing diminishing returns is part of the skill.

58. B — The lesson notes that the same diagnostic discipline used in prompt iteration applies to troubleshooting any underperforming workflow, covered in depth in Module 7.

59. C — Wrong length and tone point specifically to the constraints/format components; adjusting just those and resending is the targeted fix.

60. A — Each kind of disappointment (generic output, wrong question answered, wrong length/tone, one weak section) points back to a specific component, making the output itself a diagnostic tool.

61. B — Discarding a mostly-correct draft loses the working parts and removes the ability to identify exactly which change would have fixed the one weak section.

62. B — Targeted revision isolates and fixes the diagnosed component, while wholesale rewriting risks losing functioning parts of the prompt and obscuring cause and effect.

63. B — Round 3's prompt asked to add a subject line that signals resolution rather than just delay — the one remaining diagnosed gap.

64. B — Round 2's prompt specified keeping the note under 120 words, in an accountable, not over-apologetic tone.

65. B — The same five-component stack underlies every prompt; what changes by task type is the balance of specificity versus creative latitude — which components to tighten or loosen.

66. B — Analysis wants tight constraints and explicit criteria — what to measure, against what standard, how to handle ambiguity — with low creative latitude and high specification.

67. B — Research wants clear scope and source discipline — the question, the boundaries, and whether current sources are required — with citations requested so claims are checkable.

68. B — Drafting wants audience, tone, and format specified (medium latitude): the user controls the shape, Claude fills in the phrasing.

69. B — Brainstorming wants loose constraints and high latitude — give the goal and boundaries, then ask for volume and range before narrowing, since over-specifying kills divergence.

70. B — For Analysis, the table lists criteria, standards, and scope as what to tighten.

71. C — For Analysis, phrasing is listed as what to loosen — Claude has latitude in wording even though criteria/standards/scope stay tight.

72. B — For Research, the question, sources, and citations are what should be tightened.

73. A — For Research, the synthesis approach is listed as what to loosen, while the question, sources, and citations stay tight.

74. B — For Drafting, audience, tone, and format are what should be tightened.

75. C — For Drafting, word choice is what should be loosened — Claude has room to find the actual phrasing within the specified audience/tone/format.

76. B — For Brainstorming, only the goal and guardrails should be tightened — everything else stays loose to preserve divergence.

77. A — For Brainstorming, quantity and direction should be loosened, since the goal is a wide range of ideas before narrowing.

78. B — The Analysis demo has tight, named criteria and a clearly defined output shape (a three-row table) — no room to wander.

79. B — The lesson explicitly notes that citations produced from training memory alone can look just as confident as grounded ones, and should be independently verified — checkable citations come from a grounded source like web search or Research results.

80. B — Audience, length, and tone are fixed in the Drafting demo, while the actual phrasing is left open for Claude to determine.

81. B — This demo shows only a goal and one guardrail (range widely, don't self-edit) being set — constraints are deliberately deferred until after the range of ideas exists.

82. B — Using one fixed style across all four task types costs quality on every task it does not fit, since each rewards a different balance of specificity and latitude.

83. B — This is a Brainstorming task: loose constraints, high latitude — set the goal and guardrails, then ask for volume/range before narrowing.

84. B — The lesson concludes that the same underlying move applies every time: decide where you need control and where you need range, then set constraints accordingly.

85. B — Matching strategy to task type is explicitly named as what separates a competent prompter from someone who applies one fixed approach and gets mediocre results everywhere.

86. B — 'Make this better' leaves the task vague, provides no context on what 'better' means, no constraints, and no defined output format.

87. B — This prompt lacks constraints (scope/boundaries) and a specific task; 'everything you know' is unbounded and produces unfocused output.

88. B — 'The project' with no further detail leaves context undefined, the task vague ('a document'), and the output format unspecified.

89. B — This over-specifies a Brainstorming task with heavy constraints up front, which — per the Adapting Strategy lesson — undermines the wide range brainstorming is meant to produce.

90. B — The weak prompt produced generic themes and vague advice, disconnected from the actual data, without priorities or actionable specifics.

91. B — The product manager's real goal was the top three issues by frequency, each with a representative quote, ranked to inform this quarter's priorities.

92. B — The weak prompt lacked adequate context about the dataset, a specific enough task (top three by frequency, ranked), and a defined output format (with representative quotes) — all needed to reach the stated goal.

93. C — 'Make the summary sound professional and polished' is the fragment that does not map cleanly to Role, Context, Task, Constraints, or Output format — it's vague filler rather than a specification.

94. C — This fragment states the specific, unambiguous action to perform — identifying and ranking the top three issues — matching the Task component.

95. B — This fragment defines boundaries on what to include (a quote, an approximate share) and explicitly constrains the method (use code execution rather than estimating) — a Constraints-type specification.

96. B — Takeaway 1 states that five components — role, context, task, constraints, format — carry almost every professional prompt and should be run before sending anything that matters.

97. B — Takeaway 2 states that Claude cannot see what lives only in the user's head, and the most common cause of generic output is a context gap, not a model limitation.

98. B — Takeaway 3 states multi-stage requests succeed when each step produces a checkable result and the high-stakes foundation is built first.

99. B — Takeaway 4 states: read the output as a diagnostic, change the one thing it points to, and stop when rounds stop improving the result.

100. B — Takeaway 5 states: analysis wants constraints, brainstorming wants latitude — decide where control and where range are needed, then set constraints to fit.

101. B — Vendor evaluation is an Analysis-type task needing tight criteria/standards; applying loose Brainstorming-style constraints would work against the specificity that task type requires, undermining quality even if decomposed.

102. B — The module builds a single coherent framework: the five components define what a good prompt specifies, decomposition applies that framework stage-by-stage to complex tasks, iteration diagnoses and repairs specific component gaps, and task-type strategy adjusts the emphasis placed on each component.