

Claude Certified Associate – Foundations (CCAO-F) — Module 3 Practice Exam

Output Evaluation & Validation

100 scenario-style and recall questions, multiple choice (A/B/C/D), one correct answer each.

Covers: Discernment — Evaluating Accuracy, Completeness and Fitness · Hallucinations, Inconsistencies & Bias · Fact-Checking and Grounding Techniques · Diligence — When Human Review Is Non-Negotiable · Editing and Adapting Output for Your Audience · Choosing Output Formats (Inline, Artifacts, Structured, Code-Executed).

Instructions: Select the single best answer for each question. An answer key with brief explanations is provided at the end of this document. Correct answers are distributed evenly across A, B, C and D. This practice material is based on user-provided course screenshots and is intended for self-study only — it is not an official Anthropic exam product.

A. Module Introduction

1. The module opens by naming an asymmetry in professional AI use. What is it?
 - A. Errors are large and visible immediately, while the time saved is hidden and hard to measure
 - B. The time saved grows with the length of the output, while the risk of error shrinks
 - C. Verification always costs more time than the drafting it checks
 - D. The time saved is small and visible, while the cost of an unverified error is large and arrives later
2. In the module's cautionary tale, what happened to the five supporting statistics Claude returned for a client market-sizing deck?
 - A. Four were sound and one growth rate was fabricated; it reached the client deck and was flagged in the room by the client's own analyst
 - B. All five were fabricated, which is why the deck was rejected outright
 - C. The consultant caught the fabricated figure during review, so nothing reached the client
 - D. Claude declined to supply the statistics because no source was connected
3. According to the cautionary tale, when is the cost of a missed error actually paid?
 - A. Only if the error is large enough to be legally actionable
 - B. At the moment the time is saved, which is what makes the trade-off easy to see
 - C. Later, by someone whose trust you need — not at the moment the time is saved
 - D. Never, provided the output came from a high-capability model
4. Why does the module say output evaluation is the largest section of the certification exam?
 - A. Because evaluation requires more technical configuration than any other topic
 - B. Accountability — when you put your name on a deliverable you own every claim in it, whether you wrote the words or Claude did
 - C. Because hallucination is the only failure mode Claude exhibits
 - D. Because the exam weights each section by how many product features it covers
5. How is the Discernment competency defined in this module?
 - A. The habit of making each prompt component explicit rather than assuming inference
 - B. The skill of deciding when verification is required and taking responsibility for the result
 - C. The skill of critically evaluating output against requirements, sources, and standards
 - D. The practice of choosing the entry point before writing the prompt

6. How is the Diligence competency defined in this module?
- A. Critically evaluating output against requirements, sources, and standards
 - B. Deciding when verification is required and taking responsibility for the result
 - C. Splitting a multi-part request into discrete, ordered steps
 - D. Reviewing stored memories on a regular cadence and deleting expired entries
7. Which of the following is NOT one of the module's stated learning objectives?
- A. Configure account-level Memory so verification results persist across every future session
 - B. Identify hallucinations, inconsistencies, and biases in output
 - C. Determine when human review or additional verification is required
 - D. Select appropriate output formats and organize information for optimal results

B. Discernment: Accuracy, Completeness, and Fitness

8. How does the lesson define evaluation?
- A. A comparison of two Claude outputs to see which one reads better
 - B. An intuition that sharpens with experience and eventually replaces formal checks
 - C. A single pass to confirm the output answers the question that was asked
 - D. Not a feeling about whether output looks good, but a check against three fixed references
9. What are the three evaluation references in the Discernment protocol?
- A. Requirements, model tier, and output format
 - B. Accuracy, completeness, and formatting
 - C. Stakes, reversibility, and audience
 - D. Requirements, source material, and professional standards
10. What does checking output against the Requirements reference involve?
- A. Comparing the output against a previous output on the same topic
 - B. Asking Claude to restate the request to confirm it understood
 - C. Re-reading your own request and confirming each part is addressed, not just the easy parts
 - D. Confirming the output is formatted the way your team prefers
11. What does the Source material reference require when the output relies on documents you supplied?
- A. Re-uploading the documents and running the same prompt a second time
 - B. Tracing specific claims back to the source rather than trusting that Claude read carefully
 - C. Checking that the output cites the same number of documents you uploaded
 - D. Confirming the summary is shorter than the source material
12. Which of these is given as an example of failing professional standards even when the output reads fluently?
- A. A number without units, a recommendation without reasoning, or a citation you cannot locate
 - B. An output that is longer than the request implied
 - C. An output that uses bullet points where prose was expected
 - D. An output produced on a lower model tier than the task warranted
13. Why does the lesson insist on running the same check the same way every time?
- A. Because a varied review process consumes more of the usage limit
 - B. Because Claude produces identical output whenever a check is repeated
 - C. Because the exam requires the three references to be checked in a fixed order
 - D. So quality does not depend on how rushed you happen to be

14. According to stakes calibration, how universal are review stakes?
- A. They depend only on the length of the output
 - B. They are universal: every output deserves the same depth of review
 - C. They are domain-dependent, not universal — depth of review follows what is at stake in that field
 - D. They depend on which model tier produced the output
15. In zero-tolerance work such as legal analysis, financial figures, and compliance reporting, what does the lesson prescribe?
- A. Accuracy takes priority over speed entirely, and every claim gets verified
 - B. A lighter review, since these domains have external safeguards of their own
 - C. The same review depth as internal brainstorming, simply carried out faster
 - D. Verification of the formatting only, since the figures come from source documents
16. What risk does the stakes-calibration section name?
- A. Spending too long reviewing internal brainstorming output
 - B. Applying the same casual review to high-stakes and low-stakes work alike
 - C. Reviewing regulated work twice, which duplicates effort
 - D. Choosing an output format before the stakes are known
17. What are the three verdicts in the module's output triage?
- A. Accurate, incomplete, unfit
 - B. Ready to use, needs revision, needs human override
 - C. Ship, iterate, discard
 - D. Low stakes, medium stakes, high stakes
18. Under the triage, when does an output qualify as 'needs revision'?
- A. Any output that has not yet been read from beginning to end
 - B. The stakes or the uncertainty mean it should not go out on Claude's draft alone
 - C. It meets requirements, matches sources, and clears professional standards
 - D. It is close, but a specific gap remains — note the gap and iterate
19. When does an output need human override?
- A. Whenever the output was produced without web search enabled
 - B. Whenever a single formatting inconsistency appears anywhere in the output
 - C. When the stakes, the errors, or the uncertainty mean it should not go out on Claude's draft alone
 - D. Whenever one round of iteration has already been attempted
20. What does the 'ready to use' verdict require?
- A. It meets the requirements, matches the sources, and clears professional standards
 - B. It reads fluently and contains no obvious errors
 - C. It was produced from documents the user uploaded personally
 - D. It has been through at least three rounds of iteration
21. How do accuracy and completeness differ, and how do they fail?
- A. Completeness is a subset of accuracy, so an accurate output is complete by definition
 - B. Accuracy asks whether what is present is correct; completeness asks whether anything is missing — and they fail independently
 - C. They always fail together, which is why a single check covers both
 - D. Accuracy applies to numbers and completeness applies to prose

22. Why are missing elements harder to spot than wrong ones?
- A. Because completeness can only be assessed by a second reviewer
 - B. Because Claude deliberately omits material it has low confidence in
 - C. Because omissions only occur in outputs longer than one page
 - D. Nothing on the screen draws your eye to something that is absent
23. In the worked protocol, Claude summarized three competitors' published pricing from uploaded PDFs. Requirements were met, but the source review caught a problem. What was the verdict?
- A. Needs human override — pricing comparisons always require escalation
 - B. Ready to use — all three competitors were covered as requested
 - C. Needs revision — and the fix is a source-matched re-prompt, not a rewrite
 - D. Ready to use, provided the PDFs are re-uploaded first
24. What did the source review catch in that competitor-pricing summary?
- A. Each tier was listed as '\$40/user' when the uploaded PDF said '\$40/user, minimum 10 seats,' which changes the comparison
 - B. One competitor had been omitted from the summary entirely
 - C. The summary invented a fourth competitor that was not in the PDFs
 - D. The prices were correct but presented in the wrong currency
25. You asked for three options to reduce invoice-processing time and received three sensible options with trade-offs — an internal discussion starter, with no sources to check against. What is the verdict and why?
- A. Needs human override — process recommendations affect how the team operates
 - B. Needs revision — any output without citations requires a grounded re-prompt
 - C. Ready to use — over-verifying this one wastes the time the tool saved
 - D. Ready to use, but only after the figures are recomputed with code execution
26. Claude compared your data-handling policy against a regulation and confidently flagged four gaps — but the regulation itself was never uploaded. What is the verdict?
- A. Ready to use — four flagged gaps is a substantive, actionable result
 - B. Needs human override — Claude worked from training-data recall of a rule that may have changed, and the stakes are regulatory
 - C. Needs revision — asking Claude to double-check its own four gaps resolves the issue
 - D. Ready to use, provided the output is relabeled as a draft

C. Hallucinations, Inconsistencies & Bias

27. Why can tone or confidence not be used to flag an error in Claude's output?
- A. Claude writes fluently whether it is right or wrong, so plausible is not the same as verified
 - B. Because Claude hedges more heavily when it is uncertain, which masks the real errors
 - C. Because confidence is only an unreliable signal on lower model tiers
 - D. Because tone varies at random between runs and therefore carries no information of any kind
28. What is the practical benefit of learning the specific signatures of failure?
- A. You can identify which model tier produced a given error
 - B. You can stop reviewing outputs once you recognize the patterns reliably
 - C. You can delegate the review to Claude by naming the patterns in the prompt
 - D. You can spot failures quickly instead of reading every line with equal suspicion

29. Why are plausible-but-unsupported claims described as the most dangerous hallucination pattern?
- A. They appear only in long outputs, where reviewers stop reading carefully
 - B. They are always attached to fabricated citations, which take time to check
 - C. They sound reasonable and fit the topic, and nothing about them looks wrong
 - D. They cause Claude to contradict itself later in the same output
30. Why are fabricated specifics — invented statistics, dates, names, quotations, or citations — so persuasive?
- A. They are usually correct, which is what makes the rare exceptions dangerous
 - B. Specificity reads as authority, which is exactly why an invented precise detail convinces
 - C. They always arrive with a hedge that reviewers tend to overlook
 - D. They only appear when the prompt explicitly asks for numbers
31. What does 'confident tone masking uncertainty' mean?
- A. Claude rarely hedges in proportion to its actual certainty; a guess and a well-grounded fact arrive in the same assured voice
 - B. Claude signals its uncertainty clearly, but reviewers routinely ignore the signal
 - C. Claude hedges on everything, so hedging carries no information
 - D. Claude's tone becomes noticeably less confident as its accuracy drops
32. Where do internal contradictions most often hide, and why?
- A. In the opening paragraph, which reviewers tend to skim
 - B. In short outputs, where each individual claim carries more weight
 - C. In tables, because numbers are harder to read than prose
 - D. In long documents, because you rarely hold the whole thing in view at once
33. What is confirmation bias in framing?
- A. Claude reuses the same framing across separate conversations
 - B. Claude systematically disagrees with the premise of any leading question
 - C. If your prompt implies a preferred answer, Claude may lean toward it — so watch for output that agrees with you too readily on a genuinely open question
 - D. Claude favors whichever source appears earliest in an uploaded document
34. A professional asked Claude to compare a batch of documents and identify the differences. The output read as thorough but missed the differences in the single most important file. What is the lesson?
- A. Completeness failures concentrate where attention is lowest, and a confident summary of the easy files can mask silence on the hard one
 - B. Claude cannot compare more than two documents in a single request
 - C. Batch comparisons should always be run on a higher model tier
 - D. The output was inaccurate rather than incomplete, since the listed differences were wrong
35. Claude states 'I've emailed this to your team' or 'I've saved the file.' How should such a claim be treated?
- A. As reliable, since Claude would not report an action it had not taken
 - B. As unverified — Claude does not perform external actions it was not given a tool for, so confirm the action actually happened
 - C. As a formatting artifact with no bearing on the content of the output
 - D. As evidence that a connector was enabled earlier in the session

36. An output reads: 'Approximately 63 percent of mid-market SaaS firms adopted at least one AI tool in 2025, up from 41 percent in 2024,' with no source given. What is the tell?
- A. The subject matter, since Claude has no knowledge of SaaS markets at all
 - B. The hedge 'approximately,' which signals Claude knew it was uncertain
 - C. The two-year comparison, which Claude is unable to compute
 - D. The precision — real figures this specific arrive with a citation, so an unsourced 63 percent is a number-shaped guess
37. Asked whether a clause is enforceable in your state, Claude gives a four-sentence answer stating that it is, in the same assured voice it uses for arithmetic. What is the problem?
- A. Claude should never be asked about contract clauses in any form
 - B. The answer is too short to be trustworthy, since length signals depth of analysis
 - C. Enforceability is jurisdiction-specific and date-sensitive — exactly the kind of claim that should hedge and does not; assurance is not evidence
 - D. The answer used prose where a structured format was required
38. In a ten-page market analysis, page two states the addressable market is 'roughly \$2 billion' while page eight builds a projection on 'the \$2.6 billion market.' What does this illustrate?
- A. A paragraph-by-paragraph read is sufficient provided it is done slowly enough
 - B. Claude cannot handle figures above one billion reliably
 - C. The projection is wrong and the earlier figure is right, so page eight should simply be deleted
 - D. Long outputs need a consistency pass across the whole document, because both statements read fine in isolation
39. Why is a careful paragraph-by-paragraph read insufficient for catching internal contradictions?
- A. Because contradictions are typically placed in footnotes
 - B. Each claim is individually plausible; the conflict is only visible when the whole document is held in view
 - C. Because reading order changes the meaning of each paragraph
 - D. Because Claude marks contradictions only in the closing summary
40. How does the lesson say the failure-pattern gallery should be read?
- A. None of the outputs looks broken — you catch these by knowing the signatures and checking against sources, not by waiting for something to look wrong
 - B. Each example contains a visible error that a careful reader will notice
 - C. The gallery shows failure modes unique to older model versions
 - D. The gallery is illustrative only; these patterns are rare in real work

D. Fact-Checking and Grounding Techniques

41. Where does the lesson say the strongest verification lives?
- A. In a second review pass by a colleague after delivery
 - B. Built into the prompt, before the output exists
 - C. In the model tier selected for the task
 - D. In a line-by-line fact-check of everything produced

42. Why tell Claude explicitly that admitting uncertainty is acceptable?
- A. Without that permission, a model under pressure to answer is more likely to fill the gap with something invented
 - B. Because Claude cannot produce the phrase 'I don't know' unless the prompt contains it
 - C. Because it shortens the output, which reduces the review burden
 - D. Because it switches Claude into a grounded retrieval mode automatically
43. What does restricting Claude to provided sources accomplish?
- A. It prevents Claude from reformatting the source documents
 - B. It guarantees the output will contain no errors of any kind
 - C. It converts open-ended generation into bounded retrieval, with anything the materials do not cover flagged rather than filled in
 - D. It removes the need to review the output before use
44. What makes a citation auditable?
- A. It was produced in a session where web search had been enabled at some point
 - B. It names a well-known publication, which establishes credibility
 - C. It appears in a footnote rather than inline
 - D. It gives the specific source and location behind the claim, in a form you can check
45. The verification checklist asks four questions before you rely on an output. Which set?
- A. Did I pick the right entry point, the right model, the right Skill, and the right format?
 - B. Did I allow uncertainty, restrict to sources where appropriate, require citations I can audit, and check the high-stakes claims against something authoritative?
 - C. Is it accurate, is it complete, is it well formatted, and is it on time?
 - D. Did I define the role, the context, the task, and the constraints?
46. Why does the lesson argue for prevention over after-the-fact correction?
- A. Building these habits into the prompt is cheaper than rebuilding trust in the output afterward
 - B. Because errors caught after the fact cannot be corrected
 - C. Because post-hoc checks are only available on paid plans
 - D. Because Claude cannot revise an output once it has been produced
47. You need Claude to answer questions strictly from a contract you uploaded, without inventing anything. Which prompt technique most reduces hallucination?
- A. Use the most capable model and ask in a single sentence
 - B. Ask the question and trust Claude to stay within the document
 - C. Ask the same question five times and average the answers
 - D. Restrict the answer to the provided document, permit 'I don't know,' and require citations to specific clauses
48. A researcher asks for market statistics and receives five precise figures with no citations. Which prompt habit would most likely have prevented this?
- A. Asking for ten figures so the weakest ones stand out by comparison
 - B. Requesting a table instead of prose so the figures are easier to read
 - C. Requiring auditable citations — the specific source and location for each figure — and permitting 'I don't know'
 - D. Raising the model tier and rerunning the identical prompt

49. What is the purpose of instructing Claude to flag anything the supplied materials do not cover?

- A. It causes Claude to search the web for the missing material automatically
- B. It shortens the output by omitting any uncovered topic
- C. It surfaces the gaps explicitly instead of letting generation quietly fill them
- D. It shifts responsibility for the gaps onto the source documents

50. Which of the following is NOT one of the lesson's prompt habits for reliability?

- A. Run the prompt several times and average the numeric answers so errors cancel out
- B. Permit 'I don't know' explicitly
- C. Restrict the answer to the sources you provided
- D. Require citations you can audit

E. Diligence: When Human Review Is Non-Negotiable

51. What does Diligence mean in this module?

- A. Reviewing every output to the same depth regardless of what is at stake
- B. Knowing the escalation thresholds in advance, so the decision to escalate is made by policy rather than in the moment after something has gone wrong
- C. Iterating on a prompt until the output stops changing between rounds
- D. Recording which model tier produced each deliverable

52. What are the four risk thresholds?

- A. Cost, speed, volume, and model tier
- B. Accuracy, completeness, clarity, and tone
- C. Requirements, sources, standards, and format
- D. Stakes, reversibility, audience, and regulatory exposure

53. Which question belongs to the Stakes threshold?

- A. Does a rule, contract, or law govern this content?
- B. Can the action be undone once it has been taken?
- C. Who sees it — internal, external, executive, or regulatory?
- D. What is the cost if this is wrong? High-cost errors demand human review regardless of how confident the output appears

54. Which question belongs to the Reversibility threshold?

- A. Which audience will ultimately read the output?
- B. What is the cost if the output turns out to be wrong?
- C. Can the action be undone? An irreversible step, such as a sent client deliverable or a filed report, clears a higher bar than a draft you can revise
- D. Is the content governed by a regulation or a contract?

55. Which question belongs to the Audience threshold?

- A. How long will the output remain in circulation?
- B. Who sees it? External, executive, and regulatory audiences raise the review requirement above internal working drafts
- C. Can the deliverable be recalled after it has been sent?
- D. Does the output contain figures that need recomputing?

56. Which question belongs to the Regulatory exposure threshold?
- A. Does a rule, contract, or law govern this? Regulated content carries obligations that AI assistance does not remove
 - B. Is the reviewer qualified to interpret the regulation in question?
 - C. Has the regulation been uploaded to the Project knowledge base?
 - D. Will the output be read by more than one person?
57. Which of the following belongs on the do-not-ship-without-review list?
- A. Final client deliverables, audit-critical or financially material calculations, regulated or highly sensitive data, and public or legal communications
 - B. Any output longer than two pages, regardless of its content
 - C. Any output produced in a session where code execution was disabled
 - D. Internal meeting agendas and early-stage brainstorm
58. How should the do-not-ship-without-review list be treated?
- A. Applied only when the output contains numbers
 - B. Revisited for each deliverable, since the categories shift with workload
 - C. Decided in advance and treated as fixed, so the escalation call is never made under time pressure
 - D. Delegated to Claude, which can classify outputs against the list
59. What distinguishes productive iteration from the point where escalation is required?
- A. Escalation is required whenever the model tier changes mid-task
 - B. Iteration is productive through the fifth round, after which escalation is automatic
 - C. Escalation is required only once a visible error appears in the output
 - D. Productive iteration improves the output each round; when rounds stop improving it, the right move is human review rather than another prompt
60. What does the lesson mean by saying more prompting cannot manufacture the judgment a situation requires?
- A. Prompt quality has no effect on output quality in high-stakes work
 - B. Some situations need a person's judgment, and no further round of prompting substitutes for it
 - C. Judgment can be encoded in a prompt if the constraints are precise enough
 - D. Iteration should be avoided entirely on any task with external stakes
61. What does the module say about owning the output?
- A. Accountability passes to the reviewer once the output has been escalated
 - B. Accountability is shared between the user and the model provider
 - C. Accountability does not transfer to the tool — work produced with Claude is your work when you ship it, held to the same standard as unaided work
 - D. Accountability applies only to outputs that leave the organization
62. Claude drafts an internal meeting agenda. How do the four thresholds resolve?
- A. Low stakes, reversible, internal audience, no regulatory exposure — the routine case you can move quickly on
 - B. Low stakes but irreversible, so a second reader is still required
 - C. Internal audience but high regulatory exposure, so escalation applies
 - D. All four thresholds are ambiguous, so the default is human override

63. Claude produces a board-deck financial summary that reads cleanly. What does the module prescribe?
- A. The clean appearance is evidence the figures are sound, so it can ship as drafted
 - B. It trips three thresholds — high stakes, executive audience, partly irreversible once presented — so it goes to a human reviewer and the figures get recomputed with code execution
 - C. It should be rerun on a higher model tier and then shipped
 - D. Only the formatting needs review, since the figures came from Claude's own calculation
64. You have iterated a client proposal five times, and rounds three through five changed almost nothing. What is the signal, and what should you do?
- A. Restart the task in a fresh conversation and run five more rounds
 - B. The absence of any visible error is the signal that the proposal is ready to send
 - C. Five rounds is the maximum, so the fifth version should go out as it stands
 - D. The flat improvement curve is the signal — stop prompting and escalate to a colleague for a fresh read

F. Editing and Adapting Output for Your Audience

65. 'Claude drafts; you deliver.' What does the editing pass apply?
- A. The formatting rules attached to the output modality you selected
 - B. A second model's judgment about which draft reads better
 - C. Your professional standards and your knowledge of the audience
 - D. The verification checklist, which replaces the need to edit
66. What does 'a draft that is accurate is still not finished' mean?
- A. Editing becomes optional once the facts have been verified
 - B. Accuracy cannot be established until after editing is complete
 - C. An accurate draft only needs a formatting change to be finished
 - D. Accuracy is necessary but not sufficient; the gap between a raw draft and a finished deliverable is the editing pass
67. What are the three passes that turn a raw draft into a deliverable?
- A. Clarity, tone, and formatting
 - B. Accuracy, completeness, and fitness
 - C. Requirements, sources, and standards
 - D. Inline, artifact, and structured
68. What does the clarity pass involve?
- A. Adding detail wherever the draft is thin, since thoroughness signals rigor
 - B. Cutting hedging and tightening loose sentences — Claude tends toward thoroughness, while editing tends toward precision
 - C. Matching the register to the relationship and the occasion
 - D. Shaping the output for the way it will be read
69. What does the tone pass involve?
- A. Matching the register to the relationship and the occasion — the same content reads differently to a peer, a client, and a regulator
 - B. Adopting a uniformly formal register for every audience to stay safe
 - C. Preserving Claude's original phrasing so the draft stays internally consistent
 - D. Removing every sentence that does not earn its place

70. What does the formatting pass involve?
- A. Matching the format to the model tier that produced the draft
 - B. Applying the same template to every deliverable for consistency
 - C. Converting every deliverable into a structured table
 - D. Shaping the output for how it will be read: scannable for an executive, detailed for a working team, clean for an external recipient
71. One analysis often has to become several deliverables. How do the versions differ?
- A. Only the length changes; selection and tone stay fixed
 - B. Each version restates a different set of facts drawn from the same analysis
 - C. The underlying facts hold steady while selection, depth, and tone change — executives get decision and impact, working teams get detail and method, external communications control disclosure and framing
 - D. The executive version carries the most detail, since executives make the decisions
72. What does the lesson recommend before you start editing, when quality matters?
- A. Commit to the first response, since comparison introduces inconsistency
 - B. Generate more than once, or across models, and choose the strongest base to edit from rather than committing to the first response
 - C. Ask Claude which of its own drafts is strongest and edit that one
 - D. Compare drafts only when the first one contains a factual error
73. In the worked transformation, a raw draft on Q3 processing time became an executive version and a working-team version. What stayed the same?
- A. The facts, including the same 18 percent — only selection, depth, and tone changed
 - B. The sentence structure, which was preserved in both versions
 - C. The length, since both versions had to fit the same space
 - D. Nothing; each version reported a different figure suited to its audience
74. In that transformation, what does each cut do?
- A. The executive cut adds method, while the working cut removes it
 - B. The executive cut strips method and leads with impact; the working cut keeps the method and assigns an action with an owner and timeline
 - C. Both cuts drop the recommendation and report only the figure
 - D. The working cut is simply a shortened version of the executive cut

G. Choosing Output Formats: Inline, Artifacts, Structured, Code-Executed

75. How does the lesson frame the choice of output format?
- A. As a consequence of the model tier selected for the task
 - B. As a presentation choice that follows the reader's stated preference
 - C. As a reliability decision, not just a presentation choice — driven by what the result is for and how much the numbers have to be trusted
 - D. As a setting that should stay constant across a workstream
76. What is the Inline format for?
- A. Any output containing numbers that have to be trusted
 - B. Documents and code you intend to refine and reuse
 - C. Data that downstream tools or readers will consume directly
 - D. Conversational responses you will act on within the chat — quick and contextual, not a standalone artifact

77. What are Artifacts for?
- A. Documents and code — a separate, editable block you will refine and reuse; the right home for a deliverable rather than a reply
 - B. Quick conversational answers you act on immediately
 - C. Tables and defined schemas consumed by downstream tools
 - D. Calculations that must be computed rather than written
78. What are structured formats for?
- A. Any output that has been escalated for human override
 - B. Conversational replies that stay inside the chat
 - C. Long-form documents a recipient will open and read
 - D. Data — tables and defined schemas that downstream tools or readers can consume directly
79. What is the rule for code execution as the verified-output path?
- A. Prose generation is more reliable for arithmetic because it shows the reasoning in words
 - B. Use code execution for every output regardless of its content
 - C. When numbers must be right, have Claude compute them rather than write them — code execution returns a computed, checkable result along with the steps
 - D. Code execution removes the need for human review of numeric work
80. What limitation does the lesson attach to code execution?
- A. Code execution cannot handle files as large as a spreadsheet
 - B. Claude writes the code, so the logic can still contain a bug — the guarantee is that you can read, verify, and re-run the calculation, not that it is automatically correct
 - C. Code execution hides its steps, so the result cannot be audited
 - D. There is no limitation; a computed result is correct by definition
81. Asked for total Q3 revenue and the top three accounts from an uploaded spreadsheet, the prose path returned 'about \$4.7 million' plus three account names. What is wrong with it?
- A. It listed the accounts in the wrong order, which is the only real error
 - B. It is too short; the same figures in a longer answer would be reliable
 - C. It is fluent, fast, and unverifiable — the total is the model's best guess at summing a column it cannot reliably add, and a wrong total propagates into every downstream decision
 - D. Nothing is wrong; prose figures are acceptable for any reporting purpose
82. What does the code-executed path return for that same task?
- A. Claude writes and runs code over the actual file, returning a computed sum, the top accounts ranked by their real totals, and a chart — traceable to the rows that produced it
 - B. Claude estimates the sum and then formats it as a table for clarity
 - C. Claude retrieves a comparable figure from its training data and cites the source
 - D. Claude asks the user to confirm the total before reporting it
83. Is the prose answer simply the lazy option?
- A. No — prose and code-executed figures are equally reliable in every case
 - B. Yes — prose figures should never be produced under any circumstances
 - C. Yes — prose output indicates the wrong model tier was selected
 - D. No — it is a different kind of output, and for a low-stakes gut-check it may be fine; the reliability requirement is what points to code execution

84. What is the selection rule for output modality?
- A. Pick the output modality by how the recipient prefers to read
 - B. Pick the output modality by the reliability the task requires
 - C. Pick the output modality by the expected length of the answer
 - D. Default to structured formats, which are reliable for every purpose
85. Which three techniques does the lesson name for curating inputs?
- A. De-duplicate your sources, label and structure what you supply, and prune material irrelevant to the question
 - B. Upload as much material as possible, compress it, and let Claude select what matters
 - C. Convert every input to a table, sort it, and cite it inline
 - D. Split inputs across several conversations, one document per session
86. Why does curating inputs mean removing material as well as adding it?
- A. Because uploading more than three documents is not supported
 - B. Because noise in the input becomes noise in the output — clean, well-labeled, minimal inputs produce a cleaner result than a large undifferentiated pile
 - C. Because curation affects processing speed but not output quality
 - D. Because curation is only necessary when code execution is unavailable

H. Exercise, Quiz & Mixed Review

87. Triage exercise, Output A: an internal brainstorm of campaign ideas, clearly labeled as a starting point, with a few generic entries among strong ones. What is the verdict?
- A. Needs human override — campaign ideas carry reputational risk
 - B. Needs revision — every generic entry must be replaced before the list can be used
 - C. Ready to use — low stakes, internal, and explicitly a starting point; the generic entries do not make it unfit for its purpose
 - D. Ready to use, but only once the ideas have been checked with code execution
88. Triage exercise, Output B: a financial summary stating a total that does not match the sum of the line items above it. What is the verdict?
- A. Ready to use — the mismatch is a formatting artifact rather than a calculation error
 - B. Ready to use — the line items are present, so a reader can correct the total
 - C. Needs human override — any arithmetic error requires escalation to a person
 - D. Needs revision — the gap is specific and identified, so recompute with code execution and reconcile the total before use
89. Triage exercise, Output C: a market overview citing three statistics, none of which names a source you can check. What is the verdict?
- A. Ready to use — the statistics are consistent with the rest of the overview
 - B. Needs revision — unsourced specifics are the fabricated-specific pattern, so re-prompt with grounding and auditable citations or remove the figures
 - C. Needs human override — an output without citations can never be repaired
 - D. Ready to use, provided the figures are labeled as approximate

90. Triage exercise, Output D: a draft regulatory filing summary, accurate on the points you spot-checked, intended for submission to a regulator. What is the verdict?
- A. Needs human override — regulatory exposure, a regulator as audience, and an irreversible submission outweigh the fact that spot checks passed
 - B. Ready to use — the points you checked were accurate
 - C. Needs revision — a tone pass for the regulator is the only remaining gap
 - D. Ready to use, because labeling it a draft removes the regulatory obligation
91. A report is accurate on every figure you checked, but you suspect a relevant factor was left out entirely. What does a complete Discernment review require?
- A. Re-running the report on a more capable model
 - B. Nothing further; every figure checked out, so the report is sound
 - C. A separate completeness check against the requirements, looking for missing elements rather than wrong ones
 - D. Asking Claude whether it left anything out and trusting the answer
92. Claude extracted figures from your source files into a clean, well-formatted reconciliation table that will feed a regulatory filing, but its subtotal does not match the line items above it. What is the correct action before the table is used?
- A. Ask Claude to reformat the table so it looks cleaner, then use it
 - B. Accept the table; the formatting is clean and the other figures look right
 - C. Submit the filing now and flag the subtotal discrepancy to the regulator afterward
 - D. Recompute the figures with code execution, confirm the subtotal reconciles, and resolve the discrepancy first
93. One analysis must go to both the executive team and the working group. What is the best approach?
- A. Send the same full draft to both audiences
 - B. Produce an executive version leading with decision and impact, and a working-team version keeping detail and method
 - C. Send the executives the raw Claude output to save time
 - D. Ask Claude to decide which audience matters more
94. A deliverable depends on a multi-variable financial calculation that has to be exactly right. Which output path is most trustworthy?
- A. Use code execution so the calculation is actually run and the result is verifiable
 - B. Ask Claude to calculate it inline and review the number carefully
 - C. Request the figure formatted as a table
 - D. Use Research to find a comparable published figure
95. A learner pastes six overlapping source files — several near-duplicate drafts of the same memo — and asks for a policy summary. It comes back muddled, repeating contradictory points. What most improves output quality?
- A. Paste all six files again in a fresh conversation
 - B. Switch to a higher model tier and re-run the same six files
 - C. Ask Claude to make the summary longer and more detailed
 - D. De-duplicate and prune the sources to the single approved version, label what each remaining input is, then re-run

96. Which two behavioral indicators does the optional self-assessment variant ask you to score your own conversation against?
- A. Model tier selection and entry-point selection
 - B. Accuracy and completeness of the final deliverable
 - C. Diligence — whether you checked the output before using it — and Discernment — whether you judged that the response actually fit the question you asked
 - D. Prompt length and the number of iteration rounds
97. According to Key Takeaway 1, why is this the exam's largest section?
- A. Accountability stays with you — you own every claim in what you ship, whether you wrote it or Claude did
 - B. Because output formats change more often than any other product area
 - C. Because hallucination is the only failure mode that matters in practice
 - D. Because verification is the slowest step in any workflow
98. According to Key Takeaway 3, what do fabricated specifics, confident uncertainty, and completeness gaps have in common?
- A. They all occur only in outputs longer than one page
 - B. They all read as competent, which is why you learn the signs instead of waiting for something to look wrong
 - C. They are all eliminated by moving to a higher model tier
 - D. They are all flagged automatically before the output is displayed
99. According to Key Takeaway 5, when should the human-review line be set?
- A. After the first error reaches an external audience, so the line is evidence-based
 - B. At the point of delivery, once the audience is finally known
 - C. In advance — stakes, reversibility, audience, and regulatory exposure decide when review is mandatory, and the line is set before the moment, not after
 - D. Once a quarter, as part of a Memory curation review
100. Which best summarizes the throughline connecting the six lessons of Module 3?
- A. A case for reviewing every output to the same maximum depth regardless of context
 - B. Six unrelated topics grouped together for exam convenience
 - C. A sequence of model-selection decisions, since capability determines reliability
 - D. A single accountability loop: evaluate against three references, learn the failure signatures, build verification into the prompt, escalate at fixed thresholds, edit for the audience, and pick the format by the reliability required

Answer Key & Explanations

1. **D** — The framing is a small, visible time saving set against a large, delayed cost when an unverified error goes out.
2. **A** — Four figures were sound; one plausible fabricated growth rate survived to the client meeting and was caught by the client's analyst.
3. **C** — The ten minutes saved cost a credibility hit and a week of rebuilding trust; the bill arrives later and is paid by the relationship.
4. **B** — The stated reason is accountability: ownership of every claim in anything you ship.
5. **C** — Discernment is critical evaluation against requirements, sources, and standards — it is how you review.
6. **B** — Diligence is deciding when verification is required and owning the result — it is why you must review.
7. **A** — Memory configuration belongs to Module 1's capability layer; it is not among this module's six objectives.
8. **D** — Evaluation is a repeatable check against three fixed references, run the same way every time.
9. **D** — The three references are the requirements you set, the source material you supplied, and the professional standards of your field.
10. **C** — The requirements check is a re-read of your own request, part by part, including the parts that were harder to satisfy.
11. **B** — Specific claims get traced back to the document; you do not assume careful reading on Claude's part.
12. **A** — These three all read fluently and all fail the standards of a professional field.
13. **D** — A fixed protocol keeps review quality from tracking your available time and attention.
14. **C** — Stakes are domain-dependent, so the depth of review is calibrated to the domain rather than applied uniformly.
15. **A** — In zero-tolerance domains accuracy outranks speed and every claim is verified.
16. **B** — A lighter review suits low-stakes internal work; the risk is carrying that same casual review into work where it does not belong.
17. **B** — Each output is sorted into one of these three states, with the reasoning documented.
18. **D** — 'Needs revision' is the verdict when the output is close and the remaining gap is specific and nameable.
19. **C** — Human override is escalation to a person, triggered by stakes, error load, or uncertainty.
20. **A** — All three references have to clear before an output ships without further work.
21. **B** — An output can be entirely accurate and still omit the one factor that changes a decision.
22. **D** — Wrong content is visible; absent content has nothing on the page to catch your attention.
23. **C** — A specific, traceable source mismatch is exactly the 'needs revision' case; the repair is targeted, not a fresh start.
24. **A** — The dropped seat minimum is a source-fidelity failure that materially changes the comparison.
25. **C** — Requirements met, no sources to check, standards cleared, low stakes: verifying further would spend the saving for nothing.
26. **B** — The output is a useful prompt for a compliance expert, not a substitute for one.
27. **A** — Fluency is constant regardless of correctness, so it cannot serve as an error signal.
28. **D** — Pattern recognition makes review fast and targeted rather than uniformly slow.
29. **C** — The danger is the absence of any surface signal: the claim has no basis in the source or in fact, and still reads normally.
30. **B** — Precision looks like evidence, so a fabricated specific carries more weight than a vague claim would.
31. **A** — Hedging does not track certainty, so the voice is the same whether the claim is grounded or guessed.
32. **D** — Length is what conceals a contradiction: claims made early can conflict with claims made later.
33. **C** — A leading prompt can pull the answer with it, which is most dangerous when the question was supposed to be open.
34. **A** — The listed differences were fine; the failure was the silence on the file that mattered most.

35. **B** — This is a capability hallucination: within claude.ai Claude works with the conversation, connected tools, and uploaded files, and any other claimed action needs confirming.
36. **D** — Specificity without a citation is the signature of a fabricated specific.
37. **C** — The claim type demands hedging and citation; the confident tone is the failure signature, not a reassurance.
38. **D** — Neither page looks wrong on its own; only a whole-document consistency pass surfaces the conflict.
39. **B** — Contradiction is a relationship between distant claims, so local reading cannot detect it.
40. **A** — Fluent generation is designed to pass a casual read, so detection comes from signatures and sources rather than from an output looking off.
41. **B** — Prompt-level habits cut hallucinations at the source rather than catching them afterward.
42. **A** — Permitting 'I don't know' removes the pressure that produces invented filler.
43. **C** — The instruction bounds the answer to the supplied material and turns gaps into flags.
44. **D** — A citation you cannot trace is not a citation; the location has to be checkable.
45. **B** — These four questions are the module's pre-reliance checklist.
46. **A** — The cost compared is prompt-time effort versus the far larger cost of repairing trust later.
47. **D** — All three grounding habits apply together: bounded sources, permitted uncertainty, and clause-level citations.
48. **C** — Unsourced precise figures are the fabricated-specific pattern; the citation requirement is the direct preventive habit.
49. **C** — Flagged gaps are visible and reviewable; unflagged gaps get filled with plausible invention.
50. **A** — Averaging repeated generations is not a grounding technique; the three habits are permitted uncertainty, restricted sources, and auditable citations.
51. **B** — Diligence sets the thresholds in advance so escalation is a policy decision, not an improvisation under pressure.
52. **D** — These four questions decide whether human review is mandatory.
53. **D** — Stakes asks about the cost of being wrong, independent of how the output reads.
54. **C** — Reversibility asks whether the step can be walked back; irreversible steps carry a higher review bar.
55. **B** — Audience raises or lowers the bar: internal drafts sit below external, executive, and regulatory readers.
56. **A** — Regulated content carries obligations that using an AI tool does not remove or transfer.
57. **A** — These four categories are the fixed list, decided in advance.
58. **C** — The point of a fixed list is that the judgment is already made before the deadline arrives.
59. **D** — Diminishing returns is the signal: the next move is a person, not another prompt.
60. **B** — Recognizing that line — where prompting stops and judgment starts — is part of Diligence.
61. **C** — The professional standard is identical to work produced without assistance; Diligence is the habit of acting as though that is true.
62. **A** — This is the fast 'yes': all four thresholds sit low, so the routine path is appropriate.
63. **B** — This is the deceptive 'looks fine': the clean surface is irrelevant to the three thresholds it trips.
64. **D** — This is the slow creep: diminishing returns on a high-stakes external deliverable calls for a human, not another round.
65. **C** — The editing pass is where your standards and your audience knowledge enter the deliverable.
66. **D** — Verification and editing are separate passes; clearing one does not clear the other.
67. **A** — Three passes: clarity, tone, and formatting.
68. **B** — Clarity removes anything that does not earn its place, pulling a thorough draft toward precision.
69. **A** — Tone is register calibration, not a fixed level of formality.
70. **D** — Formatting follows the reading situation, not a house template.
71. **C** — Facts are constant; what varies is what is selected, how deep it goes, and how it is framed.

72. **B** — Comparing candidates is cheaper than reworking a weak draft, and it exposes framings you might not have prompted for.
73. **A** — Same facts, same 18 percent; the two cuts differ in what they keep and how they frame it.
74. **B** — Neither is a raw draft, and neither version would serve the other audience.
75. **C** — Format is chosen by the reliability the task demands, above all for numeric work.
76. **D** — Inline is the quick contextual reply, not a deliverable.
77. **A** — Artifacts hold deliverables you will refine and reuse.
78. **D** — Structured output exists so downstream consumers can use it without restructuring.
79. **C** — Computed numbers are traceable; generated numbers are plausible-looking guesses.
80. **B** — The guarantee is auditability and reproducibility, not automatic correctness.
81. **C** — A prose figure is a number-shaped guess, and the error travels with every decision built on it.
82. **A** — The computed figure (\$4,712,380 in the example) is traceable back to the rows that produced it.
83. **D** — Prose is fine for a low-stakes gut-check; anything that gets reported points to code execution.
84. **B** — Reliability requirement drives the modality choice.
85. **A** — De-duplicate, label, prune — these three do most of the work.
86. **B** — Organized inputs produce organized outputs; the wrong material has to come out, not just the right material go in.
87. **C** — Depth of review follows the stakes; a labeled internal starting point clears the bar as it stands.
88. **D** — A named, fixable gap is the 'needs revision' case; note that if the same summary fed a regulated filing, the Diligence thresholds would push it to override.
89. **B** — Three uncheckable statistics is a specific, repairable gap: ground them or drop them.
90. **A** — Three of the four thresholds fire at once, and passing spot checks does not substitute for review of a regulated submission.
91. **C** — Accuracy and completeness fail independently, so a clean accuracy pass does not close the completeness question.
92. **D** — An internal inconsistency in a regulated numeric deliverable is exactly the case for computed, checkable figures.
93. **B** — One analysis, two audiences: the facts hold while selection, depth, and tone are cut for each reader.
94. **A** — When numbers must be right, they get computed rather than generated.
95. **D** — Contradictory input produces contradictory output; the fix is at the input, not the model tier.
96. **C** — The variant scores a real conversation against the module's two anchoring competencies.
97. **A** — Takeaway 1 ties the section's weight directly to ownership of shipped claims.
98. **B** — Plausible is not verified: all three failure types pass a casual read, so recognition has to be deliberate.
99. **C** — Set the line before the moment, not after: that is what makes escalation a policy rather than a judgment call under pressure.
100. **D** — Discernment supplies the review protocol, the failure patterns and grounding techniques make it concrete, Diligence sets the escalation line, and editing and format selection carry the output to a standard you can put your name on.