

Claude Certified Associate – Foundations (CCAO-F) — Module 4 Practice Exam

Workflow Integration & Solution Design

100 scenario-style and recall questions, multiple choice (A/B/C/D), one correct answer each.

Covers: Analyzing Requirements and Use Cases with Claude · Research, Planning & Process Optimization · Solution Design, Development & Iteration · Delegation Mapping: Redesigning Workflows with Claude Inside · Communicating Value and Limitations to Stakeholders.

Instructions: Select the single best answer for each question. An answer key with brief explanations is provided at the end of this document. Correct answers are distributed evenly across A, B, C and D. This practice material is based on user-provided course screenshots and is intended for self-study only — it is not an official Anthropic exam product.

A. Module Introduction

1. What is the difference between 'I use Claude' and 'our workflow uses Claude'?
 - A. The first uses Chat while the second uses Research mode exclusively
 - B. The first requires a paid plan while the second requires an enterprise agreement
 - C. The first is a personal productivity habit; the second is a repeatable process, run by a team, where Claude performs specific steps every time
 - D. They are the same thing described at two different scales
2. In the two-teams story, what did the first contract-review team do?
 - A. It let Claude approve low-risk clauses unsupervised and still halved review time
 - B. It mapped the work and let Claude draft the redline while a lawyer revised every final decision; review time dropped by half and quality held
 - C. It declined to use Claude on contracts and kept the process fully manual
 - D. It replaced its legal reviewers with a configured Skill
3. In the two-teams story, what happened to the second team?
 - A. It pointed Claude at the whole process and let it approve low-risk clauses unsupervised; within a month an approved clause created an obligation nobody caught, and the team pulled the tool entirely
 - B. It used Claude only for clause extraction, which proved too limited to be worth keeping
 - C. It ran the workflow on a lower model tier, which introduced formatting errors
 - D. It succeeded, but the process took longer than the first team's
4. What was the difference between the two teams?
 - A. One team had in-house legal expertise and the other did not
 - B. The model tier each team selected for the work
 - C. One team used a Project and the other used one-off Chats
 - D. Which steps each team chose to delegate — same product, same process
5. Which AI Fluency Framework competency anchors this module?
 - A. Description — making each component of a prompt explicit
 - B. Discernment — critically evaluating output against requirements, sources, and standards
 - C. Delegation — deciding, for each step, whether the work is AI-appropriate, human-retained, or collaborative
 - D. Diligence — deciding when verification is required and owning the result

6. What is 'the deal' in this module?
- A. Workflow value comes from deciding which steps to delegate, not from automating everything
 - B. Workflow value comes from automating as many steps as the tool can handle
 - C. Workflow value comes from using the highest available model tier at every step
 - D. Workflow value comes from removing human review gates so the process stops queueing
7. Which of the following is NOT one of the module's stated learning objectives?
- A. Apply Claude to analyze requirements and use cases
 - B. Build production-grade software systems with uptime, security, and scaling guarantees
 - C. Integrate Claude into existing workflows to augment or redesign them
 - D. Communicate Claude's value and limitations to stakeholders accurately
8. When does the value of Claude in a workflow actually compound?
- A. Once the workflow has been moved out of Chat and into Research mode
 - B. Automatically, as soon as more people on the team start using Claude
 - C. Once every step of the process has been automated
 - D. Only when the shift from personal habit to repeatable team process is made deliberately

B. Analyzing Requirements and Use Cases with Claude

9. What does the lesson say most real work starts from?
- A. A dataset that can be handled with code execution alone
 - B. A clean specification that only needs reformatting
 - C. Messy inputs — a long document, a thread of half-formed emails, a verbal ask — which have to be extracted, structured, and pressure-tested before anything is built
 - D. A stakeholder briefing that already contains testable requirements
10. When you upload raw inputs, what should you ask Claude for?
- A. A structured analysis rather than a narrative summary
 - B. A narrative summary, which stakeholders find easier to read
 - C. A single-sentence restatement of the business need
 - D. A list of the model's own uncertainties and nothing else
11. What does Claude do with unstructured source material in this lesson?
- A. Verifies each requirement against an external authority automatically
 - B. Rewrites the document in a more professional register
 - C. Compresses the document to a fixed word count
 - D. Pulls the requirements out of the document, organizes them, and flags what is ambiguous or missing
12. Why is a business need stated as 'we need better reporting' not actionable?
- A. It is too long to be processed as a single requirement
 - B. It leaves the task definition unspecified — what report, for whom, how often, drawn from what data, and in what format
 - C. It names an outcome, which requirements are never permitted to do
 - D. It cannot be converted into requirements without a technical build

13. What does converting a business need into task definitions produce?
- A. A guarantee that no requirement has been overlooked
 - B. A narrative summary the whole team can read quickly
 - C. A model-selection decision for each step of the work
 - D. Requirements you can build against and check completion on
14. In the RFP worked example, what were the inputs?
- A. A 40-page RFP document and a scattered email thread of internal answers
 - B. A one-page client brief and a verbal debrief
 - C. Four quarters of ticket-volume data
 - D. A contract and the company's clause playbook
15. In the RFP worked example, what was the recurring task?
- A. Pricing the bid using code execution
 - B. Drafting the full proposal response in a single pass
 - C. Turning the raw inputs into a structured list of answerable requirements
 - D. Deciding whether the company should bid at all
16. What did the RFP prompt ask for, per requirement?
- A. A paragraph of narrative context and a recommended owner
 - B. A short label, the exact RFP section it came from, whether an answer already exists internally, and a flag on anything ambiguous that needs clarification — returned as a table
 - C. A risk score and a delegation classification
 - D. An external citation supporting each requirement
17. What did the RFP output look like?
- A. A narrative memo summarizing the RFP for executives
 - B. A requirements table the whole team can work from: each row a requirement, traced to its RFP section, marked answered or open, with ambiguities flagged for a clarifying question to the client
 - C. A finished proposal draft ready for client review
 - D. A dashboard artifact tracking bid progress
18. What entry point does the RFP workflow call for, and why?
- A. Research mode, because RFP responses require multi-source investigation
 - B. A one-off Chat, because every RFP is different
 - C. A Project rather than a one-off Chat, because the work recurs and the background context is stable
 - D. Incognito mode, because RFP content is confidential
19. In that Project setup, where do the pieces live?
- A. Standing instructions and the knowledge base live in the Project's configuration, while Skills live at the account level and apply everywhere, that Project included
 - B. All three live inside the Project and cannot be used anywhere else
 - C. All three live at the account level and apply to every conversation
 - D. Standing instructions live at the account level while Skills are Project-scoped
20. Does the RFP workflow require a technical build?
- A. No, but only on the highest available model tier
 - B. Yes — a custom integration has to be built before the workflow can run
 - C. Yes — the extraction step needs a developer to write code
 - D. No — no technical build is required, and Module 5 covers the configuration in depth

21. What is the purpose of pressure-testing the extracted requirements?
- A. Extraction is only the first pass; asking Claude to challenge its own list is what makes the output trustworthy
 - B. Pressure-testing replaces extraction with a single more reliable pass
 - C. Pressure-testing is a formatting step applied once the table exists
 - D. Pressure-testing is only needed when the requirements came from a verbal ask
22. What does the pressure-test pass surface?
- A. Formatting inconsistencies in the requirements table
 - B. Hidden requirements — the ones buried in a subordinate clause or implied by an evaluation criterion, which cost teams the bid when missed
 - C. Requirements that duplicate one another word for word
 - D. The correct delegation classification for each requirement
23. How does the requirements lesson connect to Module 3?
- A. It applies Module 1's context-management skills to long documents
 - B. It applies Module 2's Description competency to stakeholders
 - C. The structured list plus a pressure-test pass is the Discernment habit applied at the requirements stage
 - D. It removes the need for verification later in the workflow

C. Research, Planning & Process Optimization

24. Which two things does planning work usually mix, and how does Claude handle them differently?
- A. Judgment, where Claude excels, and synthesis, which stays human
 - B. Drafting, where Claude excels, and formatting, where it must be verified
 - C. Research, which is always reliable, and citation, which never is
 - D. Synthesis, where Claude excels, and calculations, where it must be verified
25. What pairing produces the strongest planning workflows?
- A. Claude's synthesis paired with code-executed analysis, so the plan rests on numbers that were computed rather than generated
 - B. Claude's synthesis paired with a higher model tier for the numeric steps
 - C. Research mode paired with Incognito mode for confidential planning
 - D. Prose calculation paired with a manual spot-check of one figure
26. What can Claude do in the research-and-synthesis part of planning?
- A. Make the choice between options once the trade-offs are clear
 - B. Guarantee that every claim in the synthesis is grounded in a checkable source
 - C. Gather the considerations, structure the options, and lay out the trade-offs across sources
 - D. Compute the underlying figures reliably in prose
27. For current information that post-dates training, how do web search and Research divide the work?
- A. Research covers quick lookups while web search supplies the deep investigation
 - B. Web search covers quick lookups while Research supplies the deeper, up-to-date inputs
 - C. Neither is needed, since planning relies only on uploaded internal documents
 - D. A higher model tier extends the training cutoff, so neither is required

28. Where can unverified claims enter a planning workflow?
- A. Nowhere, provided the sources were uploaded rather than searched
 - B. Only in code-executed steps, where a bug can change the result
 - C. Only when web search has been left disabled
 - D. In the synthesis — which is why the verification discipline from Module 3 applies throughout
29. How does the lesson illustrate the difference between a guessed and a computed input?
- A. Both approaches are equivalent, since Claude explains its reasoning either way
 - B. A plan built in prose is faster, which makes it preferable for early planning
 - C. Utilization rates should be recalled from training data for benchmarking
 - D. A staffing plan built on a guessed utilization rate is a guess; one built on a code-executed analysis of the actual timesheet data is a plan
30. What is the workflow in the capacity-plan worked example?
- A. Ask Claude to recall industry-standard throughput and build the plan around it
 - B. Upload four quarters of ticket-volume data, use code execution to compute the trend and per-analyst throughput, then have Claude synthesize a staffing recommendation from the verified figures
 - C. Have Claude draft the recommendation first and compute supporting figures afterward
 - D. Use Research to find comparable headcount plans at similar companies
31. What did the capacity-plan prompt ask code execution to calculate?
- A. Quarterly volume growth and average tickets resolved per analyst
 - B. The hiring budget and the salary bands for new analysts
 - C. The company's target resolution time
 - D. The forecast error of the previous quarter's plan
32. Why can the resulting capacity plan be defended line by line?
- A. Because the prompt asked Claude to show its assumptions, which guarantees correctness
 - B. Because the plan was produced on the highest available model tier
 - C. Because the figures came from code execution rather than prose, so the recommendation rests on computed numbers
 - D. Because the operations lead approved the plan before it was generated
33. What does the lesson say about the trustworthiness of the recommendation?
- A. The recommendation is trustworthy once it has been through three iterations
 - B. The recommendation is trustworthy provided the reasoning reads coherently
 - C. The recommendation is only as trustworthy as the figures under it
 - D. Trustworthiness depends on how many sources the synthesis cites
34. Where in a planning workflow does Claude add the most?
- A. At the judgment steps, where risk appetite and political reality are weighed
 - B. At the synthesis steps, where many considerations have to be weighed and structured
 - C. At every step equally, which is why whole planning workflows should be delegated
 - D. At the approval steps, since Claude applies criteria consistently
35. Which steps of a planning workflow stay human?
- A. The formatting steps, where the shape of the output matters
 - B. The synthesis steps, where many considerations must be structured
 - C. The data-processing steps, where files have to be transformed
 - D. The judgment steps, where a person weighs risk appetite or political reality

36. In the capacity plan, how does the leverage divide between Claude and the operations lead?
- A. Claude turns four quarters of data into a defensible recommendation; the decision to hire — against budget and hiring-freeze realities Claude cannot see — stays with the operations lead
 - B. Claude makes the hiring decision while the human formats the output
 - C. Claude owns both the analysis and the hiring approval once the figures are verified
 - D. Claude adds no leverage once the data has been uploaded
37. A planner uploads a dataset and asks Claude to estimate quarterly growth in prose, then builds a headcount plan on that figure. What is the flaw?
- A. The plan now rests on a generated figure rather than a computed one; code execution should produce the numbers first
 - B. The dataset should have been split across several conversations
 - C. The planner should have requested a narrative summary instead
 - D. The only flaw is the output format, which a table would fix

D. Solution Design, Development & Iteration

38. What does 'Claude is a design collaborator, not a vending machine' mean in practice?
- A. Each design request should be made in a fresh conversation to keep it independent
 - B. Claude should be used for design work only, never for analysis
 - C. Claude's designs always require a developer to translate them into code
 - D. The value shows up in an explicit loop — ideate, prototype, gather feedback, refine — rather than in a single request
39. What produces a solution rather than a pile of one-off drafts?
- A. Raising the model tier with each successive iteration
 - B. Generating as many independent drafts as possible and picking the strongest
 - C. Treating the work as a loop and keeping the design context stable across iterations
 - D. Keeping each draft in its own Project so the versions do not interfere
40. How does the lesson describe the iteration loop?
- A. A prototype is built first, and the requirements are derived from it afterward
 - B. Ideation produces options, a prototype makes one concrete, feedback exposes what is wrong, refinement fixes it, and the loop repeats until the solution holds
 - C. Feedback is gathered once at the end, to avoid churn during the build
 - D. Refinement comes before prototyping so the prototype is right the first time
41. Why run the design loop inside a Project?
- A. It removes the need to gather feedback between cycles
 - B. It is the only entry point where artifacts can be produced
 - C. It is the only entry point where code execution is available
 - D. It keeps the context, constraints, and prior decisions stable, so each iteration builds on the last instead of restarting

42. In the internal-process-tool example, what did the team need, and what did they do instead of commissioning a build?
- A. They needed a small internal tool to track and visualize a maintained set of metrics, so they had Claude produce it as a web artifact and iterated by asking
 - B. They needed a production reporting system, so a developer built it from Claude's specification
 - C. They needed a one-off chart, so they generated it inline in the chat
 - D. They needed a dashboard, so they used Research to find an existing product
43. What was Cycle 1 of the artifact build?
- A. Add a filter by region and a summary row with the totals
 - B. Build a simple dashboard artifact showing the five metrics from the attached data, with a chart for each
 - C. Color-code the status indicators and make the layout print cleanly
 - D. Escalate the build to Developer expertise before starting
44. What was Cycle 2?
- A. Color-code the status indicators green for improvement
 - B. Build the initial dashboard with a chart for each metric
 - C. Add a filter by region and a summary row at the top with the totals — the team refines by describing the change, not by coding it
 - D. Migrate the artifact into a production environment
45. What was Cycle 3?
- A. Hand the artifact to a developer for a security review
 - B. Add a regional filter and a totals row
 - C. Rebuild the dashboard from scratch with all feedback incorporated at once
 - D. Color-code the status indicators — green for improvement — and make the layout print cleanly
46. How much code did the team write across the three cycles?
- A. All of it; Claude supplied the design and the team implemented it
 - B. Only the chart definitions, which Claude could not produce
 - C. None — three cycles, no code written by the team
 - D. Enough to connect the artifact to a live database
47. Why did the artifact approach work in that example?
- A. Because artifacts are guaranteed to scale to any number of users
 - B. Because it served a small team's internal need
 - C. Because Developer expertise was available throughout
 - D. Because the metric definitions never changed after the first cycle
48. When has a Claude-built solution outgrown Associate scope?
- A. When it becomes a system others depend on, with uptime, security, or scaling requirements — at which point it belongs with Developer or Architect expertise
 - B. When it needs more than three iteration cycles
 - C. When it contains numeric calculations
 - D. When more than one person has seen the artifact
49. What is the escalation signal for a Claude-built solution?
- A. The moment the team stops writing code themselves
 - B. The moment the artifact grows past one screen in length
 - C. The moment a chart is added to the artifact
 - D. The moment people rely on it as infrastructure, the build is no longer a prompt-and-iterate exercise

50. A Claude-built artifact that began as one analyst's helper now runs a daily reporting step three departments depend on. What follows?
- A. Move it into Incognito mode to reduce the risk
 - B. Keep iterating by prompt, since the approach has worked so far
 - C. It has crossed into infrastructure and should be escalated toward Developer or Architect expertise rather than iterated by prompt
 - D. Switch to a higher model tier so the artifact becomes more reliable

E. Delegation Mapping: Redesigning Workflows with Claude Inside

51. What is the core skill of this module?
- A. Delegation mapping — mapping a workflow step by step and deciding, for each step, who owns it: AI, a human, or both together
 - B. Selecting the right model tier for each step of a workflow
 - C. Automating the maximum number of steps the tool can handle
 - D. Documenting a workflow so stakeholders can audit it afterward
52. What does getting the mapping right for each step determine?
- A. Whether the workflow can run without a paid plan
 - B. Whether the workflow compounds value or quietly accumulates risk
 - C. Which entry point the workflow will use
 - D. How many iteration cycles the redesign needs
53. What are the three possible classifications for a workflow step?
- A. Reversible, high-stakes, or accountable
 - B. Automated, supervised, or prohibited
 - C. Ready to use, needs revision, or needs human override
 - D. AI-appropriate, human-retained, or collaborative
54. What are the three criteria each step is judged against?
- A. Requirements, source material, and professional standards
 - B. Speed, cost, and model capability
 - C. Reversibility, cost of error, and accountability
 - D. Volume, complexity, and audience
55. What does the reversibility criterion ask?
- A. How long does the step take to complete by hand?
 - B. Can the step be undone if Claude gets it wrong? Reversible steps tolerate more delegation; irreversible ones demand human involvement
 - C. Who will be blamed if the step produces a poor result?
 - D. Does the step involve numeric computation?
56. What does the cost-of-error criterion ask?
- A. What is the cost of an error at this step? High-cost steps stay human-retained or human-reviewed
 - B. Can the step be undone after the fact?
 - C. Is the step frequent enough to justify building a Skill?
 - D. Does the step depend on current information?

57. What does the accountability criterion ask?
- A. Who is accountable for this step's outcome? Accountability does not delegate, even when the drafting does
 - B. Which team member has the highest AI literacy?
 - C. Whether the step can be reversed if it goes wrong
 - D. Whether the output will be seen outside the organization
58. Once a workflow is mapped, what gets attached at the AI-appropriate steps?
- A. A separate Project per step, so each context stays isolated
 - B. A Skill for repeatable process steps and code execution for data steps
 - C. Research mode, so every step is grounded in current sources
 - D. A higher model tier, which is what makes a step reliable
59. Why prefer a configured Skill over prompting at a repeatable step?
- A. A Skill removes the need to review the output of that step
 - B. A Skill guarantees word-for-word identical output on every run
 - C. Consistency from a configured Skill beats heroic prompting that depends on remembering the right wording every time
 - D. Prompting is not available inside a redesigned workflow
60. What happens to the human-retained steps in a redesigned workflow?
- A. They are documented but left unstaffed until an error occurs
 - B. They are removed once the AI steps have proven reliable
 - C. They are converted into collaborative steps to save time
 - D. They become explicit review gates rather than afterthoughts
61. In the contract-review map, how is 'extract clauses from the contract' classified and why?
- A. AI-appropriate — reversible, low stakes, mechanical
 - B. Human-retained — extraction determines everything downstream
 - C. Collaborative — a human must confirm each extracted clause
 - D. AI-appropriate, but only when run through code execution
62. In the contract-review map, how is 'flag departures from the company playbook' classified and why?
- A. Human-retained — playbook interpretation is a judgment call
 - B. AI-appropriate — reversible, and a Skill carries the playbook rules
 - C. Collaborative — the AI flags and the human re-reads the playbook each time
 - D. AI-appropriate only after the redline has been approved
63. In the contract-review map, how is 'draft the redline and rationale' classified?
- A. AI-appropriate with code execution, since edits have to be counted
 - B. AI-appropriate — drafting is mechanical and reversible
 - C. Human-retained — drafting carries the accountability for the contract
 - D. Collaborative — the AI drafts and the human judges each edit
64. In the contract-review map, how is 'approve or reject each change' classified and why?
- A. Collaborative — the AI approves and the human spot-checks a sample
 - B. AI-appropriate — the approval criteria are already in the playbook
 - C. Human-retained — high stakes, and accountability does not delegate
 - D. AI-appropriate, provided a Skill carries the approval rules

65. In the contract-review map, how is 'compute financial exposure of a penalty clause' classified?
- A. Human-retained — financial figures always stay with a person
 - B. AI-appropriate via code execution — the figure is numeric and must be computed, not estimated
 - C. Collaborative — the AI estimates and the human sanity-checks the number
 - D. AI-appropriate in prose, since Claude can explain its arithmetic
66. In the contract-review map, how is 'sign and send' classified and why?
- A. AI-appropriate once the redline has been approved by a human
 - B. AI-appropriate — sending is a mechanical step
 - C. Collaborative — the AI sends after a human glance at the draft
 - D. Human-retained — irreversible, external, and legally binding
67. What does the contract-review map show about the AI's role in a redesign?
- A. The AI does real work, including the redline draft, while the human owns the decisions and the irreversible steps — that split is the redesign
 - B. The AI is limited to summarizing, with all substantive work retained by the human
 - C. The AI owns every step until an error occurs, at which point steps move back
 - D. The AI and the human each perform every step so the results can be compared
68. What is over-delegation?
- A. Classifying too many steps as collaborative rather than AI-appropriate
 - B. Using a Skill where a plain prompt would have been sufficient
 - C. Giving AI more than the risk profile justifies — for instance letting Claude approve clauses or send the contract because it drafted them well
 - D. Assigning steps to a human that Claude could have handled competently
69. Why is drafting quality not a license to delegate the decision?
- A. Because drafts always contain at least one error somewhere
 - B. A good draft says nothing about who should own the approval; the classification depends on reversibility, stakes, and accountability, not on output quality
 - C. Because the approval step should be escalated to Developer expertise
 - D. Because a lower model tier should be used for decisions than for drafting
70. What is halo delegation?
- A. A collaborative step is treated as automatic because nobody reviews it
 - B. A step is handed to AI because the tool is marketed as capable of it
 - C. A step stays human because an earlier AI step failed
 - D. A step gets handed to AI because the previous step went well — each step has to be judged on its own
71. What does 'collapsing collaborative into automatic' describe?
- A. An automated step is demoted to collaborative after an error is found
 - B. A collaborative step is promoted to AI-appropriate once the AI proves reliable
 - C. 'AI drafts, human reviews' quietly becomes 'AI drafts' when the review gate is not actually staffed — a collaborative step with no real reviewer is an automated step
 - D. Two people review the same step, duplicating effort without adding safety

72. What is the 'mapping the tool, not the work' error?
- A. Teams map around the features they like — a Skill they built — rather than the actual workflow steps; map the work first, then attach features
 - B. Teams map the work in more detail than the workflow actually has
 - C. Teams map the tool's failure modes instead of its capabilities
 - D. Teams assign features before choosing a model tier

F. Communicating Value and Limitations to Stakeholders

73. What does integrating Claude into a team workflow require you to do?
- A. Certify the workflow with Anthropic before deploying it
 - B. Document the prompt library for future users of the workflow
 - C. Describe it to people who did not build it — a manager, a client, a risk function
 - D. Restrict access to the people who designed the workflow
74. Where does credibility with stakeholders come from?
- A. Accurate claims — which means communicating the limits as clearly as the value
 - B. Leading with the value and addressing limits only when asked
 - C. Demonstrating the tool live for each stakeholder group
 - D. Citing the model tier and its benchmark results
75. How do teams lose stakeholder trust, according to the lesson?
- A. By selecting a lower model tier than stakeholders expected
 - B. By describing failure modes, which invites unnecessary doubt
 - C. By naming the human review gates, which makes the tool look weak
 - D. By overstating capability, which collapses on the first visible miss
76. Which statement is accurate, and which overstates?
- A. 'Claude handles contract review' is accurate because it is concise
 - B. 'Claude drafts the first-pass redline, which a lawyer reviews' is accurate and credible; 'Claude handles contract review' overstates
 - C. Both are acceptable, depending on the audience's AI literacy
 - D. Neither should be said without legal sign-off
77. What should a high-literacy stakeholder such as a legal lead be told?
- A. That the workflow runs automatically end to end
 - B. The outcome and risk posture only, with failure modes left out
 - C. The feature detail plus the known failure mode — Claude extracts clauses, flags playbook departures, and drafts the redline; it does not approve changes, and it can miss obligations implied indirectly, so the flags prompt your read rather than replacing it
 - D. The model tier in use and the usage limits that apply
78. Across audiences, what changes and what stays the same?
- A. The human gate moves depending on who is asking
 - B. The workflow and the human gate stay the same; what changes is the detail each audience needs in order to trust it
 - C. The capability boundary is described more generously to less technical audiences
 - D. The workflow is described as automated to executives and supervised to technical stakeholders

79. What does an executive want, compared with a technical stakeholder?
- A. An executive wants the outcome, the oversight in place, and the risk posture; a technical stakeholder wants the feature detail and the failure modes
 - B. An executive wants the feature detail; a technical stakeholder wants the risk posture
 - C. Both want identical technical detail, only delivered differently
 - D. An executive wants the model tier; a technical stakeholder wants the cost
80. Why should the expectation you set match the capability boundary?
- A. So stakeholders do not ask for a different model tier
 - B. So the workflow can be approved without a review gate
 - C. So the tool can fairly be described as fully automated
 - D. So no one is surprised later
81. How does this lesson connect back to Module 2?
- A. It is the Description competency applied to stakeholders — the same precise specification of what the tool can and cannot do, now directed at people rather than at Claude
 - B. It is the Delegation competency applied to communication
 - C. It is the Discernment competency applied to stakeholder messages
 - D. It replaces the component stack when writing prompts for stakeholders
82. What does the lesson say about documenting human oversight?
- A. Replace named gates with a general assurance that outputs are reviewed
 - B. Keep the gates internal, since naming them undermines confidence in the tool
 - C. Describe the gates only to regulators and risk functions
 - D. Name the review gates that stay in place — stakeholders trust an AI workflow more, not less, when the human checkpoints are explicit
83. Which phrasing belongs in the 'accurate' column?
- A. 'Our new AI system reviews contracts automatically'
 - B. 'Claude drafts the redlines and flags playbook departures; our legal lead reviews and approves every change before anything is sent, and review time is down about half with the same approval standard'
 - C. 'Claude handles contract review end to end now'
 - D. 'The system never misses anything'
84. 'Fully automated,' 'Claude handles X,' and 'it's basically as good as a person at Y' are given as phrases that quietly overstate. What is the fix?
- A. Add a general disclaimer that AI output may contain errors
 - B. Avoid describing the workflow to non-technical audiences at all
 - C. State what the tool does, then name the human checkpoint
 - D. Treat the phrases as acceptable shorthand once the workflow is proven

G. Exercise, Quiz & Mixed Review

85. Expense-report exercise, Step 1: extract line items and amounts from submitted receipts. How should it be classified?
- A. Human-retained — receipts contain financial data
 - B. AI-appropriate — mechanical, reversible, and low stakes
 - C. Collaborative — a human must confirm every extracted line
 - D. Human-retained — extraction determines the report total

86. Expense-report exercise, Step 2: check each expense against the company travel policy. Which capability best serves this step?
- A. Memory, so the policy persists across sessions
 - B. Code execution, since policy checking is fundamentally a calculation
 - C. Research, since travel policies change over time
 - D. A Skill, which carries the policy rules consistently every run
87. Expense-report exercise, Step 3: total the report and compute the amount over or under policy limits. Which capability best serves this step?
- A. Prose calculation followed by a manual review of the result
 - B. A Skill, since the totalling procedure repeats every time
 - C. Code execution — the figures must be computed, not estimated
 - D. Research, to benchmark the total against typical expense reports
88. Expense-report exercise, Step 4: draft a summary note explaining any flagged items. How should it be classified?
- A. Collaborative — the AI drafts and a human judges the explanation before it goes anywhere
 - B. Human-retained — no draft should be generated for financial notes
 - C. AI-appropriate with code execution, since amounts appear in the note
 - D. Human-retained — written summaries carry accountability by themselves
89. Expense-report exercise, Step 5: approve or reject the report. How should it be classified?
- A. AI-appropriate once the policy Skill has proven reliable over time
 - B. AI-appropriate — the policy rules already determine the answer
 - C. Collaborative — the AI approves and a human spot-checks a sample
 - D. Human-retained — the approval decision carries the accountability and does not delegate
90. Expense-report exercise, Step 6: submit the approved report to finance for payment. How should it be classified?
- A. AI-appropriate — submission is a mechanical hand-off
 - B. Human-retained — the step is irreversible and triggers a payment
 - C. Collaborative — the AI submits and a human confirms afterward
 - D. AI-appropriate, since the approval gate has already been passed
91. Quiz: Claude extracted four requirement statements from a vendor brief. Which is most ambiguous as written — the one two engineers would build differently?
- A. 'The login page must support single sign-on through the company's existing identity provider'
 - B. 'The system must export reports in both PDF and CSV formats'
 - C. 'The system should handle a large number of concurrent users'
 - D. 'Password-reset emails must be sent within five minutes of a request'
92. Quiz: a capacity plan depends on growth and throughput figures from an uploaded dataset. How should those figures be produced?
- A. Use code execution to compute them, then have Claude synthesize the plan from the verified numbers
 - B. Ask Claude to estimate them from the data in prose
 - C. Use Research to find published industry averages instead
 - D. Have Claude recall typical figures from its training data

93. Quiz: a Claude-built dashboard artifact that started as an internal helper is now relied on daily by three departments for reporting. What does this signal?
- A. Move it into Incognito mode for safety
 - B. Nothing in particular; keep iterating by prompt as before
 - C. Switch to a more capable model and carry on
 - D. It has become a system others depend on and should be escalated toward Developer or Architect expertise
94. Quiz: in a contract-review redesign, which step should remain human-retained?
- A. Approving or rejecting each change, and signing the contract
 - B. Extracting the clauses from the document
 - C. Flagging departures from the playbook
 - D. Computing the financial exposure of a penalty clause
95. Quiz: which description of an AI-assisted contract workflow best builds stakeholder trust?
- A. 'Claude handles all the legal work now'
 - B. 'Our AI reviews contracts automatically'
 - C. 'Claude drafts the redline and flags playbook departures; our legal lead reviews and approves every change before sending'
 - D. 'The AI is fully autonomous, so we no longer need legal review'
96. According to Key Takeaway 1, where does workflow value come from?
- A. Automating every reversible step and reviewing only the irreversible ones
 - B. Delegating deliberately — choosing which steps Claude does, rather than handing Claude everything
 - C. Delegating whole workflows once the individual steps have proven reliable
 - D. Adoption — the tool supplies the value, so workflow design matters less
97. According to Key Takeaway 2, what is Claude's role in requirements work?
- A. None — requirements analysis is a judgment step that stays fully human
 - B. A requirements author that can generate a project's requirements from scratch
 - C. A reviewer that should only see requirements after a human has structured them
 - D. A requirements-analysis partner: feed it messy inputs and get structured, traceable, testable needs others can act on
98. According to Key Takeaway 3, how should plans be built?
- A. Build on Claude's prose estimates and verify only the final headline figure
 - B. Pair Claude's synthesis with code execution so the figures under a plan are computed, not generated
 - C. Use Research for numeric planning inputs, since it is grounded in sources
 - D. Prioritize the structure of the reasoning, since figures can be adjusted later
99. According to Key Takeaway 4, what decides a step's classification?
- A. Reversibility, stakes, and accountability determine whether a step is AI-appropriate, human-retained, or collaborative
 - B. Frequency, duration, and cost determine which steps to automate first
 - C. The model tier available for that step
 - D. Whether the step produces a numeric or a written output

100. Which best summarizes the throughline connecting the five lessons of Module 4?

- A. A case for automating every step a current model is capable of performing
- B. Five separate tools introduced in sequence, each replacing the one before it
- C. Requirements analysis structures the work, verified analysis grounds the plan, iteration shapes the solution, delegation mapping decides who owns each step, and accurate stakeholder communication makes the redesign defensible
- D. A model-selection framework applied to the design of workflows

Answer Key & Explanations

1. **C** — The shift from individual habit to repeatable team process is what the module is about, and it has to be made deliberately.
2. **B** — The first team delegated the drafting and retained the decisions, and got a durable win from it.
3. **A** — Over-delegating the approval step produced an uncaught obligation and cost the team the tool altogether.
4. **D** — Same product, same process; the delegation decisions were the whole difference.
5. **C** — Delegation done deliberately is what turns individual wins into workflow value.
6. **A** — The module's stated premise: choose the steps, do not hand over the workflow.
7. **B** — Production systems are exactly what the module says has outgrown Associate scope and belongs with Developer or Architect expertise.
8. **D** — The compounding requires a deliberate shift; done carelessly, teams automate the wrong steps.
9. **C** — Requirements rarely arrive structured; translating messy input into testable requirements is the first job.
10. **A** — Structure is the point: a narrative summary does not give the team something to build against.
11. **D** — Extraction, organization, and flagging of ambiguity or gaps is the described capability.
12. **B** — Converting the business need into those specifics is what turns it into something buildable.
13. **D** — Each task definition becomes a requirement with a checkable completion condition.
14. **A** — A long document plus a messy thread is the canonical messy-input case.
15. **C** — The recurring task is the translation step, not the proposal itself.
16. **B** — Label, source section, answered-or-open status, and ambiguity flags, in a table.
17. **B** — The deliverable is a traceable, actionable table rather than a summary.
18. **C** — Recurring work on stable context is the defining Project case.
19. **A** — This matches Module 1: Project configuration holds instructions and the knowledge base; Skills are account-level and invoked wherever relevant.
20. **D** — The lesson is explicit that no technical build is needed for this workflow.
21. **A** — Extraction plus a challenge pass is a far stronger foundation than either alone.
22. **B** — The expensive misses are the requirements the document states only indirectly.
23. **C** — Challenging the output against the source is Discernment, moved upstream to the requirements.
24. **D** — The lesson opens on this split and builds the workflow pattern around it.
25. **A** — Synthesis plus computed numbers is the pattern the lesson prescribes.
26. **C** — Synthesis is where Claude adds most; it does not make the call or verify itself.
27. **B** — This mirrors the Module 1 entry-point logic: quick lookup versus deep multi-source investigation.
28. **D** — Synthesis is useful and is also where unsupported claims slip in.
29. **D** — The character of the whole plan follows the character of the numbers underneath it.
30. **B** — Compute first, synthesize second: the recommendation is built on verified figures.
31. **A** — Those two computed figures are what the headcount recommendation is then built on.
32. **C** — Traceable figures are what make the plan defensible in front of a challenge.
33. **C** — Everything above the numbers inherits the reliability of the numbers.
34. **B** — Not every step benefits equally; the synthesis-heavy steps are where to spend.
35. **D** — A quick scan for synthesis-heavy steps also shows where the call must stay with a person.
36. **A** — Claude supplies the defensible analysis; the decision rests on context Claude does not have.
37. **A** — This is the guessed-utilization-rate failure: a plan built on a prose number is a guess.
38. **D** — The loop, not the single request, is where design value comes from.

- 39. C — Loop plus stable context is the stated condition for converging on a solution.
- 40. B — Four stages, run repeatedly, each feeding the next.
- 41. D — Stable context across iterations is what stops each cycle from starting over.
- 42. A — Three cycles of describing changes replaced a commissioned build.
- 43. B — Cycle 1 produced a working artifact from the attached data.
- 44. C — Refinement by description is the point of the cycle.
- 45. D — Cycle 3 handled presentation: color coding and clean printing.
- 46. C — The example's headline is three iterations with no code written by the team.
- 47. B — Scope is what made it appropriate: a small internal need, not shared infrastructure.
- 48. A — Dependency plus operational requirements marks the boundary.
- 49. D — Reliance as infrastructure is the signal, not size or complexity.
- 50. C — This is the quiz scenario: cross-department daily reliance is the escalation trigger.
- 51. A — The lesson names delegation mapping as the core skill, to be done before redesigning anything.
- 52. B — The same tool either compounds value or accumulates risk depending on the map.
- 53. D — Every step is classified into one of these three.
- 54. C — Reversibility, stakes, and accountability decide the classification.
- 55. B — Reversibility sets how much delegation the step can absorb.
- 56. A — High-cost steps do not become AI-appropriate just because they are mechanical.
- 57. A — Drafting can move to AI; accountability cannot.
- 58. B — Features are attached after the map, matched to what each step is.
- 59. C — The Skill removes the dependency on someone recalling the right phrasing; it does not remove variance or review.
- 60. D — Naming them as gates is what makes the redesign defensible.
- 61. A — All three criteria point the same way on a mechanical, reversible step.
- 62. B — A flag is reversible and the rules are exactly what a Skill can carry consistently.
- 63. D — Real AI work with human judgment on every edit is the collaborative pattern.
- 64. C — This is the step the second team in the introduction delegated, at the cost of the whole tool.
- 65. B — Numeric steps are AI-appropriate specifically through code execution.
- 66. D — Irreversibility and legal effect put this firmly with a person.
- 67. A — The map is not a timid one: it gives the AI genuine work and keeps the decisions human.
- 68. C — When the map hands an irreversible or high-accountability step to AI, that is over-delegation.
- 69. B — Quality at one step is not evidence about the risk profile of the next.
- 70. D — Success at step three is not evidence about step four.
- 71. C — The classification is only real if someone is actually doing the review.
- 72. A — The order matters: the work defines the map, and features attach to it afterward.
- 73. C — Integration is also a communication problem, aimed at people outside the build.
- 74. A — Limits stated as clearly as value is what makes the claim credible.
- 75. D — Overstatement is the named failure: the first visible error exposes it.
- 76. B — The accurate version names the work and the gate; the short version invites a question your first error answers badly.
- 77. C — High AI literacy earns the feature-level detail and the explicit failure mode.
- 78. B — Same workflow, same gate, different depth of explanation.

79. **A** — Calibrate to AI literacy and to what the audience actually has to decide.
80. **D** — Expectation-setting at the boundary is what prevents the later surprise.
81. **A** — Same precision, different audience: Description pointed at people instead of at the model.
82. **D** — 'Every output destined for a client passes human review' is the control that makes the workflow defensible.
83. **B** — Value and limits in one breath, with the gate and the measured result both named.
84. **C** — Each phrase replaces a defensible bounded claim with an inflated one; the repair is the same every time.
85. **B** — Extraction mirrors the contract-review clause step: mechanical and reversible, so it delegates cleanly.
86. **D** — This is the analogue of the playbook-departure flag: a repeatable rule check is what a Skill is for.
87. **C** — Numeric steps go to code execution; a Skill defines procedure, not verified arithmetic.
88. **A** — This matches the redline-and-rationale step: real AI drafting with human judgment on the wording.
89. **D** — Approval is the accountability step, exactly as in the contract-review map.
90. **B** — This is the 'sign and send' pattern: irreversible and external, so it stays human.
91. **C** — 'A large number' is unquantified, so there is nothing to build against or check completion on; the other three state a checkable condition.
92. **A** — Compute first, synthesize second — the capacity-plan pattern exactly.
93. **D** — Reliance as infrastructure is the escalation signal; it is no longer a prompt-and-iterate exercise.
94. **A** — Approval and signature are the high-stakes, irreversible, accountability-bearing steps; the other three are AI-appropriate.
95. **C** — It states what the tool does and names the human checkpoint, which is the pattern the lesson prescribes.
96. **B** — Takeaway 1: delegate deliberately, do not automate indiscriminately.
97. **D** — Messy inputs in, structured and traceable requirements out.
98. **B** — Takeaway 3: build plans on verified numbers.
99. **A** — Takeaway 4: map every step against the three criteria.
100. **C** — The module moves from structuring the work to grounding it, designing it, dividing it, and finally describing it honestly.