

Claude Certified Associate – Foundations (CCAO-F) — Module 5 Practice Exam

Configuration & Knowledge Management

100 scenario-style and recall questions, multiple choice (A/B/C/D), one correct answer each.

Covers: Configuring Claude Projects · Connectors and Uploaded Knowledge · System-Level Instructions That Stick · Maintaining Configurations.

Instructions: Select the single best answer for each question. An answer key with brief explanations is provided at the end of this document. Correct answers are distributed evenly across A, B, C and D. This practice material is based on user-provided course screenshots and is intended for self-study only — it is not an official Anthropic exam product.

A. Module Introduction

1. What is the line between 'using Claude' and 'operating Claude'?
 - A. The two terms describe the same activity at different points in a project
 - B. Using it means working in Chat while operating it means working through the API
 - C. Using it is for individuals while operating it requires an Enterprise plan
 - D. Using it means typing a good prompt today; operating it means building an environment where the right context, instructions, and procedures are already in place
2. Why is configuration described as leverage?
 - A. You set it up once and benefit from it on every conversation that follows
 - B. It lowers the usage cost of each individual message
 - C. It unlocks a higher model tier for the Project
 - D. It removes the need to review output before use
3. What does configuration do for a team?
 - A. It guarantees identical wording in every response
 - B. It restricts the Project so only its author can work in it
 - C. It turns individual skill into team capability — two people asking the same question get the same quality of answer
 - D. It removes the need for standing instructions once the team is trained
4. What happens in a team without configuration?
 - A. Claude refuses to answer without a knowledge base
 - B. Everyone re-invents context daily and the answers drift
 - C. The Project reverts to plain Chat behavior automatically
 - D. Output quality is unaffected as long as individual prompts are good
5. What is the second half of the configuration discipline?
 - A. Maintenance — configurations age, and stale instructions, superseded documents, or a drifted Skill quietly degrade output
 - B. Permissions — deciding who is allowed to open the Project
 - C. Model selection — matching a tier to each Project's workload
 - D. Escalation — deciding when a workflow needs developer involvement

6. How does the module describe a well-configured environment?
 - A. Built once and then left alone so it stays stable
 - B. Built deliberately and reviewed on a cadence
 - C. Rebuilt from scratch at the start of each quarter
 - D. Maintained automatically by Claude as the work changes
7. Which of the following is NOT one of Module 5's stated learning objectives?
 - A. Create effective system-level instructions
 - B. Configure Claude Projects with instructions and knowledge sources
 - C. Manage uploaded knowledge and connectors such as Google Drive and Gmail
 - D. Choose the model tier best suited to each Project's workload
8. What is 'the deal' in this module?
 - A. Consolidate every workstream into a single well-configured Project
 - B. Automate the configuration so it updates itself as processes change
 - C. Set up once, benefit on every conversation, then maintain it so it stays true
 - D. Replace prompting skill with configuration

B. Configuring Claude Projects

9. What is the actual skill in configuring a Project?
 - A. Writing the longest and most detailed standing instructions you can
 - B. Uploading as much reference material as possible into the knowledge base
 - C. Putting each piece of a recurring need into the right slot
 - D. Deciding between a Project and a Chat for each individual task
10. What are the four configuration mechanisms of a Project?
 - A. Instructions, connectors, artifacts, and Memory
 - B. Standing instructions, the knowledge base, Skills, and scoped Memory
 - C. Chat, Projects, Artifacts, and Research
 - D. Role, context, task, and constraints
11. How do the four mechanisms divide the work?
 - A. Instructions govern behavior, the knowledge base holds facts, Skills carry procedures, and scoped Memory keeps continuity
 - B. Instructions hold facts, the knowledge base governs behavior, Skills store Memory, and Memory carries procedures
 - C. All four hold facts and differ only in how long each persists
 - D. Instructions and knowledge are interchangeable, and Skills and Memory are optional extras
12. What belongs in the standing instructions slot?
 - A. The stakeholder names and preferences that accrue over time
 - B. The documents and reference files Claude should draw on when answering
 - C. Repeatable multi-step procedures for a particular task type
 - D. How Claude should behave across every conversation in the Project: tone, format defaults, verification habits — behavior, not facts

13. What belongs in the knowledge base?
- A. Decisions settled in earlier conversations within the Project
 - B. Tone and format defaults that should apply to every conversation
 - C. Procedures Claude should follow consistently for a given task type
 - D. The documents, policies, and reference files Claude should draw on without re-uploading — facts and reference, not behavior
14. What belongs in a Skill?
- A. Reference documents Claude consults while answering
 - B. Behavior rules such as 'always cite your sources'
 - C. Repeatable procedures Claude should follow consistently for a task type — procedure, not one-off instruction
 - D. Continuity of settled facts within a single Project
15. Where do Skills live relative to a Project?
- A. Inside the Project where they were created, and nowhere else
 - B. At the account level under Customize, not inside any one Project — a Skill you build is reusable across any Project that needs it
 - C. In the knowledge base of every Project that uses them
 - D. In scoped Memory, which is what makes them persist
16. What belongs in scoped Memory?
- A. Continuity within the Project, kept separate from your other Projects so context does not bleed between workstreams
 - B. A shared account-level store applied identically across all Projects
 - C. A full backup of every conversation held in the Project
 - D. Uploaded documents that change too often for the knowledge base
17. A rule such as 'always cite sources' belongs in which mechanism?
- A. A standing instruction, because it is a rule about behavior
 - B. The knowledge base, because it concerns source documents
 - C. A Skill, because it has to apply on every task
 - D. Scoped Memory, because it needs to persist across conversations
18. A fact such as 'our brand palette is these hex codes' belongs in which mechanism?
- A. A standing instruction, because it governs how output looks
 - B. Knowledge, because it is a fact Claude needs and cannot infer
 - C. A Skill, because branding is applied repeatedly
 - D. Scoped Memory, because it applies to one client only
19. A multi-step procedure such as 'format findings into our standard report template' belongs where?
- A. In a Memory entry recorded after the first successful report
 - B. In the standing instructions of each Project that produces reports
 - C. In each Project's knowledge base as a template document
 - D. In a Skill, built once at the account level under Customize and reused across any Project that needs it
20. What is named as the most common configuration mistake?
- A. Creating more than one Project for the same client
 - B. Enabling Memory before the knowledge base is complete
 - C. Putting a procedure into the instructions, or a behavior rule into the knowledge base — which makes the Project harder to maintain
 - D. Writing instructions that are too precise for general use

21. In the client account workspace, what were the standing instructions?
- A. Write in a formal register, always cite the source document for any factual claim, and flag anything you are unsure of rather than guessing
 - B. Use the firm's status-report formatter on every response
 - C. Remember the client's stakeholder names across all sessions
 - D. Draw on the brand guide and statement of work when answering
22. In that same workspace, what went into the knowledge base?
- A. The firm's status-report formatting procedure
 - B. The stakeholder names and the client's standing preferences
 - C. The client's brand guide, the current statement of work, and the last three status reports
 - D. The instruction to cite a source document for every factual claim
23. In that same workspace, what role did the Skill play?
- A. It was a Memory entry recording how the previous report had been formatted
 - B. It was a Project-specific procedure that could not be used for other clients
 - C. It was a document in the knowledge base describing the report format
 - D. The firm's status-report formatter lives under Customize at the account level and is available to any Project, so every weekly update comes out in the same structure
24. In that same workspace, what went into scoped Memory?
- A. The client's brand guide and current statement of work
 - B. The client's stakeholder names and standing preferences, which never appear in Client B's Project
 - C. The firm's status-report formatting procedure
 - D. The instruction to write in a formal register
25. What does the lesson say about Team or Enterprise use of such a Project?
- A. Sharing requires the knowledge base to be emptied first
 - B. Projects cannot be shared; each team member has to build their own
 - C. Sharing a Project merges its Memory into every member's account-level Memory
 - D. The Project can be shared with the engagement team under permissions, so everyone works from the same configured baseline
26. What makes the one-Project-per-client pattern safe?
- A. Using a different model tier for each client
 - B. Incognito mode applied to every client conversation
 - C. Project-scoped Memory, which guarantees Client A's content does not appear in a Client B conversation
 - D. Clearing the knowledge base between client sessions
27. How should Memory be regarded inside a Project?
- A. As a configuration slot like any other, not an afterthought
 - B. As a backup feature that is best disabled inside Projects
 - C. As a replacement for the knowledge base once enough conversations have happened
 - D. As something relevant only outside Projects
28. What does scoped Memory hold that the other mechanisms do not?
- A. The uploaded documents that support the Project's work
 - B. What the Project has already settled — stakeholder names, standing preferences, and decisions from earlier conversations you do not want to restate
 - C. The step-by-step procedures the Project runs regularly
 - D. The behavior rules that apply to every conversation

29. Why does the word 'scoped' matter?
- A. It means Memory only applies when the knowledge base is empty
 - B. It means Memory is capped at a fixed number of entries per Project
 - C. It means Memory expires at the end of each conversation
 - D. A Project's Memory is sectioned off from your other Projects, so content built up for one client or workstream never appears in another's
30. What is the simple rule for deciding between Memory and knowledge?
- A. Recent material is Memory and older material is moved into knowledge
 - B. Anything a client sends is Memory; anything produced internally is knowledge
 - C. Uploaded documents are knowledge; what accrues in conversation and gets decided in the room is Memory
 - D. Facts belong in Memory and documents belong in the instructions
31. What is the most common pattern when a need spans two mechanisms?
- A. A behavior rule plus the facts it acts on: 'always cite the source document' is a standing instruction, while the documents it cites live in the knowledge base
 - B. Two standing instructions, one covering tone and one covering format
 - C. A Skill duplicated into two Projects so both teams can use it
 - D. A Memory entry mirrored into the knowledge base as a backup
32. Why does neither half of that pairing work alone?
- A. The instruction overrides the documents, which makes them redundant
 - B. The instruction has nothing to cite without the documents, and the documents get used inconsistently without the instruction
 - C. The documents cannot be read unless a Skill parses them first
 - D. Neither functions unless Memory is enabled to connect them
33. How do Skills pair with the other mechanisms?
- A. A Skill contains both the procedure and all the reference material it needs
 - B. A status-report Skill carries the procedure while the brand guide it formats against sits in the knowledge base
 - C. A Skill replaces the knowledge base for any Project that uses it
 - D. A Skill and a knowledge base cannot both be active in one Project

C. Connectors and Uploaded Knowledge

34. What do connectors do?
- A. Extend Claude's reach into the data you already work in, such as Google Drive and Gmail
 - B. Replace the knowledge base with a live feed of every file you own
 - C. Grant Claude permission to act on external systems without limits
 - D. Increase the context window available inside a Project
35. What keeps connectors useful rather than frustrating?
- A. Refreshing each connection at the start of every conversation
 - B. Connecting every available source so nothing is ever missing
 - C. Managing them as curated sources and knowing exactly what each one can and cannot do
 - D. Using them only on the highest available model tier

36. What does a connector let Claude reach?
- A. Other users' Projects within the same account
 - B. Any system in your organization, whether or not you authorized it
 - C. The public web, replacing the need for web search
 - D. An external system you authorize — for example searching your Drive for a document or finding a relevant email
37. How should the set of connected sources be managed?
- A. Rotate connectors weekly to limit exposure
 - B. Connect everything available so Claude never lacks context
 - C. Connect nothing until an administrator approves each conversation
 - D. Keep what is accessible deliberate rather than connecting everything by default
38. What is a capability boundary?
- A. A limit that expands automatically as the model tier increases
 - B. A cap on how many documents a connector can return per request
 - C. Each connector has a defined limit on what it can do — a mail connector may let Claude search and read messages but not send them
 - D. A restriction that applies only on Team and Enterprise plans
39. What happens when you expect an action a connector cannot perform?
- A. It produces a confusing failure rather than a clear error, so learn each connector's boundaries before building a workflow on it
 - B. Claude returns an explicit permissions error naming the missing capability
 - C. The connector escalates the request to an administrator for approval
 - D. The action succeeds but the result is not saved anywhere
40. What is the 'wrong path to add a connector' pitfall?
- A. Adding a connector twice creates duplicate entries that conflict with each other
 - B. The obvious add-connector path can route to a public directory rather than your organization's vetted connectors
 - C. Connectors added from inside a Project apply only to that Project
 - D. The add path expires after a period and has to be regenerated
41. Where is confirming the connector path most important, and with whom?
- A. Only when the Project's knowledge base is still empty
 - B. On the free plan, where connectors are otherwise unsupported
 - C. Only when the connector handles email rather than files
 - D. On Team or Enterprise, confirmed with your admin, so you connect the approved source and not a look-alike
42. What is the 'boundary confusion' pitfall?
- A. Boundary limits silently delete material from the knowledge base
 - B. Hitting a connector's capability limit can look like a bug, so reports go to the wrong team and the fix stalls; knowing where the capability ends keeps the issue from being mislabeled
 - C. Boundary confusion is a documentation problem with no practical cost
 - D. Boundary limits change without notice, so they cannot usefully be learned

43. How are the two connector pitfalls characterized?
- A. As patterns reported from field use rather than documented product behavior
 - B. As documented behavior specified in the connector reference
 - C. As defects Anthropic has confirmed and scheduled for repair
 - D. As problems specific to third-party connectors only
44. What care does uploaded knowledge need?
- A. Weekly re-uploading of every file to keep it fresh
 - B. Only an initial upload, since documents do not age
 - C. The same care as a connected source: recent, relevant, and free of duplication
 - D. Consolidation into a single combined document
45. What is the risk of a knowledge base holding three versions of the same policy?
- A. It invites Claude to cite the wrong one
 - B. It improves accuracy by giving Claude more context to compare
 - C. It has no effect, because Claude always selects the newest file
 - D. It causes the Project to reject any further uploads
46. How should a knowledge base be curated?
- A. By moving reference documents into scoped Memory instead
 - B. By archiving nothing, so the full version history stays available
 - C. By keeping only the single most recent file of any type
 - D. The way you would a shared drive — removing deprecated versions as you add new ones
47. A team notices Claude quoting a travel policy that was replaced two months ago. The new policy is also in the knowledge base. What is the fix?
- A. Switch to a higher model tier so Claude selects the newer file
 - B. Add a standing instruction telling Claude which year to prefer
 - C. Remove the superseded version — a knowledge base holding both invites the wrong citation
 - D. Move both versions into scoped Memory
48. A team plans a workflow that assumes a connector can create calendar events. What should they do first?
- A. Build the workflow and handle any failures as they appear
 - B. Confirm the connector's capability boundary before building the workflow on it
 - C. Assume the capability exists, since the connector already reaches the calendar
 - D. Request a higher model tier to unlock the action

D. System-Level Instructions That Stick

49. What question do teams keep asking that persistent instructions answer?
- A. Whether Claude can write its own standing instructions
 - B. Whether the guardrails can be baked in once instead of re-typed every conversation
 - C. Whether instructions can be shared between separate accounts
 - D. Whether a higher model tier removes the need for instructions
50. What gets written once and then inherited?
- A. The Skills library, inherited by every account in the organization
 - B. The knowledge base contents, inherited by every Project on the account
 - C. The verification behavior, format defaults, and tone — inherited by every conversation in the Project
 - D. The conversation history, inherited by each new chat

51. Which instructions are the highest-value ones to write?
- A. The ones you would otherwise re-type constantly
 - B. The ones that describe the Project's purpose for new team members
 - C. The ones that specify which model tier to use
 - D. The ones that list every document in the knowledge base
52. What is the example verification instruction given in the lesson?
- A. Prefer the most recent document whenever two sources conflict
 - B. Always run code execution before answering any question
 - C. Escalate any uncertain answer to a human reviewer automatically
 - D. Cite the source document for every factual claim, and say 'I don't know' rather than guessing when the documents do not cover something
53. What does making that a standing instruction achieve?
- A. It guarantees that no factual claim will ever be wrong
 - B. The discipline applies to every conversation without anyone remembering to ask for it
 - C. It removes the need to review the output before use
 - D. It applies the same rule to every Project on the account
54. What does it mean to anticipate the use cases?
- A. Instructions should be written only after the first ten conversations reveal patterns
 - B. Instructions should stay general so they never constrain an unusual request
 - C. Good standing instructions embed format, tone, and guardrail guidance ahead of need, so the first draft lands closer to a final deliverable
 - D. Instructions should list the corrections made in previous conversations
55. Why are vague instructions especially dangerous?
- A. They apply only to the first message of each conversation
 - B. They trigger an error that halts the conversation
 - C. They cause Claude to ignore the knowledge base as well
 - D. They do not announce that they failed; they just quietly do not work
56. How does the lesson contrast 'be professional' with a precise instruction?
- A. 'Be professional' gives Claude almost nothing to act on, while 'use a formal register, define any acronym on first use, and keep paragraphs under four sentences' is precise enough to change the output
 - B. 'Be professional' is preferable because it adapts to each context
 - C. The two are equivalent, since Claude infers the same intent from either
 - D. The precise version is too restrictive to use in a Project
57. What is the test of a standing instruction?
- A. Whether it references at least one knowledge base document
 - B. Whether it fits on a single line
 - C. Whether Claude can restate it accurately when asked
 - D. Whether two different people would read it the same way
58. What was the 'vague' version in the worked example?
- A. 'Lead each report with a one-sentence headline' — too narrow to matter
 - B. 'Make the reports good and accurate' — output quality varies conversation to conversation and nothing concrete changed
 - C. 'Mark unverified figures rather than including them' — too strict to apply
 - D. 'Cite the source for every figure' — impossible to satisfy in practice

59. What was the 'precise' version in the worked example?
- A. 'For every figure in a report, state its source; if a figure is not in the provided data, mark it unverified rather than including it; lead each report with a one-sentence headline'
 - B. 'Make the reports good and accurate'
 - C. 'Write in whatever style best suits the audience'
 - D. 'Follow the company's usual reporting standards'
60. What changed once the precise instruction was in place?
- A. Claude began invoking code execution automatically
 - B. The reports became shorter but no more accurate
 - C. The verification behavior became consistent and the headline appeared every time
 - D. The knowledge base no longer required maintenance
61. A Project's standing instructions include 'keep it concise.' What is wrong with it?
- A. It is not testable — two readers would apply it differently, so it will silently fail to change output
 - B. It is too long to work as a standing instruction
 - C. Conciseness cannot be specified at the Project level
 - D. It belongs in the knowledge base rather than in the instructions
62. A Project produces client deliverables that always need the same two corrections. What is the right move?
- A. Raise the model tier so the corrections become unnecessary
 - B. Keep correcting each draft manually, since instructions cannot control format
 - C. Move the deliverable format into scoped Memory
 - D. Write the format and register into the standing instructions so the first draft lands closer to final
63. A Project produces meeting-summary notes. Which standing instruction is most likely to actually change the output?
- A. 'Write up the notes nicely'
 - B. 'Capture the meeting clearly and thoroughly'
 - C. 'List each decision as its own bullet; record every action item as owner plus due date; flag any unassigned action as owner TBD'
 - D. 'Summarize each meeting as best you can'

E. Maintaining Configurations

64. What does it mean that configurations are living assets?
- A. Configurations stay valid until Claude reports an error against them
 - B. Instructions, knowledge, Skills, and Memory all drift toward stale, and stale configuration degrades output quietly with no error to alert you
 - C. Only the knowledge base ages; instructions and Skills are static
 - D. Drift shows up immediately as failed or refused responses
65. Why schedule maintenance rather than react to problems?
- A. It resets the context window for long-running conversations
 - B. It is required to stop the Project from being archived
 - C. It is how you catch decay before it reaches a deliverable
 - D. It is the only way to add new documents to a knowledge base

66. What does the review cadence check?
- A. Whether the model tier is still the cheapest one that works
 - B. Whether the standing instructions still match the current process, whether the knowledge base is free of superseded documents, and whether the right Skills are enabled
 - C. Whether the Project has been used within the past week
 - D. Whether Memory has approached its storage limit
67. What cadence does the lesson recommend?
- A. A review only once a stakeholder complains
 - B. An annual review is sufficient for any Project
 - C. A review after every conversation
 - D. A monthly pass for active Projects catches most drift
68. What is the signal that you waited too long between reviews?
- A. Output quality slipping for no visible reason
 - B. An error message naming the stale component
 - C. The Project refusing to accept new uploads
 - D. A noticeable drop in response speed
69. How does Skills versioning work?
- A. Skills update only when the Project that uses them is reviewed
 - B. All Skills update automatically, including custom uploads
 - C. No Skills update automatically; every Skill has to be re-uploaded
 - D. Anthropic-built and organization-provisioned Skills update automatically, while your own custom-uploaded Skills change only when you re-upload them
70. A Skill silently produces a slightly off format on every run. How should this be classified?
- A. As a prompting problem, fixed by rewording the request each time
 - B. As a maintenance problem, not a prompting problem
 - C. As evidence that the model tier is too low for the task
 - D. As a reason to delete the Skill and replace it with a standing instruction
71. How should Memory be treated in a Project?
- A. Like a cache, cleared at the end of every session
 - B. Like an archive, preserving every entry for auditability
 - C. Like a working file: review it periodically, edit or delete entries that have gone stale, and export it as a backup before a major change
 - D. Like a knowledge base, with documents uploaded directly into it
72. When is a full Memory reset the right call?
- A. When a Project's Memory has accumulated enough outdated content to mislead
 - B. Whenever a single stale entry is discovered
 - C. At the start of every quarter, regardless of content
 - D. Never — a reset loses continuity permanently
73. How does the lesson weigh accuracy against volume in Memory?
- A. Volume matters more, since more retained context produces better answers
 - B. The accuracy of what is stored matters more than the volume
 - C. Both matter equally and should be balanced against each other
 - D. Neither matters once the knowledge base is current

74. What does it mean that configuration is also additive?
- A. New Skills should be added rather than existing ones edited
 - B. Every new conversation should contribute a new standing instruction
 - C. Knowledge bases should only grow, never shrink
 - D. When Claude has not automatically captured a fact that matters — a constraint, a preference, a piece of background — add it explicitly to standing instructions or Memory rather than re-supplying it each session
75. In the degraded-setup audit, what was wrong with the standing instruction?
- A. It still referenced a metric the team had renamed last quarter
 - B. It had been written too precisely for the current process
 - C. It duplicated a rule already present in the knowledge base
 - D. It had been deleted when the Project was shared with the team
76. In that same audit, what was wrong with the knowledge base?
- A. It contained documents from a different Project
 - B. It was empty, so Claude had nothing to draw on
 - C. It held two versions of the template
 - D. It had grown too large for the Project to index
77. In that same audit, what was wrong with Memory?
- A. A Memory entry recorded a stakeholder who had since left
 - B. Memory had been disabled for the Project
 - C. Memory had been shared across all of the account's Projects
 - D. Memory contained no entries at all
78. What was the fix in the degraded-setup example?
- A. Rewriting the prompts the team uses each week
 - B. Maintenance, not a new prompt: update the instruction, remove the old template, and delete the stale Memory entry
 - C. Rebuilding the Project from scratch with fresh configuration
 - D. Raising the model tier until output returned to standard
79. What was the result after that fix?
- A. Memory had to be disabled to prevent a recurrence
 - B. Output improved, but only after the team adopted new prompts
 - C. The Project had to be re-shared with the team to take effect
 - D. Output returned to standard without changing how anyone prompts
80. Two people get noticeably different quality of answer from the same Project. What does the module suggest?
- A. Memory should be disabled so neither has an advantage
 - B. One of them is working on a lower model tier
 - C. The configuration is not carrying the load — a configured environment is what makes the same question produce the same quality of answer
 - D. Each of them should build a separate Project
81. A custom Skill built six months ago now produces a slightly outdated format. What is the fix?
- A. Add a standing instruction that corrects the format afterward
 - B. Wait for it to update automatically, as Anthropic-built Skills do
 - C. Re-upload the updated Skill, since custom-uploaded Skills change only when you re-upload them
 - D. Delete the Project and rebuild it from scratch

82. A Project's Memory has become a mix of current and long-outdated facts, and a large clean-up is planned. What does the lesson suggest?
- A. Move the Memory contents into the knowledge base before editing
 - B. Export Memory as a backup first, then edit or delete the stale entries, resetting fully if the outdated content is enough to mislead
 - C. Leave it alone, since Memory self-corrects as the Project is used
 - D. Turn Memory off permanently for that Project
83. Why is stale configuration described as especially dangerous?
- A. It degrades output quietly, with no error to alert you
 - B. It causes conversations in the Project to fail outright
 - C. It locks the Project until an administrator intervenes
 - D. It affects only numeric outputs, which are hardest to check

F. Quiz & Mixed Review

84. A team wants every report in the Project to follow the same multi-step formatting procedure. Which configuration mechanism fits?
- A. A Memory entry recording how reports should look
 - B. A standing instruction describing the format in prose
 - C. A document in the knowledge base showing the format
 - D. A Skill that carries the repeatable formatting procedure
85. An equity analyst covers two competing retailers and must ensure each company's confidential modelling assumptions never surface in work on the other. What is the correct setup?
- A. A different model tier for each company's analysis
 - B. A single Project covering both companies, separated by careful prompting in each chat
 - C. Turning Memory off entirely so nothing is retained for either company
 - D. A separate Project for each company, each with its own scoped Memory holding that company's assumptions
86. A user expects a mail connector to send emails on their behalf, but it only searches and reads. What is the most accurate understanding?
- A. Each connector has a defined capability boundary, and sending is outside this one's scope
 - B. The connector is broken and should be reported as a bug
 - C. A more capable model would enable the sending action
 - D. The connector needs to be re-added through the public directory
87. A monthly-report Project has started producing subtly outdated figures with no error message. The standing instruction pins the report to the FY25 template and FY25 targets, the knowledge base holds the FY26 targets and FY26 template, and a Memory entry states the reporting period is FY25. The period is now FY26. What is driving the drift?
- A. The report template, which should be deleted so Claude stops mixing versions
 - B. The knowledge base, which should have the brand guide removed as unrelated to the figures
 - C. The standing instruction and the Memory entry — both still pin the report to FY25 while the knowledge base already holds the FY26 versions, so update both to FY26
 - D. Nothing in the setup — subtle drift with no error means the model tier is too low

88. What does that FY25/FY26 scenario illustrate about maintenance?
- A. Knowledge base updates always override instructions and Memory
 - B. Uploading current documents is not enough — instructions and Memory still pointing at the old period keep the drift alive
 - C. Memory takes precedence over the knowledge base by design
 - D. Drift of this kind produces an error message identifying the cause
89. According to Key Takeaway 1, what is configuration?
- A. A feature that matters mainly for Enterprise accounts
 - B. A one-time setup task that does not need revisiting
 - C. Leverage — set up once, benefit on every conversation; a configured environment separates operating Claude from merely using it
 - D. A substitute for prompting skill
90. According to Key Takeaway 2, how should needs be matched to mechanisms?
- A. Instructions for facts, knowledge for procedures, Skills for continuity, Memory for behavior
 - B. Instructions for behavior, knowledge for facts, Skills for procedures, scoped Memory for continuity
 - C. All four are interchangeable provided each is written carefully
 - D. Use whichever mechanism is quickest to edit at the time
91. According to Key Takeaway 3, what should you know about every connector?
- A. Its boundary — connectors extend Claude's reach, but each has a capability edge, and knowing it prevents frustration and misdirected fixes
 - B. Its latency, so workflows can be timed correctly
 - C. Its storage limit, so the knowledge base does not overflow
 - D. Its model tier requirement, since capabilities vary by tier
92. According to Key Takeaway 4, what distinguishes an instruction that works?
- A. Placement — instructions matter less than what is in the knowledge base
 - B. Brevity — short instructions outperform long ones in every case
 - C. Freshness — instructions should be rewritten before each new conversation
 - D. Precision — vague standing instructions silently fail, while precise, testable ones change output every conversation
93. According to Key Takeaway 5, what happens if configurations are not maintained?
- A. Decay announces itself through slower responses
 - B. They stay stable once tested and rarely need review
 - C. They age, and output degrades with no warning unless reviews of instructions, knowledge, Skills, and Memory are scheduled
 - D. Only Memory is affected, so only Memory needs review
94. A Project's answers have drifted, though nobody has changed how they prompt. What is the most likely cause?
- A. The configuration has aged — an instruction, a document, a Skill, or a Memory entry has gone stale
 - B. The model tier changed automatically without notice
 - C. The context window was exceeded within a single conversation
 - D. The Project needs to be re-shared with the team

95. Which of these belongs in scoped Memory rather than the knowledge base?
- A. The client's signed statement of work
 - B. The stakeholder preferences and decisions settled in earlier conversations in that Project
 - C. The company's brand guide
 - D. The last three status reports
96. Which of these belongs in a Skill rather than a standing instruction?
- A. A requirement to flag uncertainty rather than guess
 - B. A rule that every factual claim must cite its source
 - C. A preference for a formal register in all written output
 - D. A multi-step procedure for turning findings into the firm's standard report template
97. Which of these belongs in a standing instruction rather than the knowledge base?
- A. A rule that Claude should flag anything it is unsure of rather than guessing
 - B. The current statement of work
 - C. The brand palette hex codes
 - D. The last three status reports
98. Why does one Project per client work for a consultant with competing clients?
- A. Knowledge bases are automatically encrypted per client
 - B. Each Project can be set to a different model tier
 - C. Memory is disabled by default inside client Projects
 - D. Project-scoped Memory isolates each client's continuity, so nothing bleeds between workstreams
99. A behavior rule and the facts it acts on have to be configured together. Which pairing is correct?
- A. Both the rule and the documents recorded as Memory entries
 - B. 'Always cite the source document' in the knowledge base, with the documents written into the standing instructions
 - C. 'Always cite the source document for factual claims' as a standing instruction, with the documents it cites in the knowledge base
 - D. Both the rule and the documents packaged inside a single Skill
100. Which best summarizes the throughline connecting the four lessons of Module 5?
- A. Consolidate every workstream into a single well-maintained Project
 - B. Match each recurring need to the right mechanism, extend reach through connectors whose boundaries you know, write instructions precise enough to be testable, and review the whole configuration on a cadence so it does not decay
 - C. Prefer configuration over prompting, since a configured Project needs no prompting skill
 - D. Automate the configuration so that maintenance becomes unnecessary

Answer Key & Explanations

1. **D** — Operating Claude means every conversation starts from a configured baseline instead of a blank slate.
2. **A** — One setup investment pays out across every later conversation in that space.
3. **C** — Curated Project, instructions, and knowledge are what make quality reproducible across people.
4. **B** — Without a shared configured baseline, context gets rebuilt every session and quality varies by person.
5. **A** — Building the environment is half the job; keeping it true is the other half.
6. **B** — Deliberate construction plus scheduled review is the standard the module sets.
7. **D** — Model selection belongs to Module 1; this module covers configuration, connectors, instructions, and maintenance.
8. **C** — Configure instructions, knowledge, and scoped Memory; manage connectors; write instructions that hold; and run the maintenance reviews.
9. **C** — A Project has several configuration slots, and the skill is matching each need to the right one.
10. **B** — These four slots are what a Project offers, and each one takes a different kind of content.
11. **A** — Behavior, facts, procedure, continuity — one mechanism each.
12. **D** — Instructions are the behavior slot.
13. **D** — Knowledge is the facts-and-reference slot.
14. **C** — Skills are the procedure slot.
15. **B** — Account-level and reusable: this is why a Skill is not a per-Project configuration item.
16. **A** — Scoped Memory is the continuity slot, isolated per Project.
17. **A** — Behavior rules are instructions, regardless of what they are about.
18. **B** — Facts go in the knowledge base.
19. **D** — Multi-step procedure is the definition of a Skill, and it is built once rather than per Project.
20. **C** — Slot confusion is the common error, and its cost shows up at maintenance time.
21. **A** — Three behavior rules: register, citation, and uncertainty handling.
22. **C** — Documents and reference material, uploaded once.
23. **D** — Account-level Skill, reused inside this Project among others.
24. **B** — Continuity that accrues in use, isolated to that client's Project.
25. **D** — Shared Projects are how a configured baseline becomes a team asset.
26. **C** — The separation between Projects is what makes client-specific continuity safe to keep.
27. **A** — Memory is a first-class mechanism alongside instructions, knowledge, and Skills.
28. **B** — Memory carries what accrues in use, as distinct from what was uploaded or written up front.
29. **D** — That isolation is what makes Memory safe for sensitive or client-specific continuity.
30. **C** — Source of the content decides the slot: uploaded versus accrued.
31. **A** — Instruction plus knowledge is the classic pairing.
32. **B** — Each half depends on the other to have any effect.
33. **B** — Procedure in the Skill, reference in knowledge — the same two-slot pattern.
34. **A** — Connectors are powerful reach into existing systems, and they have boundaries.
35. **C** — Curation plus knowledge of the boundary is the stated condition.
36. **D** — Authorization defines the reach; you manage what is accessible.
37. **D** — Deliberate rather than default-everything is the guidance.
38. **C** — Every connector has a defined edge, and the edge is capability, not just data volume.

- 39. **A** — The failure is confusing rather than clear, which is why the boundary has to be learned up front.
- 40. **B** — The look-alike public entry is the trap; the vetted organizational source is the target.
- 41. **D** — Organizational plans are where a vetted-versus-public mismatch actually bites.
- 42. **B** — A capability edge misread as a defect wastes the team's time on the wrong fix.
- 43. **A** — The lesson flags them as field-observed, which is why the guidance is to confirm current behavior.
- 44. **C** — Uploaded material is a curated source too, not a one-time deposit.
- 45. **A** — Duplicate versions are a citation hazard, not extra context.
- 46. **D** — Add and remove together; the base stays current only if superseded material leaves.
- 47. **C** — This is the duplicate-versions problem exactly; curation is the repair.
- 48. **B** — Reach is not the same as capability; expecting an unsupported action produces a confusing failure.
- 49. **B** — Persistent instructions are exactly that mechanism.
- 50. **C** — Behavior written once at the Project level applies to every conversation inside it.
- 51. **A** — Anything you keep re-typing is a candidate for a standing instruction.
- 52. **D** — It combines the citation requirement and permitted uncertainty from Module 3, made persistent.
- 53. **B** — Persistence is the gain: nobody has to remember to request it.
- 54. **C** — Specifying format and register up front removes the same corrections being made every time.
- 55. **D** — Silent failure is the hazard: nothing signals that the instruction had no effect.
- 56. **A** — Precision is what converts an instruction into a behavior change.
- 57. **D** — Shared readings mean the instruction is testable rather than aspirational.
- 58. **B** — It sounds like a standard but specifies nothing, so nothing changes.
- 59. **A** — Three testable behaviors: source per figure, unverified marking, and a headline.
- 60. **C** — Consistency across conversations is the observable result.
- 61. **A** — Vague instructions do not announce their failure; the two-reader test catches them.
- 62. **D** — Anticipating the use case is exactly what standing instructions are for.
- 63. **C** — Only one of these is precise enough that two people would produce the same result from it.
- 64. **B** — Silent degradation is the defining property, which is why maintenance has to be scheduled.
- 65. **C** — Scheduling is what puts the check ahead of the deliverable.
- 66. **B** — Three checks: instructions against process, knowledge against currency, Skills against need.
- 67. **D** — Monthly for active Projects is the stated cadence.
- 68. **A** — Unexplained quality slippage is the symptom, because nothing else announces the decay.
- 69. **D** — The split matters: your custom Skills are the ones that silently go out of date.
- 70. **B** — The fix is in the configuration, not in how anyone phrases the request.
- 71. **C** — Working file, not archive: reviewed, edited, and backed up before big changes.
- 72. **A** — A reset is a real option once the accumulated staleness outweighs the continuity.
- 73. **B** — This carries the Module 1 Memory guidance into the Project context.
- 74. **D** — Maintenance removes stale content and also adds what was never captured.
- 75. **A** — The instruction had aged out of alignment with the process it described.
- 76. **C** — Duplicate template versions are the same citation hazard flagged in the connectors lesson.
- 77. **A** — A stale entry that was accurate when written is exactly the misleading case.
- 78. **B** — Three configuration repairs, none of them a prompting change.
- 79. **D** — That is the point of the example: the problem was never in the prompting.
- 80. **C** — Reproducible quality across people is the stated benefit of configuration.

- 81. **C** — Only Anthropic-built and organization-provisioned Skills update on their own.
- 82. **B** — Export before a major change is the stated safeguard.
- 83. **A** — No alert fires, so the only symptom is quality that slips for no visible reason.
- 84. **D** — Multi-step repeatable procedure is the definition of a Skill.
- 85. **D** — Project-scoped Memory is the mechanism that keeps one workstream's context out of another's.
- 86. **A** — This is a boundary, not a defect — and misreading it is the pitfall the lesson names.
- 87. **C** — Two of the three slots still point at the old period; the knowledge base is already current.
- 88. **B** — Maintenance has to cover every slot, because the stale ones override the current ones in practice.
- 89. **C** — Takeaway 1: configuration is leverage.
- 90. **B** — Takeaway 2: match each need to the right mechanism.
- 91. **A** — Takeaway 3: know each connector's boundary.
- 92. **D** — Takeaway 4: write instructions precisely.
- 93. **C** — Takeaway 5: maintain, or watch quality decay.
- 94. **A** — Quality slipping with no visible cause is the named signal of overdue maintenance.
- 95. **B** — Uploaded documents are knowledge; what accrues and gets decided in conversation is Memory.
- 96. **D** — Procedure goes in a Skill; the other three are behavior rules.
- 97. **A** — Behavior in instructions; facts and documents in knowledge.
- 98. **D** — The scoping is the safety mechanism that makes the pattern viable.
- 99. **C** — Rule in instructions, facts in knowledge — two slots working together.
- 100. **B** — Build the environment deliberately, know its edges, state its rules precisely, and keep it true over time.