

Claude Certified Associate – Foundations (CCAO-F) — Module 6 Practice Exam

Governance, Risk & Responsible Use

100 scenario-style and recall questions, multiple choice (A/B/C/D), one correct answer each.

Covers: Appropriate vs Inappropriate Use Cases · Skill Trust and Feature-Level Risk · Data Sensitivity, Privacy & Feature Controls · Organizational Policies and Diligence as a Habit · Ethical Implications: Bias, Fairness, Transparency.

Instructions: Select the single best answer for each question. An answer key with brief explanations is provided at the end of this document. Correct answers are distributed evenly across A, B, C and D. This practice material is based on user-provided course screenshots and is intended for self-study only — it is not an official Anthropic exam product.

A. Module Introduction

1. According to the module's opening line, what can freeze an entire organization's AI program?
 - A. A model tier downgrade mid-project
 - B. A single failed prompt in Chat
 - C. A missed monthly maintenance review
 - D. One inappropriate use case
2. Which triggers does the module name as capable of causing that kind of freeze?
 - A. Sensitive data uploaded to the wrong place, an untrusted Skill granted broad access, or a quiet policy violation at the wrong moment
 - B. A conversation exceeding its context window
 - C. A Project left without standing instructions
 - D. A Skill that updates automatically
3. How is governance exercised, according to the module?
 - A. By the model itself, which refuses inappropriate requests automatically
 - B. By a compliance team that reviews every conversation before it is sent
 - C. By practitioners, one decision at a time, not by a policy binder on a shelf
 - D. By an annual policy review conducted at the account level
4. Which two AI Fluency Framework competencies govern the module's decisions?
 - A. Discernment and Description
 - B. Diligence and Delegation
 - C. Diligence and Discernment
 - D. Delegation and Description
5. What does Diligence contribute to this module's framing?
 - A. Ownership and verification
 - B. The criteria for judging whether a use case is appropriate at all
 - C. The habit of matching strategy to task type
 - D. The habit of making prompt components explicit
6. What does Delegation supply to this module's framing?
 - A. The four-mechanism configuration model
 - B. Ownership and verification
 - C. The criteria for judging whether a use case is appropriate at all
 - D. The three-reference evaluation protocol

7. According to 'the deal in this module,' what does responsible use actually consist of?
- A. A fixed list of banned use cases published by the organization
 - B. A single onboarding certification completed once
 - C. A checklist an administrator runs before granting account access
 - D. A set of judgments you make every day: classifying use cases, vetting Skills, handling data by sensitivity, applying policy as a habit, and evaluating output for bias and fairness

B. Appropriate vs Inappropriate Use Cases

8. How does the lesson describe deciding whether a use case is appropriate?
- A. A one-time approval granted by an administrator
 - B. A structured evaluation, not a gut feeling
 - C. A judgment best left to the most senior person in the room
 - D. A decision that only applies to high-stakes work
9. What do the same Delegation criteria used to map workflow steps also do?
- A. Determine which entry point to use
 - B. Set the model tier for a task
 - C. Screen whole use cases
 - D. Decide whether a Skill should be enabled
10. What does the Reversibility criterion ask?
- A. Who is answerable for the outcome?
 - B. What is the cost if the output is wrong?
 - C. Does the task require judgment, relationship, or care that AI cannot supply?
 - D. Can a wrong output be caught and undone before it causes harm? Irreversible consequences raise the bar sharply
11. What does the Consequence of error criterion ask?
- A. Can a wrong output be caught and undone before it causes harm?
 - B. What is the cost if the output is wrong? Higher consequence demands more human control or rules the use case out
 - C. Does the task require human creativity or empathy?
 - D. Who is answerable for the outcome, and can that be exercised over an AI result?
12. What does the Need for human creativity or empathy criterion ask?
- A. Does the task require judgment, relationship, or care that AI cannot supply? Some work should stay human regardless of capability
 - B. What is the cost if the output is wrong?
 - C. Can a wrong output be caught and undone?
 - D. Who is answerable for the outcome?
13. What does the Accountability criterion ask?
- A. Can the output be caught and undone before it causes harm?
 - B. Does the task require judgment or care that AI cannot supply?
 - C. What is the cost if the output is wrong?
 - D. Who is answerable for the outcome, and can that accountability be exercised over an AI-produced result?

14. Are the four criteria a checklist where any single failure ends the discussion?
- A. Yes — but only Accountability can end the discussion by itself
 - B. Yes — failing any one criterion makes the use case automatically inappropriate
 - C. No — they interact, and a task can pass one criterion cleanly while still failing on another
 - D. No — only Reversibility and Consequence interact; the other two are independent
15. What is the practical rule for handling the four criteria's interaction?
- A. Apply only the criterion with the highest stakes and ignore the rest
 - B. Run all four, then ask which one is load-bearing for this specific use case
 - C. Apply the criteria in a fixed order and stop at the first failure
 - D. Average the four criteria into a single risk score
16. What is the load-bearing criterion?
- A. The one that, if it were changed, would move the use case between classifications
 - B. The criterion listed first in the table
 - C. The criterion that always applies regardless of use case
 - D. The criterion a compliance reviewer names, rather than the practitioner
17. What are the three classifications a use case sorts into?
- A. Ready to use, needs revision, needs human override
 - B. Fully appropriate, appropriate with human review, and inappropriate
 - C. AI-appropriate, human-retained, collaborative
 - D. Green, yellow, red
18. What defines 'fully appropriate'?
- A. Any task a Skill has already been built for
 - B. Useful for AI assistance, but stakes or accountability require a human gate
 - C. The consequence or irreversibility means AI should not perform this at all
 - D. Reversible, low consequence, no special human element — delegate with normal review
19. What defines 'inappropriate'?
- A. Useful for AI assistance, but stakes or accountability require a human gate
 - B. Reversible, low consequence, no special human element
 - C. The consequence, irreversibility, or human-element requirement means AI should not perform this — articulate why and name the human role that must own it
 - D. Any task involving personal or regulated data
20. Which classification do people most often get wrong, and why?
- A. 'Appropriate with human review' — because they stop at the label and never specify the gate
 - B. 'Fully appropriate' — because they underestimate reversibility
 - C. 'Inappropriate' — because they overestimate consequence
 - D. All three are equally likely to be misclassified
21. What does a properly defined review gate establish?
- A. Who built the Skill, what data it accesses, and when it was last updated
 - B. Who reviews (the role with accountability, not whoever is free), what they verify (the specific risk the review exists to catch), and when in the workflow the review happens (before the output is used, not after)
 - C. Which model tier is required and how long review should take
 - D. Whether the use case is reversible and what its consequence tier is

22. Which of these is an example of a properly defined gate?
- A. 'Someone will check it before it goes out'
 - B. 'We will keep a human in the loop'
 - C. 'A manager reviews the shortlist for adverse-impact patterns before any candidate is contacted'
 - D. 'The team will review outputs periodically'
23. What should you conclude if you cannot state a gate in who/what/when form?
- A. The gate can be left implicit and added later
 - B. The use case should default to fully appropriate
 - C. The use case should be escalated to a higher model tier
 - D. The use case is not yet ready to run
24. In the worked use-case portfolio, how was 'Draft internal FAQ from approved policy docs' classified?
- A. Fully appropriate — reversible, low stakes, factual grounding available
 - B. Appropriate with human review — factual accuracy needs checking
 - C. Inappropriate — policy documents are sensitive by default
 - D. Appropriate with human review — FAQs face customers directly
25. In that same portfolio, how was 'Summarize candidate resumes to a shortlist' classified?
- A. Appropriate with human review, load-bearing criterion is reversibility
 - B. Fully appropriate — resume summarization is purely mechanical
 - C. Inappropriate — hiring decisions cannot involve AI at any stage
 - D. Appropriate with human review, and only with bias safeguards — consequential and fairness-sensitive, with accountability resting with the hiring manager
26. In that same portfolio, how was 'Generate a final medical or legal determination' classified, and why?
- A. Inappropriate — irreversible consequence and a professional accountability that cannot transfer to a tool
 - B. Appropriate with human review — a licensed professional can sign off afterward
 - C. Fully appropriate — the determination can be corrected if wrong
 - D. Appropriate with human review, load-bearing criterion is reversibility
27. In that same portfolio, what was the load-bearing criterion for 'Draft customer-facing responses to billing complaints,' the case described as sitting on the line?
- A. Reversibility — a sent response can always be retracted
 - B. Accountability — the company owns what it tells a customer about their money
 - C. Need for human creativity or empathy — billing responses require warmth
 - D. Consequence of error — billing errors are always catastrophic

C. Skill Trust and Feature-Level Risk

28. What is a Skill, according to this lesson?
- A. A model tier optimized for a specific task type
 - B. A static document describing a procedure with no ability to act
 - C. Software — it can access what you give Claude access to during a session and can take actions through code execution
 - D. A configuration setting that only affects tone and format

29. What is the practical answer to 'Skills are a black box,' according to the lesson?
- A. Testing each Skill in Incognito mode before wider use
 - B. Disabling all third-party Skills across the organization
 - C. Running every Skill on the lowest model tier as a precaution
 - D. Not blind trust or a blanket ban, but a repeatable trust evaluation
30. What are the three questions of the Skill trust evaluation?
- A. Who published this Skill; what could it reach in the sessions where it runs, and is that exposure proportional; and is this Skill the right tool for the task at all
 - B. How many users have enabled this Skill; what is its average rating; and how often is it updated
 - C. What model tier does it require; how much does it cost to run; and is it compatible with every entry point
 - D. Who is legally liable if it fails; how is it licensed; and has it passed a security certification
31. Why is 'what permissions does this Skill request' the wrong question to ask?
- A. It is not the wrong question; permissions are exactly what should be checked
 - B. Skills do not request permissions — a Skill inherits whatever access the session it runs in already has
 - C. Because permissions can only be checked by an administrator
 - D. Because permission requests apply only to connectors, not Skills
32. What should you ask instead of 'what permissions does this Skill request'?
- A. What model tier does this Skill require to function?
 - B. How many other users have already enabled this Skill?
 - C. What could this Skill reach in the sessions where it will run, and is that exposure proportional to the task?
 - D. How often does this Skill's publisher release updates?
33. Why is 'internal' not the same as 'vetted'?
- A. Internal Skills are technically incapable of broad access
 - B. Internal Skills are automatically reviewed by IT before anyone can enable them
 - C. Internal Skills cannot access data outside their own team's Project
 - D. A Skill built by another team inside your own organization may have been given broad permissions for that team's convenience, or built against an older policy
34. How should you treat an internal Skill built by a different team?
- A. As automatically higher-risk than a vetted third-party Skill
 - B. As automatically trustworthy, since it originated inside the organization
 - C. The way you would treat software from a sister department — confirm with the publisher what it accesses and why, and check its permissions still match current policy
 - D. As exempt from review, since only external Skills need vetting
35. What is the trust question the lesson reduces all of this to?
- A. Do I know what it does, and does its access match the job?
 - B. Who is legally liable if this Skill fails?
 - C. How much does this Skill cost to run?
 - D. Is this Skill built on the newest model tier?
36. What are the three outcomes of a trust check?
- A. Enable, monitor, or archive
 - B. Enable, escalate, or decline
 - C. Approve, flag, or delete
 - D. Trust, verify, or reject

37. When does a Skill get enabled outright?
- A. When its permissions are unknown but the task is low-stakes
 - B. When it has been used successfully by at least one other team
 - C. When the source, permissions, and appropriateness are all clear
 - D. When it comes from any Anthropic-provided source, regardless of the task
38. When should a Skill be escalated rather than enabled or declined?
- A. When it requires code execution to function
 - B. When it is clearly inappropriate and should never be used
 - C. When it has already been vetted by another team
 - D. When it is useful but the source is unknown or the permissions look broad — route it to an admin or security function for review
39. When should a Skill be declined outright?
- A. When the permissions are clearly disproportionate or the source cannot be established, and no review would change that
 - B. Whenever the Skill was not built by Anthropic
 - C. Whenever the Skill has not yet been used by another team
 - D. Whenever the Skill requires more than one permission
40. In the worked example, how was the Anthropic-provided document formatter evaluated?
- A. Source trusted, but permissions too broad for the task — escalate it
 - B. Source trusted, permissions match the task, appropriate for the need — enable it
 - C. Source unknown, but the task is low-stakes — enable it
 - D. Source trusted, but a better tool exists for the task — decline it
41. In the worked example, how was the third-party 'analytics booster' evaluated?
- A. Unknown source, but permissions match the task — enable it with monitoring
 - B. Unknown source, disproportionate permissions for the stated purpose — do not enable without organizational review
 - C. Known source, disproportionate permissions — escalate for a permissions trim
 - D. Unknown source, minimal permissions — enable it cautiously
42. How does the lesson say the Skill-vetting rule generalizes?
- A. Apply Skill vetting only to Skills from outside the organization
 - B. Apply Skill vetting only to Skills that request code execution
 - C. Evaluate a Skill's source and permissions like any software you were about to install on a work machine
 - D. Treat Skill vetting as a one-time step that does not need repeating
43. Before enabling any feature that can read your data or act on it — a connector, a tool, an integration — what three questions should you ask?
- A. What data does it store, how long is it retained, and can it be exported?
 - B. What does it cost, how fast does it run, and which model tier does it require?
 - C. Who published it, how many users have it enabled, and is it open source?
 - D. Who provides it, what does it access, and is that access proportional to what you need?
44. What principle does the lesson name for granting access to any capability?
- A. Least privilege — grant the narrowest access that lets the job get done, and revisit it when the job changes
 - B. Maximum utility — grant the broadest access so the capability never becomes a bottleneck
 - C. Default trust — grant full access unless a specific reason exists to restrict it
 - D. Fixed grant — set access once and leave it unchanged regardless of how the job evolves

D. Data Sensitivity, Privacy & Feature Controls

45. What should you know before any data enters any feature?
- A. Its file format
 - B. Its sensitivity
 - C. Its file size
 - D. Its upload date
46. What is data classification described as?
- A. A setting configured once at the account level
 - B. A one-time compliance certification
 - C. The habit that prevents leakage
 - D. A feature exclusive to Enterprise plans
47. How do feature-specific controls relate to data classification?
- A. They determine the classification automatically
 - B. They replace the need for classification entirely
 - C. They apply only after a leak has already occurred
 - D. They are how you act on the classification when persistence or processing is not appropriate
48. What is the safe default when you are unsure whether data is safe to upload?
- A. Ask before uploading
 - B. Upload it and correct course if a problem appears
 - C. Redact it and upload the redacted version regardless of the task
 - D. Use Incognito mode, which resolves any uncertainty
49. What defines the Green tier?
- A. Internal documents not meant to leave the company
 - B. Published material, anonymized or aggregated data, or internal documents already cleared for wide sharing
 - C. Regulated data such as health, financial, or government records
 - D. Anything covered by a confidentiality obligation to a third party
50. What defines the Yellow tier?
- A. Credentials and secrets
 - B. Published material or anonymized data
 - C. Internal documents not meant to leave the company, anything with names or contact details, or draft material tied to a deal or product not yet announced
 - D. Regulated health, financial, or government data
51. What defines the Red tier?
- A. Regulated data (health, financial, government), credentials and secrets, or anything covered by a confidentiality obligation to a third party
 - B. Any internal document not yet published externally
 - C. Any document containing a customer's name
 - D. Anonymized survey data used for trend analysis

52. What is the redaction habit's purpose?
- A. Shortening documents so they fit within a Project's knowledge base
 - B. Converting Red-tier data into Green-tier data so it can be uploaded freely
 - C. Replacing the need to classify data before uploading it
 - D. Removing names, account numbers, or identifiers before uploading lets you get the analytical value without exposing the data, when the task's substance does not require the specifics
53. What is the example given of redaction that loses nothing the analysis needs?
- A. Asking for trends across a customer list does not require the customers' names; replacing them with 'Customer 1, Customer 2' preserves the analysis
 - B. Replacing every number in a financial report with a placeholder
 - C. Removing the dates from a time-series dataset
 - D. Deleting the outcome column from a results table
54. What is the 'partial redaction' failure mode?
- A. Redacting so much data that the task can no longer be completed
 - B. Removing the name but leaving an account number, a rare job title, or a specific date can still identify someone, especially in a small population
 - C. Redacting data that was already public
 - D. Redacting a field that was never sensitive in the first place
55. What is the 'redaction that breaks the task' failure mode?
- A. Applying redaction to Green-tier data unnecessarily
 - B. Redacting too few fields, so identification remains possible
 - C. Redacting data before confirming its sensitivity tier
 - D. If the work genuinely needs the sensitive specifics, redaction is not the answer — confirm an approved path or keep the data out entirely
56. What does code execution's sandboxed environment mean for uploaded files?
- A. Files bypass the organization's data-retention policy entirely
 - B. Files are automatically deleted after each session
 - C. Files are processed there, so you should review what you upload before running analysis
 - D. Files are shared across every Project in the account
57. What does Memory do that matters for sensitive work?
- A. It encrypts uploaded files automatically
 - B. It carries information across sessions — persistence that may be exactly what you do not want for sensitive material
 - C. It restricts a Skill's access to only what the task needs
 - D. It applies the same redaction rules as the knowledge base
58. What does Incognito mode do?
- A. Automatically classifies uploaded data by sensitivity
 - B. Exempts the session from the organization's data-retention policy
 - C. Keeps the session out of your chat history and Memory
 - D. Restricts which connectors can be used in that session

59. Does Incognito mode override an organization's underlying data-retention policy?
- A. No, but only when a connector is active in the session
 - B. Yes — Incognito sessions are excluded from all organizational data retention
 - C. Yes, but only on Team plans
 - D. No — Incognito chats still follow the organization's data-retention policies and can appear in organizational data exports
60. How do Team and Enterprise plans differ in controlling memory features at the organization level?
- A. Team plans do not have organization-level controls for memory features; on Enterprise, Owners and Primary Owners hold the org-wide Memory controls, including disabling memory for the organization
 - B. Both Team and Enterprise give every user individual control over org-wide Memory
 - C. Only Team plans allow disabling memory at the organization level
 - D. Neither plan allows disabling memory at the organization level
61. What does Incognito mode actually control, versus what it does not confirm?
- A. It controls both memory and data-retention policy simultaneously
 - B. It controls whether the data is regulated; it does not control whether it gets remembered
 - C. It controls whether something gets remembered; it does not confirm whether the data was allowed here in the first place
 - D. It controls redaction; it does not control Memory
62. For regulated data, which question comes first?
- A. Whether a Skill will process the data
 - B. Whether Incognito mode is enabled
 - C. Whether the data has been redacted
 - D. Whether the data is allowed in this entry point at all
63. What is the plain rule for handling different claude.ai entry points, according to the lesson?
- A. Memorize the technical data-retention details of every entry point
 - B. Know your data's sensitivity first, then match the feature and its controls to it
 - C. Assume all entry points handle retention identically
 - D. Use the entry point your team already uses, regardless of the data
64. In the worked example, how was a confidential M&A document for summarization handled?
- A. Yellow — needs review against policy first, and Incognito so it does not enter Memory or chat history, remaining subject to organizational data retention
 - B. Red — M&A material is always regulated data
 - C. Green — internal documents are safe once cleared for sharing
 - D. Yellow, but Incognito is unnecessary since the document is internal
65. In the worked example, how was a patient-records healthcare summary handled, and why does Incognito not solve it?
- A. It is Red data; the first question is whether the entry point is approved for protected health information at all, and if not, the correct action is to stop and escalate to an admin
 - B. It is Yellow data; Incognito mode fully resolves the sensitivity
 - C. It is Red data, but redacting patient names makes it safe to proceed
 - D. It is Green data once the records are anonymized

E. Organizational Policies and Diligence as a Habit

66. What does the lesson say is not governance?
- A. A policy that has been reviewed within the past year
 - B. A policy that is followed only when someone is watching
 - C. A policy applied consistently to routine decisions
 - D. A policy documented in a shared knowledge base
67. Where does risk live, according to the lesson?
- A. In the gap between Chat and Research mode
 - B. In the gap between two consecutive model releases
 - C. In the gap between a Project's instructions and its knowledge base
 - D. In the gap between what the policy says and what people do
68. What is Diligence, in this lesson's framing?
- A. The habit of deciding whether a use case is AI-appropriate
 - B. The habit of reviewing output against requirements, sources, and standards
 - C. The habit of applying the governance framework consistently, and of auditing real usage against it to close gaps
 - D. The habit of writing precise, testable standing instructions
69. What is the standard for applying the governance framework?
- A. Apply it on routine, low-visibility decisions, not only on the obvious high-stakes ones
 - B. Apply it only when a use case is classified inappropriate
 - C. Apply it only during scheduled quarterly reviews
 - D. Apply it only to decisions involving regulated data
70. What does auditing usage against policy involve?
- A. Reviewing every conversation before it is sent
 - B. Periodically comparing what the team does or plans to do against what the policy requires
 - C. Testing every Skill against a fixed benchmark each quarter
 - D. Comparing the team's usage against a competitor's policy
71. What counts as a Diligence gap?
- A. A Project whose knowledge base has not been updated in a month
 - B. A conversation that runs longer than expected
 - C. A data type being uploaded that should not be, a review step being skipped, or a Skill enabled without vetting
 - D. A use case classified as fully appropriate
72. Why does staying current with both policy and product capability matter?
- A. Staying current is required only for Enterprise accounts
 - B. Product capabilities never change once a feature ships
 - C. Policy updates apply only to new Projects, not existing ones
 - D. A practice that was compliant last quarter may not be after a policy update or a new feature
73. In the mini usage audit, what were the three gaps found?
- A. An unreleased product spec uploaded to a non-approved entry point, a Skill enabled without a source check, and a recurring client report that skipped its required human-review gate twice under deadline pressure
 - B. A Skill granted broad access deliberately, a client report sent without any review process, and a Project shared with an unauthorized team
 - C. A data breach, a compliance violation, and a missed model-tier requirement
 - D. A stale standing instruction, a duplicate knowledge base document, and an expired Memory entry

74. How does the lesson characterize the three gaps found in the audit?
- A. All three were deliberate policy violations requiring disciplinary action
 - B. None was malicious; all were drift
 - C. Only one was serious; the other two were negligible
 - D. All three stemmed from a single Skill misconfiguration
75. What were the fixes applied after the audit?
- A. Escalating all three gaps to Anthropic for review
 - B. Revoking the team's access to Claude until retraining is complete
 - C. Habit-level fixes — a reminder on approved entry points, a Skill-vetting step added to Project setup, and a non-negotiable review gate on client deliverables
 - D. Switching the team to a lower model tier to reduce risk

F. Ethical Implications: Bias, Fairness, Transparency

76. How does the lesson describe ethical risk?
- A. It always triggers a visible error or refusal
 - B. It only appears in use cases already classified inappropriate
 - C. It is limited to hiring and legal work
 - D. It does not announce itself; it hides in ordinary outputs
77. Which of these is given as an example of a routine output carrying ethical weight?
- A. A knowledge base holding two versions of the same document
 - B. A dashboard artifact that has outgrown Associate scope
 - C. A Skill enabled without a source check
 - D. A summary that quietly favors one group, a recommendation built on a biased framing, or an AI-assisted document presented as fully human-authored
78. Where does the lesson say evaluating for bias, fairness, and transparency belongs?
- A. In a dedicated ethics review conducted quarterly
 - B. In routine review, not a separate ethics exercise
 - C. In the Diligence audit only, not in day-to-day review
 - D. Only in use cases classified as appropriate with human review
79. Where can AI-assisted work products carry bias from, according to the lesson?
- A. The prompt, the framing, or the underlying patterns in how language is generated
 - B. The model tier selected for the task
 - C. The entry point used to produce the output
 - D. The length of the conversation before the output was generated
80. In which kind of work does the lesson say to specifically check whether output treats people fairly?
- A. Any work produced through code execution
 - B. Any work involving numeric calculations
 - C. People-facing work — hiring, evaluation, communications to specific groups
 - D. Any work stored in a shared Project

81. Where is bias and fairness risk highest, according to the lesson?
- A. Where the output is longest
 - B. Where the stakes for individuals are highest
 - C. Where the most sources were consulted
 - D. Where the task recurs most frequently
82. What does the transparency and disclosure guidance say about disclosing AI assistance?
- A. Disclosure is required only when the output is inaccurate
 - B. Disclosure is never required, since AI assistance is now standard practice
 - C. Disclosure is always required regardless of context
 - D. The obligation depends on the setting and audience; some contexts and policies require it, others treat it as routine tooling
83. What is the responsible default when you are unsure whether to disclose AI assistance?
- A. Disclose only if a regulator requests it
 - B. Conceal unless directly asked
 - C. Disclose rather than conceal
 - D. Ask Claude to decide whether disclosure is warranted
84. What structured approach does the lesson recommend for ambiguous ethical cases?
- A. Name who is affected, what could go wrong, what the fair outcome looks like, and what disclosure the situation calls for
 - B. Escalate every ambiguous case to the organization's ethics function immediately
 - C. Default to the most conservative classification available
 - D. Apply the same four Delegation criteria used for use-case screening
85. When does structured reasoning stop being enough on its own?
- A. Whenever the task involves Red-tier data
 - B. Whenever the use case has already been classified appropriate with human review
 - C. Whenever more than one person is involved in the decision
 - D. When the affected population is large, the potential harm is significant, or the question touches areas your team lacks standing to resolve
86. Why is escalating a question with documented reasoning more useful than bringing a verdict?
- A. It shows you applied the framework, identified where it ran out, and flagged the gap for the right reviewer
 - B. It removes the escalator's accountability for the outcome
 - C. It guarantees the ethics function will approve the use case
 - D. It satisfies the disclosure obligation automatically
87. In the performance-review worked example — a manager using Claude to draft summaries from their own notes — what is the lesson's conclusion?
- A. Fully appropriate, since the manager owns the underlying notes
 - B. Appropriate with human review and a fairness check, not a blanket yes or no; the reasoning, not the verdict alone, is what makes the decision defensible
 - C. Inappropriate, since performance reviews affect individual employees
 - D. Appropriate with human review, and the load-bearing criterion is reversibility

G. Quiz, Key Takeaways & Mixed Review

88. Quiz: A team proposes using Claude to produce final, unreviewed determinations on benefits eligibility. How should this use case be classified?
- A. Appropriate with light human review
 - B. Fully appropriate, since it saves time
 - C. Inappropriate: irreversible consequence and non-transferable accountability require a human to own the determination
 - D. Appropriate if a more capable model is used
89. Quiz: A colleague shares a Skill they found in a public forum that converts meeting notes into a summary. Skills do not request permissions; once enabled, a Skill runs with whatever access your session already has, and nothing in its bundled instructions limits it to meeting notes. What is the right action?
- A. Enable it and watch what it does for the first week
 - B. Enable it; a colleague vouched for it, so the source is trusted
 - C. Enable it; a Skill can only use the access its own task actually needs
 - D. Evaluate the source and think through what the Skill could reach in your session; the publisher is unknown and a Skill inherits the session's full access, so do not enable it without organizational review
90. Quiz: You need to analyze a confidential document and nothing about the session should persist. Which control fits?
- A. A more capable model
 - B. Standard Chat with Memory on
 - C. Incognito mode, after confirming the entry point is approved for this data
 - D. Upload it to a shared Project so the team can reuse it
91. Quiz: A quarterly review of a team's AI use finds that several members routinely paste draft client deliverables into a personal Claude account because the company-approved workspace feels slower to log into. What does this represent?
- A. A Diligence gap between policy and practice; the friction in the approved path is driving people to an unapproved one, and the fix is to remove that friction and make the compliant route the easy default
 - B. An acceptable workaround, since the work still gets done faster
 - C. A reason to ban the team from using Claude altogether
 - D. A problem that only the individuals who took the shortcut are responsible for
92. Quiz: A hiring coordinator uses Claude to screen incoming résumés and produce a shortlist, then forwards that shortlist to the hiring manager as 'the candidates who qualified.' The manager interviews only those candidates and never sees the ones Claude dropped. Which ethical concern is most directly raised by this use?
- A. Disclosure: the applicants were not told that Claude was used in screening
 - B. Bias: an automated screen filters people out with no human review of the exclusions, so a systematic disadvantage to some groups can go undetected
 - C. Intellectual property: the résumés may contain copyrighted material
 - D. Accountability transfer: Claude, not the coordinator, is now responsible for the hiring decision
93. According to Key Takeaway 1, what is governance?
- A. A setting an administrator configures once for the whole account
 - B. A compliance function that operates independently of daily work
 - C. A one-time certification completed at onboarding
 - D. A practitioner skill — responsible use is exercised one decision at a time, by you, not by the policy binder

94. According to Key Takeaway 2, what classifies a use case?
- A. The model tier assigned to the task
 - B. Reversibility, consequence, human element, and accountability — sorting it as appropriate, appropriate-with-review, or inappropriate
 - C. Whether a Skill has already been built for it
 - D. Whether the task has previously been performed by a human
95. According to Key Takeaway 3, what is a Skill?
- A. A feature available only on Enterprise plans
 - B. A configuration file with no ability to act on data
 - C. Software — evaluate its source and permissions before enabling, the way you would any installed software
 - D. A document that requires no more scrutiny than a knowledge base upload
96. According to Key Takeaway 4, what should happen before data enters a feature?
- A. Classify it first, then use Incognito, Memory controls, and redaction to match the handling to its sensitivity
 - B. Encrypt it manually before every upload
 - C. Assume all data is safe unless a policy names it as regulated
 - D. Redact every field regardless of what the task requires
97. According to Key Takeaway 5, where does ethical risk hide?
- A. Only in work involving regulated data
 - B. Only in use cases already classified inappropriate
 - C. Only in outputs that have not been through Discernment review
 - D. In ordinary outputs — evaluate for bias, fairness, and disclosure as part of routine review, and reason ambiguous cases through
98. Per the module's sources note, what should be treated as a 'confirm before finalizing' item rather than settled course content?
- A. The three-tier data classification system itself
 - B. Data retention behavior by entry point and any compliance-scope language such as HIPAA, FedRAMP, or ZDR
 - C. The Delegation criteria used to screen use cases
 - D. The three outcomes of a Skill trust check
99. A billing-response workflow is classified 'appropriate with human review,' with a defined gate: an agent reads each drafted response against the account record before sending. Over time, the agent starts sending drafts unedited because they are consistently accurate. What has happened to the classification?
- A. The defined gate has collapsed in practice — a review step nobody actually performs is functionally no review at all, regardless of the label still attached to it
 - B. The use case has correctly graduated to fully appropriate based on demonstrated accuracy
 - C. Nothing has changed, since the classification label is unaffected by daily practice
 - D. The load-bearing criterion has shifted from accountability to reversibility

100. Which best summarizes the throughline connecting the five lessons of Module 6?

- A. Treat governance as a compliance function separate from daily practitioner decisions
- B. Automate every use case a current model is technically capable of performing, then correct errors as they surface
- C. Screen use cases against the Delegation criteria, vet Skills like any software before granting them access, classify data by sensitivity before it enters a feature, apply policy as a sustained habit rather than a one-time acknowledgment, and evaluate routine output for bias and fairness as part of ordinary review
- D. Apply a single fixed risk tier to every use case, Skill, and dataset regardless of context

Answer Key & Explanations

1. **D** — The module opens on this stakes-setting line to frame why governance matters.
2. **A** — Any of these three triggers can cost a team the productivity gains they had made.
3. **C** — Governance is framed as something practitioners exercise in the moment, decision by decision.
4. **B** — The module names these two competencies explicitly as its anchors.
5. **A** — Diligence is about ownership and verification, as in Module 3.
6. **C** — Delegation supplies the criteria used throughout the module to screen use cases.
7. **D** — This is the module's opening summary of what responsible use actually requires day to day.
8. **B** — The lesson's opening line frames appropriateness screening as structured, not intuitive.
9. **C** — The four criteria apply at the use-case level, not only at the workflow-step level.
10. **D** — Reversibility asks whether a mistake can be caught and undone.
11. **B** — Consequence of error asks about the cost of being wrong.
12. **A** — This criterion asks whether the work needs something AI cannot supply, independent of capability.
13. **D** — Accountability asks who is answerable, and whether that answerability can attach to an AI-produced result.
14. **C** — The four criteria interact rather than functioning as an all-or-nothing checklist.
15. **B** — The practical rule is to run all four and then identify the one doing the real work in this case.
16. **A** — The load-bearing criterion is defined by what would happen if it changed.
17. **B** — These are the module's three use-case classifications, each with documented rationale.
18. **D** — Fully appropriate use cases clear all four criteria at a low level.
19. **C** — Inappropriate use cases require an articulated reason and a named human owner.
20. **A** — The middle classification is the one people leave unfinished by not defining the gate.
21. **B** — A defined gate answers who, what, and when — not just that review will occur.
22. **C** — This example names the role, the specific risk checked, and the timing — the who/what/when form.
23. **D** — An unspecified gate means the classification is unfinished, not that the risk has been handled.
24. **A** — All four criteria sit low for this case, so it delegates with normal review.
25. **D** — Fairness sensitivity and hiring accountability are what push this into the review-gated category.
26. **A** — The determination itself carries a professional accountability that a tool cannot hold.
27. **B** — The company's ownership of what it tells a customer is what makes this case appropriate with human review rather than fully appropriate.
28. **C** — Treating a Skill as software, not a document, is the premise of the whole lesson.
29. **D** — The lesson rejects both extremes and prescribes a repeatable evaluation instead.
30. **A** — These three questions make up the repeatable trust check.
31. **B** — There is no permission request to inspect; the access is inherited from the session.
32. **C** — This reframes the question around actual session access rather than a nonexistent permission request.
33. **D** — The hardest case is the internal Skill that feels safe but was never actually vetted against current policy.
34. **C** — The 'sister department' framing is the lesson's stated way to treat an unvetted internal Skill.
35. **A** — This single question is the lesson's summary of the whole evaluation.
36. **B** — These are the three named outcomes the judgment sorts into.
37. **C** — All three checks clearing is what supports outright enabling.
38. **D** — Escalation is for the genuinely ambiguous case: useful, but unresolved on source or permissions.
39. **A** — Decline is reserved for cases where further review would not resolve the disproportionate risk.

- 40. **B** — All three trust questions clear cleanly for this Skill.
- 41. **B** — This is the escalate case: unresolved source plus broad access for the stated purpose.
- 42. **C** — The install-on-a-work-machine framing is the generalized version of the Skill rule.
- 43. **D** — The same three-question pattern from Skill vetting applies to any capability you turn on.
- 44. **A** — Least privilege, revisited as the job changes, is the closing principle of the lesson.
- 45. **B** — The lesson opens on this rule: sensitivity comes before anything else.
- 46. **C** — Classification is framed as an ongoing habit, not a one-time setup step.
- 47. **D** — Controls are the action taken once classification has already decided what the data needs.
- 48. **A** — Asking first is the stated default when the tier is unclear.
- 49. **B** — Green is the safe-to-use tier: public, anonymized, or already-cleared material.
- 50. **C** — Yellow is the review-first tier: internal, identifiable, or not-yet-announced material.
- 51. **A** — Red is the keep-out-unless-approved tier: regulated data, secrets, and third-party confidentiality obligations.
- 52. **D** — Redaction preserves analytical value while removing the exposure, specifically when the specifics are not needed.
- 53. **A** — Trend analysis on a customer list is the lesson's example of redaction that costs nothing analytically.
- 54. **B** — Partial redaction leaves enough identifying detail that re-identification is still possible.
- 55. **D** — When the task depends on the specifics, redacting them just breaks the task instead of protecting it.
- 56. **C** — The sandbox is where uploaded files get processed, which is why review happens before upload.
- 57. **B** — Cross-session persistence is precisely the property that makes Memory a concern for sensitive work.
- 58. **C** — Incognito's function is specific: chat history and Memory exclusion, nothing broader.
- 59. **D** — Memory exclusion and data retention are explicitly named as separate controls.
- 60. **A** — This is a plan-specific distinction worth knowing which plan you are on.
- 61. **C** — This distinction is the direct explanation for why Incognito does not make Red data safe.
- 62. **D** — For regulated data, permission to use the entry point is settled before any handling question.
- 63. **B** — You do not need every technical detail memorized — you need the habit of classifying first.
- 64. **A** — This is the Yellow case where Incognito keeps the analysis out of persistent memory without exempting it from retention policy.
- 65. **A** — Patient records are the case Incognito cannot solve: the entry-point question has to be settled before any handling question.
- 66. **B** — The lesson's opening line names selective compliance as the failure to avoid.
- 67. **D** — The policy-practice gap is named directly as where risk accumulates.
- 68. **C** — This lesson's Diligence is specifically about consistent application plus auditing, distinct from Module 3's output-review Diligence.
- 69. **A** — Routine decisions are exactly where the standard is easiest to let slip, which is why it is named explicitly.
- 70. **B** — The audit is a periodic comparison of practice against policy, not a review of every individual conversation.
- 71. **C** — These three examples are named directly as what a gap between policy and practice looks like.
- 72. **D** — Both policy and capability evolve, so judgment has to keep pace with both.
- 73. **A** — These are the three specific gaps the team lead's audit surfaced.
- 74. **B** — The audit's finding is framed as drift, not intent — which shapes the kind of fix that follows.
- 75. **C** — The fixes match the drift diagnosis: habit-level corrections, not punitive measures.
- 76. **D** — The lesson's opening line frames ethical risk as hidden in routine outputs, not flagged automatically.
- 77. **D** — These three examples show how ethical weight can sit inside ordinary, unremarkable-looking output.

78. **B** — Ethical evaluation is folded into ordinary review rather than treated as a separate process.
79. **A** — Three named sources of bias: the prompt, the framing, and generation patterns.
80. **C** — People-facing work is called out as the category where fairness checks matter most.
81. **B** — The lesson ties risk level directly to individual stakes, not to output characteristics.
82. **D** — Disclosure obligations are context-dependent rather than fixed in either direction.
83. **C** — The stated default under uncertainty favors disclosure.
84. **A** — This four-part structure is the lesson's tool for reasoning through cases with no clean rule.
85. **D** — These three conditions mark when individual judgment needs to hand off to an organizational function.
86. **A** — Documented reasoning is what makes the escalation useful to the reviewer, not just a handoff of the problem.
87. **B** — The example works through who is affected, the risk, and the fair outcome to reach a reasoned, not reflexive, classification.
88. **C** — Benefits determinations carry the same irreversible, non-transferable accountability as the medical/legal determination case.
89. **D** — An unknown publisher plus inherited session access is the escalate case, not the enable case.
90. **C** — Incognito addresses persistence, but the entry-point approval question still has to be confirmed first.
91. **A** — This mirrors the mini-audit pattern: unintentional drift caused by friction, fixed at the habit level.
92. **B** — The unreviewed exclusions are the direct fairness risk, since nobody checks who got filtered out or why.
93. **D** — Takeaway 1: governance is a practitioner skill.
94. **B** — Takeaway 2: screen use cases with the Delegation criteria.
95. **C** — Takeaway 3: a Skill is software.
96. **A** — Takeaway 4: know data sensitivity before it enters a feature.
97. **D** — Takeaway 5: ethical risk hides in ordinary outputs.
98. **B** — The sources note flags entry-point retention behavior and compliance-scope language as items to confirm with the Anthropic account team.
99. **A** — A gate that exists on paper but is not actually performed is the same failure the module warns about: the classification only holds if the gate is real.
100. **C** — Each lesson supplies one part of the same practitioner judgment: screen the use case, vet the tool, classify the data, sustain the policy, and watch for ethical risk in ordinary output.