

# Claude Certified Associate – Foundations

(CCAO-F) — Module 1 Practice Exam

## Claude Platform & Model Foundations

100 scenario-style and recall questions, multiple choice (A/B/C/D), one correct answer each.

Covers: How Claude Behaves · Core Entry Points (Chat, Projects, Artifacts, Research) · Capability Layer (Skills, Code Execution, Memory) · Choosing Models (Haiku, Sonnet, Opus) · Context Management.

Instructions: Select the single best answer for each question. An answer key with brief explanations is provided at the end of this document. This practice material is based on user-provided course screenshots and is intended for self-study only — it is not an official Anthropic exam product.

## A. How Claude Behaves

1. You ask Claude the exact same question twice in separate conversations and receive two different answers. What does this demonstrate?
  - A. Claude is malfunctioning and needs to be reset
  - B. Outputs are produced probabilistically rather than retrieved from a fixed database, so variation is expected
  - C. One of the two answers must be factually wrong
  - D. The context window has been exceeded
2. A fabricated statistic and a verified statistic are both written by Claude in the same confident, fluent tone. What does this illustrate?
  - A. Claude only fabricates information when explicitly asked to
  - B. Fluency signals nothing about accuracy — confident tone does not indicate correctness
  - C. Claude flags fabricated statistics automatically
  - D. This only happens with the Haiku model tier
3. What is the recommended safeguard against Claude's fluent-but-not-necessarily-accurate output style?
  - A. Always use the Opus model tier
  - B. Build a verification/review habit into the workflow
  - C. Avoid using Claude for any factual tasks
  - D. Ask Claude to repeat the answer twice
4. As a long conversation approaches its working-memory (context window) limit, what does claude.ai do automatically on paid plans with Code Execution enabled?
  - A. It ends the conversation and starts a new one
  - B. It deletes the oldest messages permanently
  - C. It automatically summarizes earlier messages so the session can continue
  - D. It switches to a smaller model tier
5. A user gives Claude detailed instructions in the first 10 minutes of a 2-hour session. Ninety minutes later, responses stop reflecting that detail. What is the most likely cause?
  - A. Claude intentionally ignored the instructions
  - B. The account's usage limit was reached
  - C. Earlier context was compressed through summarization as the context window filled, causing detail to lose force
  - D. The user must have switched Projects mid-session
6. Claude is asked about an event that occurred after its training data cutoff. What is true by default?
  - A. Claude will always know about it because it has real-time internet access by default
  - B. The event is outside Claude's reliable knowledge unless a current source is connected (e.g., web search or Research)
  - C. Claude will refuse to answer any question about recent events
  - D. This only affects the Haiku model tier

**7.** For questions involving recent events, current regulations, or up-to-date market data, what should a user do?

- A. Assume Claude's default knowledge is always current
- B. Turn on web search in chat, use Research, or otherwise verify/connect a current source
- C. Switch to the Opus model, which has no knowledge cutoff
- D. Nothing — Claude auto-detects staleness and refuses to answer

**8.** A Skill is configured to generate a weekly report in a consistent format. What is still true about its output?

- A. The output will be identical, word-for-word, every time it runs
- B. It will still produce different prose each time it is run, and review stays part of the workflow
- C. It eliminates the need for any output review going forward
- D. It converts Claude's underlying model to a deterministic system

**9.** Which statement best summarizes the relationship between Skills and output variation?

- A. Skills eliminate variation entirely
- B. Skills increase variation to improve creativity
- C. Skills reduce variation but do not remove the need to check output
- D. Skills only affect variation on the Opus tier

**10.** Which of the following is NOT one of the five foundational behavioral properties of generative AI covered in the module?

- A. Outputs are probabilistic, not retrieved from a fixed database
- B. Fluency does not indicate accuracy
- C. Every conversation has a working-memory (context) limit
- D. Claude automatically bills usage at a fixed hourly rate

**11.** Why does understanding these five behavioral properties matter before choosing entry points and models?

- A. It doesn't matter — entry point selection is unrelated to these properties
- B. Understanding them prevents the most common sources of frustration and misaligned expectations
- C. It is only relevant for developers using the API
- D. These properties only apply to the free plan

**12.** A user notices Claude gives a highly confident answer that turns out to be fabricated. What should the user conclude?

- A. Claude never fabricates information when it sounds confident
- B. This is a rare bug specific to that session
- C. Confidence of tone is not a reliable signal of correctness, so verification habits matter
- D. Only Haiku produces fabricated content

## B. Core Entry Points

- 13.** Which entry point is best suited for a one-off question with no plan to reuse the context?
- A. Chat
  - B. Projects
  - C. Artifacts
  - D. Research
- 14.** Which entry point is best suited for recurring work that has stable context requirements?
- A. Chat
  - B. Projects
  - C. Artifacts
  - D. Research
- 15.** A user needs to produce a deliverable that a recipient will open and read (e.g., a polished document). Which entry point fits best?
- A. Chat
  - B. Projects
  - C. Artifacts
  - D. Research
- 16.** A task requires deep multi-source investigation and synthesis. Which entry point is appropriate (as opposed to a quick current-information lookup, which fits web search in Chat)?
- A. Chat
  - B. Projects
  - C. Artifacts
  - D. Research
- 17.** What are the three things a Project holds, according to the course material?
- A. Chat history, billing settings, and model tier
  - B. Standing instructions, a knowledge base, and its own conversation history
  - C. A Skill, a Code Execution sandbox, and Memory
  - D. A prompt template, a export button, and a public link
- 18.** Within a single Project, how do the individual conversations relate to one another?
- A. They fully share context with each other automatically
  - B. They share the Project's instructions and knowledge base, but do not share context with each other
  - C. They are merged into a single conversation thread
  - D. They cannot access the Project's knowledge base individually
- 19.** What core productivity problem do Projects solve?
- A. The need to switch between different Claude accounts
  - B. Re-explaining the same background context every session
  - C. The requirement to use Code Execution for all tasks
  - D. The lack of a mobile app

**20.** According to the 'what makes a Project worth building' guidance, when does a Project save more time than it costs to build?

- A. Only when the task will be run more than 100 times
- B. When the task recurs AND the output format is inconsistent
- C. When the task recurs and the background context is stable — if two of these conditions hold
- D. Never — Projects always cost more time than they save for individual users

**21.** A project coordinator opens a brand-new Chat every Monday and re-types the same team structure, stakeholder list, and format requirements before drafting a status report. What is the main cost of this approach?

- A. Chat cannot produce status reports at all
- B. Significant recurring time is spent on context-loading every single session
- C. Chat conversations cannot be more than 5 minutes long
- D. Reports produced in Chat are always lower quality than those produced in a Project

**22.** In the worked comparison scenario, after rebuilding the weekly report workflow as a Project, what changed?

- A. Report quality significantly improved while time increased
- B. Report quality stayed the same, but the weekly setup time was eliminated, cutting session time from 40 to 25 minutes
- C. The report had to be manually reformatted every week
- D. The Project could no longer retain the stakeholder list

**23.** What are the three diagnostic questions used to test whether a Project is worth building?

- A. Does the task recur? Is the background context the same across sessions? Is the output format consistent?
- B. Is the task urgent? Is the output long? Does it involve numbers?
- C. Does the task need Code Execution? Is Memory enabled? Is a Skill active?
- D. Is the user on a paid plan? Is Research needed? Is the model Opus?

**24.** If the answer is 'yes' to two or more of the three Project-worthiness questions, what is the guidance?

- A. Use Research instead
- B. Build the Project — the setup time is typically recovered within two or three sessions
- C. Avoid building a Project because it wastes time
- D. Use Chat exclusively going forward

**25.** Claude appears in one interface but offers how many distinct working entry points?

- A. Two
- B. Three
- C. Four
- D. Five

- 26.** Why does choosing the right entry point before starting work matter?
- A. It has no real effect on outcomes
  - B. It only affects billing, not workflow quality
  - C. It determines how efficiently the session runs and whether today's context carries forward to tomorrow
  - D. It is required only for Enterprise accounts
- 27.** A knowledge base within a Project is best described as:
- A. A place to store standing instructions only
  - B. Documents, policies, and reference files uploaded once so Claude can draw on them without re-uploading each session
  - C. A separate chat history synced across all Projects
  - D. An automatically generated summary of all past conversations
- 28.** 'Standing instructions' in a Project refer to:
- A. One-time instructions used only in the very first conversation
  - B. What Claude should know and do consistently across every conversation in that Project space
  - C. Instructions that only apply to Code Execution tasks
  - D. A list of files uploaded to the knowledge base
- 29.** Which entry-point choice would be inappropriate for a quick current-information lookup, such as checking today's exchange rate?
- A. Web search in Chat
  - B. Full Research mode for a single quick fact
  - C. Either could technically work, but Research is overkill for a single quick fact
  - D. None of the above — quick lookups are never appropriate for Claude
- 30.** A recurring analytical workflow with stable background context but where each output is a client-facing deliverable would most likely combine which two entry-point elements?
- A. Chat only, repeated every session
  - B. A Project (for stable context) producing Artifacts (for deliverables)
  - C. Research only, since deliverables always require investigation
  - D. None — Artifacts and Projects cannot be used together
- 31.** What happens to a Project's standing instructions and knowledge base when a brand-new conversation is started within that same Project?
- A. They must be manually re-entered in the new conversation
  - B. They carry forward automatically to the new conversation
  - C. They are permanently deleted after each conversation ends
  - D. They only apply if Memory is also enabled

## C. Capability Layer

**32.** In the four-layer capability model, which layer answers 'what background knowledge and standing instructions apply to this workstream'?

- A. Skills
- B. Standing instructions / Project configuration
- C. Code Execution
- D. Memory

**33.** In the four-layer capability model, which layer answers 'how a specific task should be executed, consistently, every time'?

- A. Skills
- B. Standing instructions
- C. Code Execution
- D. Memory

**34.** In the four-layer capability model, which layer answers 'when the result must be correct, not merely plausible'?

- A. Skills
- B. Standing instructions
- C. Code Execution
- D. Memory

**35.** In the four-layer capability model, which layer answers 'relevant facts carry forward across sessions without re-entry'?

- A. Skills
- B. Standing instructions
- C. Code Execution
- D. Memory

**36.** Are the four capability layers (standing instructions, Skills, Code Execution, Memory) mutually exclusive?

- A. Yes, only one can be used per task
- B. No — they are independent and can be used in any combination based on what the task requires
- C. Yes, but only on the Opus model tier
- D. No, but only two can be combined at once

**37.** What are Skills, as defined in the course material?

- A. One-time prompts that cannot be reused
- B. Reusable procedures — Anthropic provides built-in ones for common professional tasks, and custom Skills can be added
- C. A billing feature for enterprise accounts
- D. A synonym for Projects

**38.** Where do Skills 'live' relative to Projects?

- A. Skills exist only inside a single Project and cannot be used elsewhere
- B. Skills live at the account level, not inside any one Project, and Claude can invoke them automatically inside or outside a Project
- C. Skills must be manually re-enabled in every new conversation
- D. Skills replace the need for a Project entirely

**39.** What are the two things to keep in mind about Skills, per the course material?

- A. They cost extra money, and they require Opus
- B. They reduce but do not eliminate variance, and they require a trust evaluation before enabling third-party Skills
- C. They only work with Code Execution enabled, and they expire after 30 days
- D. They cannot be used inside Projects, and they require Research mode

**40.** Before enabling a third-party Skill, what should a user do?

- A. Nothing — all Skills are automatically vetted by Anthropic
- B. Review its source and the permissions it requests
- C. Switch to Haiku for lower risk
- D. Disable Memory first

**41.** What is Code Execution, as described in the course material?

- A. A feature that lets the user write and submit their own code for Claude to review
- B. Claude's sandboxed computation environment where Claude writes and runs code internally, then returns the result
- C. A separate paid add-on unrelated to Claude's core capabilities
- D. A tool for generating images only

**42.** Why does the distinction between generated prose and Code Execution matter for numeric tasks?

- A. Prose generation is always more accurate for numbers
- B. Claude generates prose by producing the most probable next text, which can look plausible but be wrong for computation; Code Execution produces a verified result by actually running the calculation
- C. Code Execution is slower and less reliable than asking Claude to calculate directly
- D. There is no meaningful difference between the two approaches

**43.** Which of the following is NOT listed as a reason to use Code Execution instead of asking Claude to calculate directly?

- A. Any numeric output will be used or reported
- B. Data needs to be transformed or cleaned
- C. The output needs to be a chart or visualization, or a real downloadable file
- D. The task is a simple one-off creative writing request

**44.** A user needs Claude to produce an actual downloadable .xlsx file. Which capability should handle this?

- A. Memory
- B. Code Execution
- C. Standing instructions only
- D. Research



- 45.** What does Memory retain, according to the course material?
- A. Only the text of the last message sent
  - B. Work-relevant facts across sessions, removing the need to re-enter the same context each time
  - C. A full transcript of every conversation ever held
  - D. Only information entered into a Project's knowledge base
- 46.** According to Memory curation guidance, how often should stored memories be reviewed for active users?
- A. Once a year
  - B. At least once per month
  - C. Only when the account is created
  - D. Every single session
- 47.** Why is Memory curation important?
- A. Memory automatically deletes itself after 24 hours, so curation is unnecessary
  - B. A memory that was accurate last quarter and hasn't been reviewed since can be actively misleading
  - C. Uncurated Memory causes Claude to stop functioning
  - D. Curation is only relevant for Enterprise accounts
- 48.** What does 'Project-scoped Memory' mean?
- A. Memory is shared identically across all Projects
  - B. Memory contexts are kept separate for each Project, so context from one client doesn't appear in another's sessions
  - C. Memory only works if Code Execution is also enabled
  - D. Memory cannot be used within Projects at all
- 49.** What does Incognito mode do?
- A. Permanently deletes an organization's underlying data retention records
  - B. Keeps a standalone session out of Memory and chat history, for sensitive or exploratory work with confidential inputs
  - C. Automatically upgrades the user to Opus for that session
  - D. Applies retroactively to Projects created before it was enabled
- 50.** Does Incognito mode override an organization's underlying data retention policy?
- A. Yes, it always overrides organizational retention
  - B. No — it keeps the session out of Memory and chat history, but does not override the org's underlying data retention
  - C. Only on the Enterprise plan
  - D. Only when Code Execution is disabled
- 51.** As of the course material (June 2026), what is the status of importing memories from other AI platforms?
- A. A fully stable, generally available feature on all plans including Enterprise
  - B. An experimental feature, documented for Free, Pro, Max, and Team plans, but not Enterprise
  - C. Not available on any plan
  - D. Only available through the API, never in `claude.ai`

**52.** If the automatic memory-import path is not available in a user's account, what is the documented fallback?

- A. Route the facts into a Project's knowledge base instead
- B. Add the key facts to Memory manually through memory settings
- C. Wait for the feature to become available with no workaround
- D. Contact billing support to force-enable the feature

**53.** A recurring analytical workflow that needs stable background context, a consistent execution procedure, verified calculations, AND continuity of facts across sessions would most reasonably use:

- A. Only Code Execution, since numbers matter most
- B. All four capability layers together — standing instructions, a Skill, Code Execution, and Memory
- C. Only Chat, since Projects add unnecessary complexity
- D. Only Memory, since it covers all other layers automatically

**54.** In the regulatory tracking report scenario, what specifically moved to Code Execution when the analyst rebuilt her workflow in month three?

- A. The portfolio background context
- B. The numeric calculations
- C. The report output format
- D. The prior months' reports

**55.** In that same scenario, what happened to the verification step after the workflow was rebuilt using the capability layer?

- A. It was removed since the new workflow was assumed to be error-free
- B. It still ran, and no errors were found in months three through eight
- C. It was replaced entirely by the Skill's built-in accuracy guarantee
- D. It could no longer be performed because Code Execution hides its calculations

**56.** In the contract-review checkpoint exercise, which capability layer best fits the instruction: 'Flag any clause that departs from the playbook and explain the departure in one sentence'?

- A. Knowledge base
- B. Standing instructions
- C. Code Execution
- D. Memory

**57.** In that same exercise, which capability layer best fits: 'The company's standard contract playbook (acceptable terms, required clauses, risk thresholds)'?

- A. Standing instructions
- B. Knowledge base
- C. Skill
- D. Code Execution

**58.** In that same exercise, which capability layer best fits: 'A procedure that extracts flagged clauses and formats them into the company's standard review report template'?

- A. Knowledge base
- B. Skill
- C. Code Execution
- D. Standing instructions

**59.** In that same exercise, which capability layer best fits: 'A calculation of the financial exposure from a penalty clause with a variable rate structure'?

- A. Skill
- B. Standing instructions
- C. Code Execution
- D. Knowledge base

## **D. Choosing Models**

**60.** What does the model tier primarily determine, as distinguished from the capability layer?

- A. What Claude can do within an entry point
- B. How well Claude performs a task, and at what cost in speed
- C. Which entry point is available to the user
- D. Whether Memory is enabled

**61.** Which model tier is described as the fastest and most efficient in the Claude family?

- A. Opus
- B. Sonnet
- C. Haiku
- D. They are all equally fast

**62.** For which task profile is Haiku most appropriate?

- A. Complex judgment on ambiguous, multi-layered inputs
- B. Routine, structured extraction or classification at volume
- C. Most professional drafting, synthesis, and analysis
- D. Deep multi-source research synthesis

**63.** For which task profile is Sonnet most appropriate?

- A. Routine, structured extraction or classification at volume
- B. Most professional drafting, synthesis, and analysis
- C. Only mathematical calculations
- D. Only tasks requiring Research mode

**64.** For which task profile is Opus most appropriate?

- A. Routine, structured extraction at volume
- B. Simple one-off formatting tasks
- C. Complex judgment, high-stakes output, or ambiguous/multi-layered inputs
- D. Tasks where speed is the only consideration

- 65.** A legal team needs to classify 500 standard vendor agreements as low/medium/high risk using a clear, unambiguous rubric, completing the task in one session. Which model tier is most appropriate?
- A. Opus, for the highest possible accuracy regardless of cost
  - B. Sonnet, as a generic default for all tasks
  - C. Haiku, for fast and efficient structured classification at volume
  - D. The model choice is irrelevant since selection cannot be adjusted
- 66.** A financial planner needs Claude to calculate compound interest across five input variables. Which approach produces the most trustworthy result?
- A. Asking Claude to calculate it in chat and reviewing the answer manually
  - B. Uploading the variables to a Project and asking Claude to estimate the figure
  - C. Using Code Execution to run the calculation and return a verified result
  - D. Using Research to find comparable portfolio scenarios from public sources
- 67.** According to the 'usage and cost' guidance, what is the general trade-off with Opus?
- A. It produces better output on complex tasks but runs slower
  - B. It is always cheaper than Haiku on a per-call basis
  - C. It has no meaningful difference in output quality from Sonnet
  - D. It cannot be used for high-stakes tasks
- 68.** On metered or usage-budgeted plans (including API access), how does the speed-versus-capability trade-off change?
- A. It doesn't change at all — cost is identical across all tiers
  - B. It also becomes a per-call cost trade-off, since a higher-capability tier like Opus consumes more usage per call
  - C. Higher tiers always cost less because they finish tasks faster
  - D. Cost only matters for the Haiku tier
- 69.** On the standard claude.ai subscription surface (not metered/API), what should be treated as the practical proxy for cost?
- A. Efficiency (speed and usage headroom)
  - B. The exact dollar price per token
  - C. The number of Projects created
  - D. The size of the knowledge base
- 70.** Per the course material's note on model lineup accuracy, what should a learner do regarding model tier names, defaults, and switching behavior?
- A. Assume the three-tier Haiku/Sonnet/Opus frame never changes
  - B. Verify the current model lineup, default model, and automatic switching behavior in the product before relying on it, since this varies by plan and changes over time
  - C. Ignore any update notices, since certification content is always current
  - D. Only verify model names if using the API, never in claude.ai

- 71.** As explicitly noted in the course material (as of June 2026), what should be treated as a 'verify-at-delivery' item rather than exam-relevant content?
- A. The three-tier Haiku/Sonnet/Opus decision logic itself
  - B. The existence of a fourth tier above Opus and any change in Opus-tier latency characterization
  - C. The definition of Code Execution
  - D. The definition of a Project
- 72.** Which task profile description best matches Sonnet rather than Haiku or Opus?
- A. High-volume routine classification where speed is the primary requirement
  - B. General professional drafting, synthesis, and analysis representing most everyday knowledge work
  - C. A task requiring the absolute highest quality ceiling regardless of speed
  - D. A task with no clear structure and extremely high stakes
- 73.** What is the risk of using Opus for simple, high-volume, well-structured tasks?
- A. There is no risk; Opus is always the safest choice
  - B. Over-engineering routine work by using more capability (and cost/speed) than the task requires
  - C. Opus cannot process structured data at all
  - D. Opus tasks cannot be reviewed afterward
- 74.** What is the risk of using Haiku for complex, high-stakes, ambiguous analytical work?
- A. There is no risk; Haiku's speed compensates for any complexity
  - B. Under-resourcing high-stakes analysis with a tier optimized for speed over depth of judgment
  - C. Haiku cannot be used for text-based tasks at all
  - D. Haiku automatically upgrades itself to Opus when needed

## **E. Context Management**

- 75.** What determines how long Claude 'can do it well' within a single session, as distinguished from model selection and capability configuration?
- A. The Project's knowledge base size
  - B. Context management
  - C. The number of Skills enabled
  - D. The billing plan tier alone
- 76.** Which of the following is listed as a sign that a conversation needs intervention?
- A. Claude responds faster than usual
  - B. Claude stops following instructions it followed correctly earlier in the same session
  - C. The user has asked more than 10 questions
  - D. A Skill was used more than once
- 77.** Which of the following is also listed as a sign a conversation needs intervention?
- A. Responses address only the most recent exchange without reference to earlier decisions or context
  - B. The conversation has lasted exactly 10 minutes
  - C. The user switched from Chat to a Project
  - D. Code Execution was used for a calculation

**78.** What is the 'Restart' response to context degradation?

- A. Ending the account subscription
- B. Starting a new conversation; the Project's standing instructions and knowledge base carry forward automatically but the conversation thread does not
- C. Switching to a different model tier only
- D. Deleting the Project entirely

**79.** When is restarting the conversation the right call?

- A. Only when the user wants to delete all prior work
- B. When the current session has drifted beyond recovery, or when beginning a genuinely new task within the same workstream
- C. Every single time a new question is asked
- D. Only on the Haiku model tier

**80.** What is the 'Summarize' response to context degradation?

- A. Deleting the Project's knowledge base to free up space
- B. Asking Claude to produce a summary of the current state (decisions, work in progress, open questions) and pasting it into a new conversation
- C. Switching to Research mode automatically
- D. Reducing the model tier to Haiku

**81.** Why does summarizing before restarting help preserve thread continuity?

- A. It permanently merges the two conversations into one
- B. It preserves continuity without carrying a degraded context window into the next session
- C. It automatically updates the Project's standing instructions
- D. It disables the context window limit entirely

**82.** What is the 'Persist' response to context degradation?

- A. Repeating the same information every session manually
- B. Saving information that should be available across all future sessions to Memory or the Project's knowledge base
- C. Switching to Incognito mode
- D. Increasing the model tier to Opus

**83.** A consultant has worked in a single Claude session for two hours on a complex analysis. In the last 30 minutes, responses have stopped reflecting the analytical framework defined at the start. What is the most appropriate response?

- A. Increase the model tier to Opus to restore output quality
- B. Repeat the analytical framework in the current conversation to restore context
- C. Start a new conversation, paste a summary of decisions and progress, and continue from there
- D. Enable Research to give Claude access to additional external context

- 84.** A consultant works on two separate client accounts using Claude and wants to ensure context from Client A never appears in Client B sessions. What is the correct configuration?
- A. Use Incognito mode for all sessions to prevent any Memory persistence
  - B. Clear Memory manually at the start of each session before switching clients
  - C. Set up separate Projects for each client account and use Project-scoped Memory
  - D. Use a different model tier for each client to maintain separate contexts
- 85.** How is the accuracy versus volume of stored memories weighted, per the course material?
- A. Volume of stored memories matters more than accuracy
  - B. The accuracy of stored memories matters more than the volume
  - C. Neither accuracy nor volume matters once Memory is enabled
  - D. Only volume matters on paid plans
- 86.** How should Memory be treated according to the context management guidance?
- A. As a permanent, unchangeable historical archive
  - B. Like a working file — reviewed on a regular cadence, with expired entries deleted and changed facts updated
  - C. As something that only needs to be set up once and never revisited
  - D. As a feature exclusive to the Research entry point
- 87.** Regarding usage limits, what does the course material say about how Claude's limits operate?
- A. There is a single unlimited usage window on all plans
  - B. Limits operate on more than one time window: a short rolling session window, plus weekly limits on paid plans (with a separate weekly pool for the highest model tier)
  - C. Usage limits apply only to Research mode
  - D. Usage limits are identical regardless of plan and never change
- 88.** For intensive tasks that may run into a usage limit before completion, what is the more efficient approach than working around an interrupted session?
- A. Ignore the limit and hope it does not trigger
  - B. Break the large task into segments, save interim progress to the knowledge base, and restart from a summary
  - C. Always downgrade to Haiku to avoid any limits
  - D. Disable Memory to free up additional usage
- 89.** Which statement correctly distinguishes model selection/capability configuration from context management?
- A. They are the same thing described with different terms
  - B. Model selection and capability configuration determine what Claude can do in a session; context management determines how long it can do it well
  - C. Context management only matters for Research mode
  - D. Capability configuration replaces the need for context management entirely

**90.** True or False: A well-maintained Project with current knowledge base entries reduces the impact of individual session context limits.

- A. True
- B. False
- C. Only true for the Opus tier
- D. Only true if Incognito mode is also enabled

## F. Mixed & Scenario Review

**91.** An operations analyst runs a weekly inventory reconciliation: uploads a new inventory file, compares it against a reorder-threshold document that hasn't changed in a year, and produces a formatted exception report. Which configuration best fits this workflow?

- A. A new Chat conversation each week, with the threshold document and inventory file pasted in
- B. A Project with the threshold document in the knowledge base, standing instructions for the exception format, and Code Execution for the comparison calculations
- C. Research with the inventory file uploaded each week
- D. Artifacts saved from each week's session as the primary continuity mechanism

**92.** Which single guiding idea connects entry-point selection, the capability layer, and model selection?

- A. All three should always default to the same setting regardless of task
- B. Each layer of decision-making should match the specific requirements of the task at hand rather than being applied uniformly
- C. Only the model tier matters; entry points and capability layers are optional
- D. Capability layers replace the need to choose an entry point

**93.** According to the five-takeaways recap, what should be selected before writing the prompt?

- A. The model tier only
- B. The entry point
- C. The Skill only
- D. The billing plan

**94.** According to the five-takeaways recap, what do the four capability layers represent?

- A. Four names for the same underlying feature
- B. Four distinct problems: Projects carry context, Skills define procedures, Code Execution verifies computations, Memory persists continuity
- C. Four different pricing tiers
- D. Four separate entry points

**95.** According to the five-takeaways recap, what do model tiers reflect?

- A. A random naming convention with no functional difference
- B. A speed-capability trade-off, where the tier should match what the task demands
- C. A difference in the entry points available to the user
- D. A distinction that only matters for API users



**96.** According to the five-takeaways recap, what is said about context?

- A. Context is unlimited and never needs management
- B. Context is a budget; long conversations degrade as context fills, and managing it deliberately is more efficient than working around drift after it accumulates
- C. Context only matters for the Research entry point
- D. Context management is only relevant on the Haiku tier

**97.** According to the five-takeaways recap, what is said about variation?

- A. Variation is a bug that should never occur
- B. Variation is inherent to every Claude feature including configured Skills, and review of outputs is structural, not optional
- C. Variation only occurs on the free plan
- D. Variation can be permanently eliminated with the right Skill configuration

**98.** A recurring, high-stakes monthly report requires: (1) consistent background context, (2) a repeatable formatting procedure, (3) verified numeric calculations, and (4) continuity of prior report data across sessions. Which combination of tools most directly maps to these four needs, in order?

- A. Chat, Research, Artifacts, Incognito mode
- B. Project standing instructions, a Skill, Code Execution, Memory/knowledge base
- C. Opus, Sonnet, Haiku, Research
- D. Artifacts, Chat, Projects, Skills

**99.** A user complains that a Claude-generated financial projection 'looked right but the numbers were actually wrong.' What foundational lesson does this best illustrate?

- A. Claude should never be used for any numeric task
- B. Fluency doesn't indicate accuracy, and unverified numeric generation should have used Code Execution instead
- C. The user should have used the Haiku model instead of Sonnet
- D. This is a rare and unexplainable error with no underlying cause

**100.** A user wants Claude to remember their job title and standard report format across all future sessions, regardless of which Project they're working in. What is the appropriate mechanism?

- A. Project-scoped Memory only, since it applies everywhere by default
- B. Account-level Memory, since Project-scoped Memory keeps facts separate per Project
- C. Incognito mode, since it persists facts long-term
- D. Re-typing the information at the start of every session going forward

# Answer Key & Explanations

- 1. B** — Generative AI produces the most probable next sequence of text rather than retrieving a single fixed answer, so variation between runs is expected and normal.
- 2. B** — Claude writes with consistent fluency regardless of whether the answer is actually correct, so tone cannot be used as a proxy for accuracy.
- 3. B** — Because fluency doesn't signal accuracy, building review and verification into the workflow is the correct mitigation.
- 4. C** — As a conversation approaches its context limit, claude.ai automatically summarizes earlier messages (on paid plans with Code Execution enabled) so the session can continue.
- 5. C** — Summaries compress detail, which is why instructions given early in a long session can lose force later — not because Claude is deliberately ignoring them.
- 6. B** — Claude's training data has a cutoff date; information after that date is unreliable unless the user enables web search, connects a source, or uses Research.
- 7. B** — When recency/relevance matters, the user should verify or connect a current source, such as enabling web search or using Research.
- 8. B** — Skills reduce output variance but do not eliminate it. Review remains necessary regardless of how carefully a Skill is configured.
- 9. C** — Even a well-configured Skill-based workflow produces different outputs each time it's run; checking the output is still required.
- 10. D** — Billing/pricing mechanics are not one of the five behavioral properties covered; they concern probability, fluency vs. accuracy, context limits, knowledge cutoff, and Skill variance.
- 11. B** — The five properties apply to Claude regardless of which feature is used, and understanding them before building workflows prevents common frustration and misaligned expectations.
- 12. C** — A fabricated statistic reads with the same assurance as a verified one — tone alone cannot be trusted as a signal of accuracy.
- 13. A** — Per the selection logic, a one-off question or quick task with no plan to reuse belongs in Chat.
- 14. B** — Recurring work with stable context requirements is the defining case for a Project.
- 15. C** — When the output is a deliverable the recipient will open and read, Artifacts is the appropriate entry point.
- 16. D** — Deep multi-source investigation or synthesis calls for Research; quick current-info lookups can use web search directly in Chat.
- 17. B** — A Project holds standing instructions (what Claude should know/do consistently), a knowledge base (uploaded reference material), and its own conversation history.
- 18. B** — Each Project maintains its own conversation list; conversations share the Project's instructions and knowledge base but do not share context with one another.
- 19. B** — Projects solve the common productivity drain of re-explaining the same background every session by setting context once in standing instructions and the knowledge base.
- 20. C** — If the task recurs, the background context is stable, and the output format is consistent — and at least two of these conditions hold — a Project saves more time than it costs to build.
- 21. B** — In the worked comparison, about 12 of 40 minutes per session were spent purely on re-loading context that a Project would have retained automatically.

- 22. B** — The report quality was the same; the entry-point decision eliminated the weekly setup tax, reducing the session from 40 to 25 minutes.
- 23. A** — The three-question test is: Does this task recur? Is the background context the same across sessions? Is the output format consistent?
- 24. B** — If yes to two or more of the three questions, build the Project; the setup time is recovered within two or three sessions.
- 25. C** — Claude offers four distinct working entry points: Chat, Projects, Artifacts, and Research Mode.
- 26. C** — Choosing the right entry point determines how efficiently the session runs and whether the context built today carries forward to future sessions.
- 27. B** — The knowledge base holds documents, policies, and reference files uploaded once so Claude can draw on them without needing to be re-uploaded each session.
- 28. B** — Standing instructions define what Claude should know and do consistently across every conversation within that Project.
- 29. C** — The guidance distinguishes quick current-information lookups (suited to web search in Chat) from deep multi-source investigation or synthesis (suited to Research).
- 30. B** — A Project supplies the stable, recurring context while Artifacts is the appropriate output format when the result is a deliverable the recipient will open and read.
- 31. B** — A Project's standing instructions and knowledge base carry forward automatically to every conversation created within it.
- 32. B** — The first layer covers what background knowledge and standing instructions apply to the workstream — this is the Project's own configuration.
- 33. A** — Skills define how a specific task should be executed, consistently, every time it applies.
- 34. C** — Code Execution exists for cases where the result must be verifiably correct, not just plausible-sounding text.
- 35. D** — Memory persists relevant facts across sessions so they don't need to be re-entered each time.
- 36. B** — The layers are independent; a one-off question may need none of them, while a recurring analytical workflow might use all four.
- 37. B** — Skills are reusable procedures. Anthropic provides built-in Skills for tasks like creating/editing Excel, Word, PowerPoint, and PDF files, and custom Skills can be added for task-specific workflows.
- 38. B** — Skills live at the account level, not inside any one Project, and Claude invokes them automatically when relevant, whether inside or outside a Project.
- 39. B** — First, Skills reduce variance but don't eliminate it — review stays necessary. Second, Skills require a trust evaluation, since a Skill can access whatever you give Claude access to during the session.
- 40. B** — Before enabling a third-party Skill, review its source and the permissions it requests; Anthropic-provided and organization-approved Skills are the lower-risk starting point.
- 41. B** — Code Execution is Claude's sandboxed computation environment. Claude writes and runs code internally, then returns the result — the user does not write the code themselves.
- 42. B** — For computation, Claude's default text generation can produce plausible-looking but potentially inaccurate numbers, whereas Code Execution actually runs the calculation and returns a verified result.
- 43. D** — Code Execution is recommended for numeric outputs, data transformation/cleaning, charts/visualizations, and real downloadable files — not for simple creative writing requests.
- 44. B** — Producing a real downloadable file (.xlsx, .pptx, .docx, .pdf) is one of the listed reasons to use Code Execution.

- 45. B** — Memory retains work-relevant facts across sessions, removing the need to re-enter the same context each time.
- 46. B** — The guidance recommends reviewing stored memories periodically, at least once per month for active users.
- 47. B** — Memory is most useful when actively curated; stale entries that were accurate in the past can become actively misleading if not reviewed and updated.
- 48. B** — Project-scoped Memory keeps Memory contexts separate for each Project — context from Client A does not appear in Client B sessions.
- 49. B** — Incognito mode keeps a session out of Memory and chat history (for standalone chats, outside Projects); it does not override an organization's underlying data retention.
- 50. B** — The course material explicitly notes that Incognito mode does not override an organization's underlying data retention.
- 51. B** — Importing memories from other AI platforms is described as experimental as of June 2026, documented for Free, Pro, Max, and Team plans (not Enterprise).
- 52. B** — The documented fallback is to add key facts to Memory manually through memory settings, rather than routing them into a Project's knowledge base.
- 53. B** — Because the four layers are independent, a sufficiently complex recurring workflow may reasonably combine all four: standing instructions, a Skill, Code Execution, and Memory.
- 54. B** — Numeric calculations moved to Code Execution, while portfolio context and format instructions went into the Project, and prior reports went into the knowledge base.
- 55. B** — The verification step still ran even after the rebuild; no errors were found in months three through eight, illustrating that review remains part of the workflow.
- 56. B** — A rule that should apply consistently across every review in this workstream is a standing instruction.
- 57. B** — Stable reference material that Claude should draw on without re-uploading each session belongs in the knowledge base.
- 58. B** — A defined, repeatable procedure for producing a consistently formatted output is a Skill.
- 59. C** — A computation that must be verifiably correct, not merely plausible, belongs in Code Execution.
- 60. B** — The capability layer determines what Claude can do; the model tier determines how well Claude does it, and at what cost in speed.
- 61. C** — Haiku is the fastest and most efficient model in the Claude family.
- 62. B** — Haiku fits routine, structured extraction or classification at volume, where speed matters and the cost of an imperfect output is low.
- 63. B** — Sonnet is positioned for most professional drafting, synthesis, and analysis work.
- 64. C** — Opus is reserved for complex judgment, high-stakes output, and ambiguous or multi-layered inputs, where the quality ceiling genuinely matters.
- 65. C** — This is routine, structured classification at volume with an unambiguous rubric — the definition of a Haiku-appropriate task.
- 66. C** — For computation, Code Execution produces a verified result by actually running the calculation, rather than Claude generating a plausible-looking estimate.
- 67. A** — Opus produces better output on complex tasks, with the tradeoff of running slower; for most knowledge work, Sonnet handles the task well.

- 68. B** — On metered/usage-budgeted plans, a higher-capability tier such as Opus consumes more usage per call than Haiku or Sonnet, so the trade-off becomes a cost trade-off as well as a speed trade-off.
- 69. A** — On the standard claude.ai subscription surface, efficiency — speed and usage headroom — is treated as the practical proxy for cost; the selection logic itself does not change.
- 70. B** — The model picker lineup, default model, and automatic switching behavior vary by plan and change over time; the course explicitly instructs verifying all three in-product before publishing/acting on it.
- 71. B** — The course notes that a fourth tier above Opus exists in the shipped model family and Opus latency is now characterized differently, but instructs treating tier names/characteristics beyond the three-tier frame as a verify-at-delivery item, not exam-relevant.
- 72. B** — Sonnet is positioned as the tier that covers most professional drafting, synthesis, and analysis — the bulk of everyday knowledge work.
- 73. B** — Matching the tier to the task avoids over-engineering routine work with an unnecessarily high-capability, slower, costlier tier.
- 74. B** — Matching the tier to the task also avoids under-resourcing high-stakes analysis by using a tier that prioritizes speed over the depth of judgment such work requires.
- 75. B** — Model selection and capability configuration determine what Claude can do; context management determines how long it can do it well within a session.
- 76. B** — One listed sign is that Claude stops following instructions it followed correctly earlier in the same session.
- 77. A** — A listed sign is that responses address only the most recent exchange without referencing earlier decisions or context — consistent with earlier context being compressed or lost.
- 78. B** — Restarting means starting a new conversation. The Project's standing instructions and knowledge base carry forward automatically, but the conversation thread itself does not.
- 79. B** — Restarting is appropriate when the session has drifted beyond recovery or when starting a genuinely new task within the same workflow.
- 80. B** — Summarizing means asking Claude to produce a summary of the current state before starting a new conversation, then pasting that summary as context in the new conversation.
- 81. B** — Summarizing preserves thread continuity without carrying a degraded (overly compressed) context window forward into the next session.
- 82. B** — Persisting means saving information that should be available across all future sessions to Memory or updating the Project's knowledge base, rather than re-entering it repeatedly.
- 83. C** — This matches the 'Summarize then restart' pattern: the session has degraded, so producing a summary and continuing in a fresh conversation is the appropriate response.
- 84. C** — Project-scoped Memory keeps Memory contexts separate per Project; setting up separate Projects per client keeps their contexts from crossing over.
- 85. B** — The course explicitly states that the accuracy of stored memories matters more than the volume — Memory should be treated like a working file that is kept current.
- 86. B** — Memory should be treated like a working file: reviewed on a regular cadence, with expired entries deleted and changed facts updated.
- 87. B** — Claude's usage limits operate on more than one time window — a short rolling session window plus weekly limits on paid plans, with a separate weekly pool for the highest model tier.
- 88. B** — Planning ahead — breaking large tasks into segments, saving interim progress to the knowledge base, and restarting from a summary — is more efficient than working around an interrupted session.

- 89. B** — This is the explicit framing given in the course material: capability determines what can be done, context management determines how long it can be done well.
- 90. A** — The course material states that a well-maintained Project with current knowledge base entries reduces the impact of individual session context limits.
- 91. B** — This is recurring work with stable context (the threshold document) and a need for verified comparison calculations — a Project plus Code Execution is the correct combination.
- 92. B** — Across all three areas, the course material consistently teaches matching the choice (entry point, capability layer, model) to what the specific task actually requires.
- 93. B** — The first takeaway is to select the entry point before writing the prompt, since it shapes every subsequent choice in the session.
- 94. B** — The recap frames the four capability layers as solving four distinct problems: context, repeatable procedure, verified computation, and cross-session continuity.
- 95. B** — Model tiers reflect a speed-capability trade-off: Haiku for structured/high-volume, Sonnet for most professional work, Opus for complex/high-stakes work.
- 96. B** — Context is described as a budget: long conversations degrade as context fills, and deliberate management (restart, summarize, persist) is more efficient than working around drift afterward.
- 97. B** — Every Claude feature, including configured Skills, produces different outputs run to run; review before outputs leave your hands is a structural part of the workflow, not an optional extra.
- 98. B** — This maps directly onto the four-layer capability model: standing instructions/Project for background context, a Skill for repeatable formatting, Code Execution for verified numbers, and Memory or the knowledge base for continuity.
- 99. B** — This combines two module lessons: fluent output is not a reliable signal of accuracy, and Code Execution — not generated prose — should be used for verified numeric results.
- 100. B** — Because Project-scoped Memory is deliberately kept separate per Project, information meant to apply across all Projects should be stored at the account level rather than scoped to a single Project.